

facebook

Robotron: Top-down Network Management atfacebook Scale

Yu-Wei Eric Sung, Xiaozheng Tie, Starsky H.Y. Wong, Hongyi

ZengSIGCOMM 2016

August 25, 2016

Scale of Facebook Community



1.7 Billion
on Facebook Monthly



1 Billion
on Whatsapp Monthly



500 Million
on Instagram Monthly

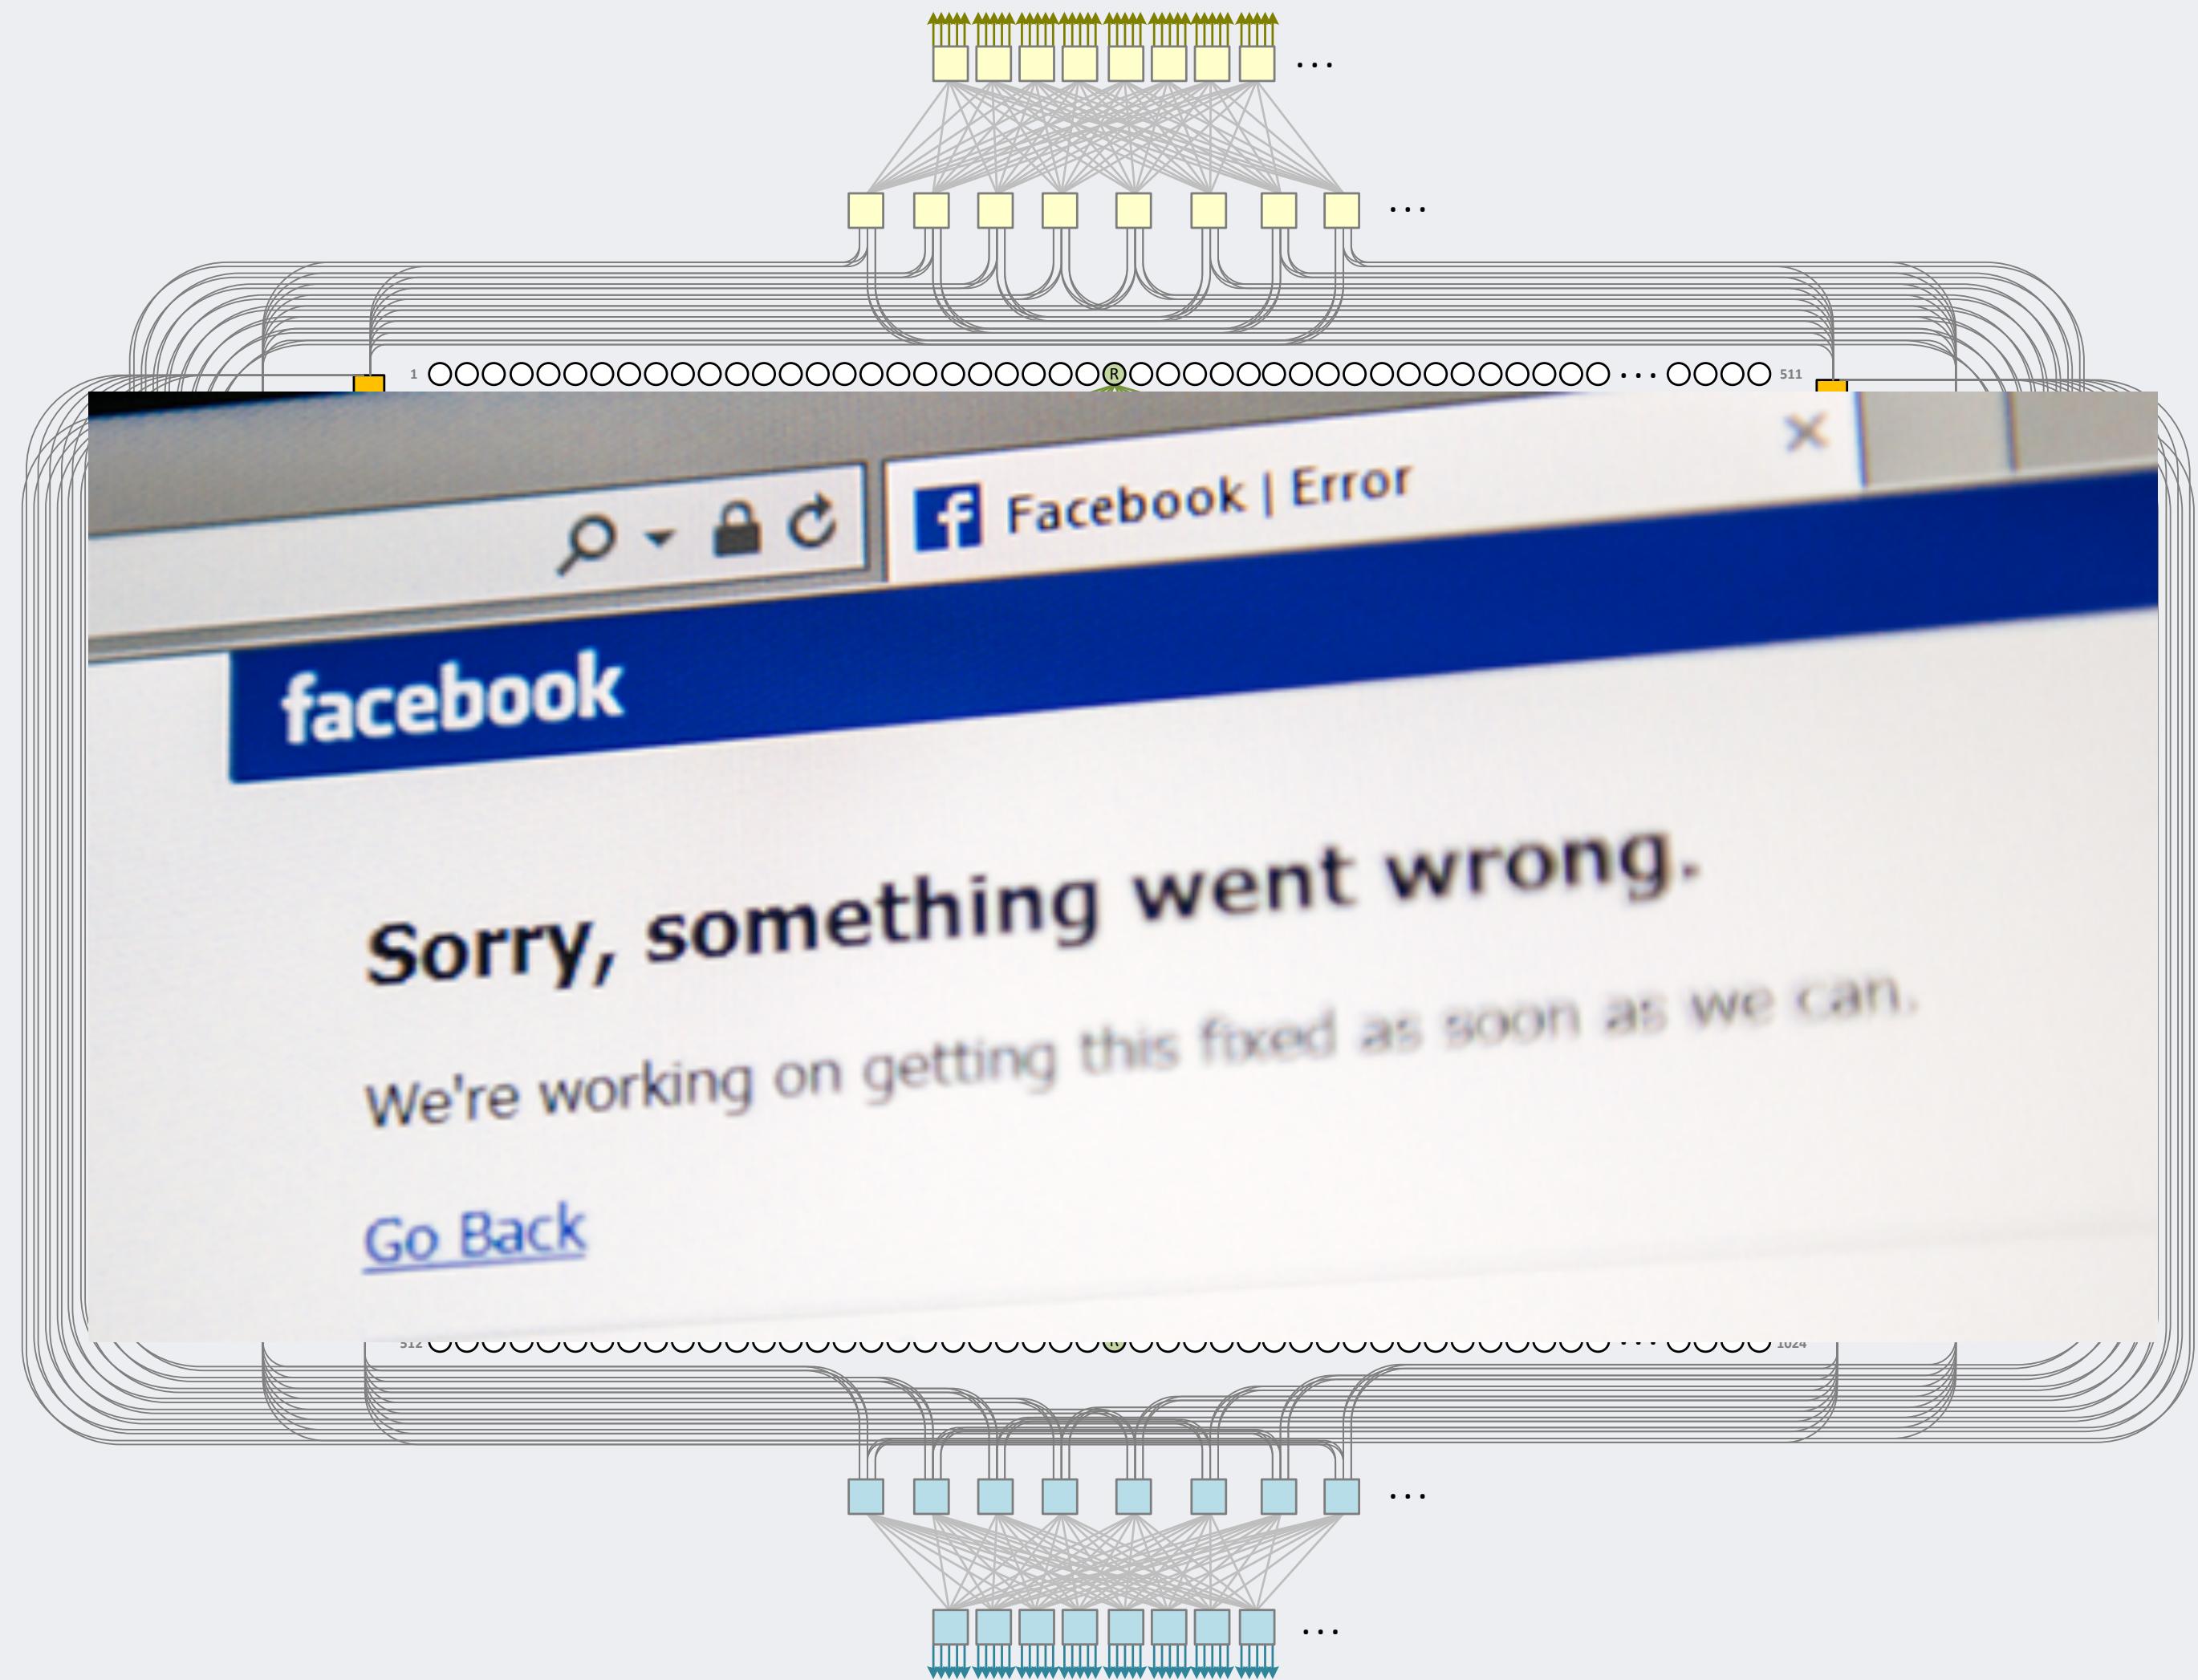


1 Billion
on Messenger Monthly

Network Management at Facebook

What's involved?

- Goals: Build and evolve FB network
- Example tasks: circuit/device turnup, network monitoring
- Human interactions -> outages



Network Management at Facebook

Why is it hard?

- Distributed Configurations
- Multiple Domains
- Versioning
- Dependency
- Vendor Differences



Network Management at Facebook

Early days...

2004-2007

2008

2009

2010

2011

2012

2013

2014

2015

Manual Configuration and
Monitoring with ad-hoc scripts

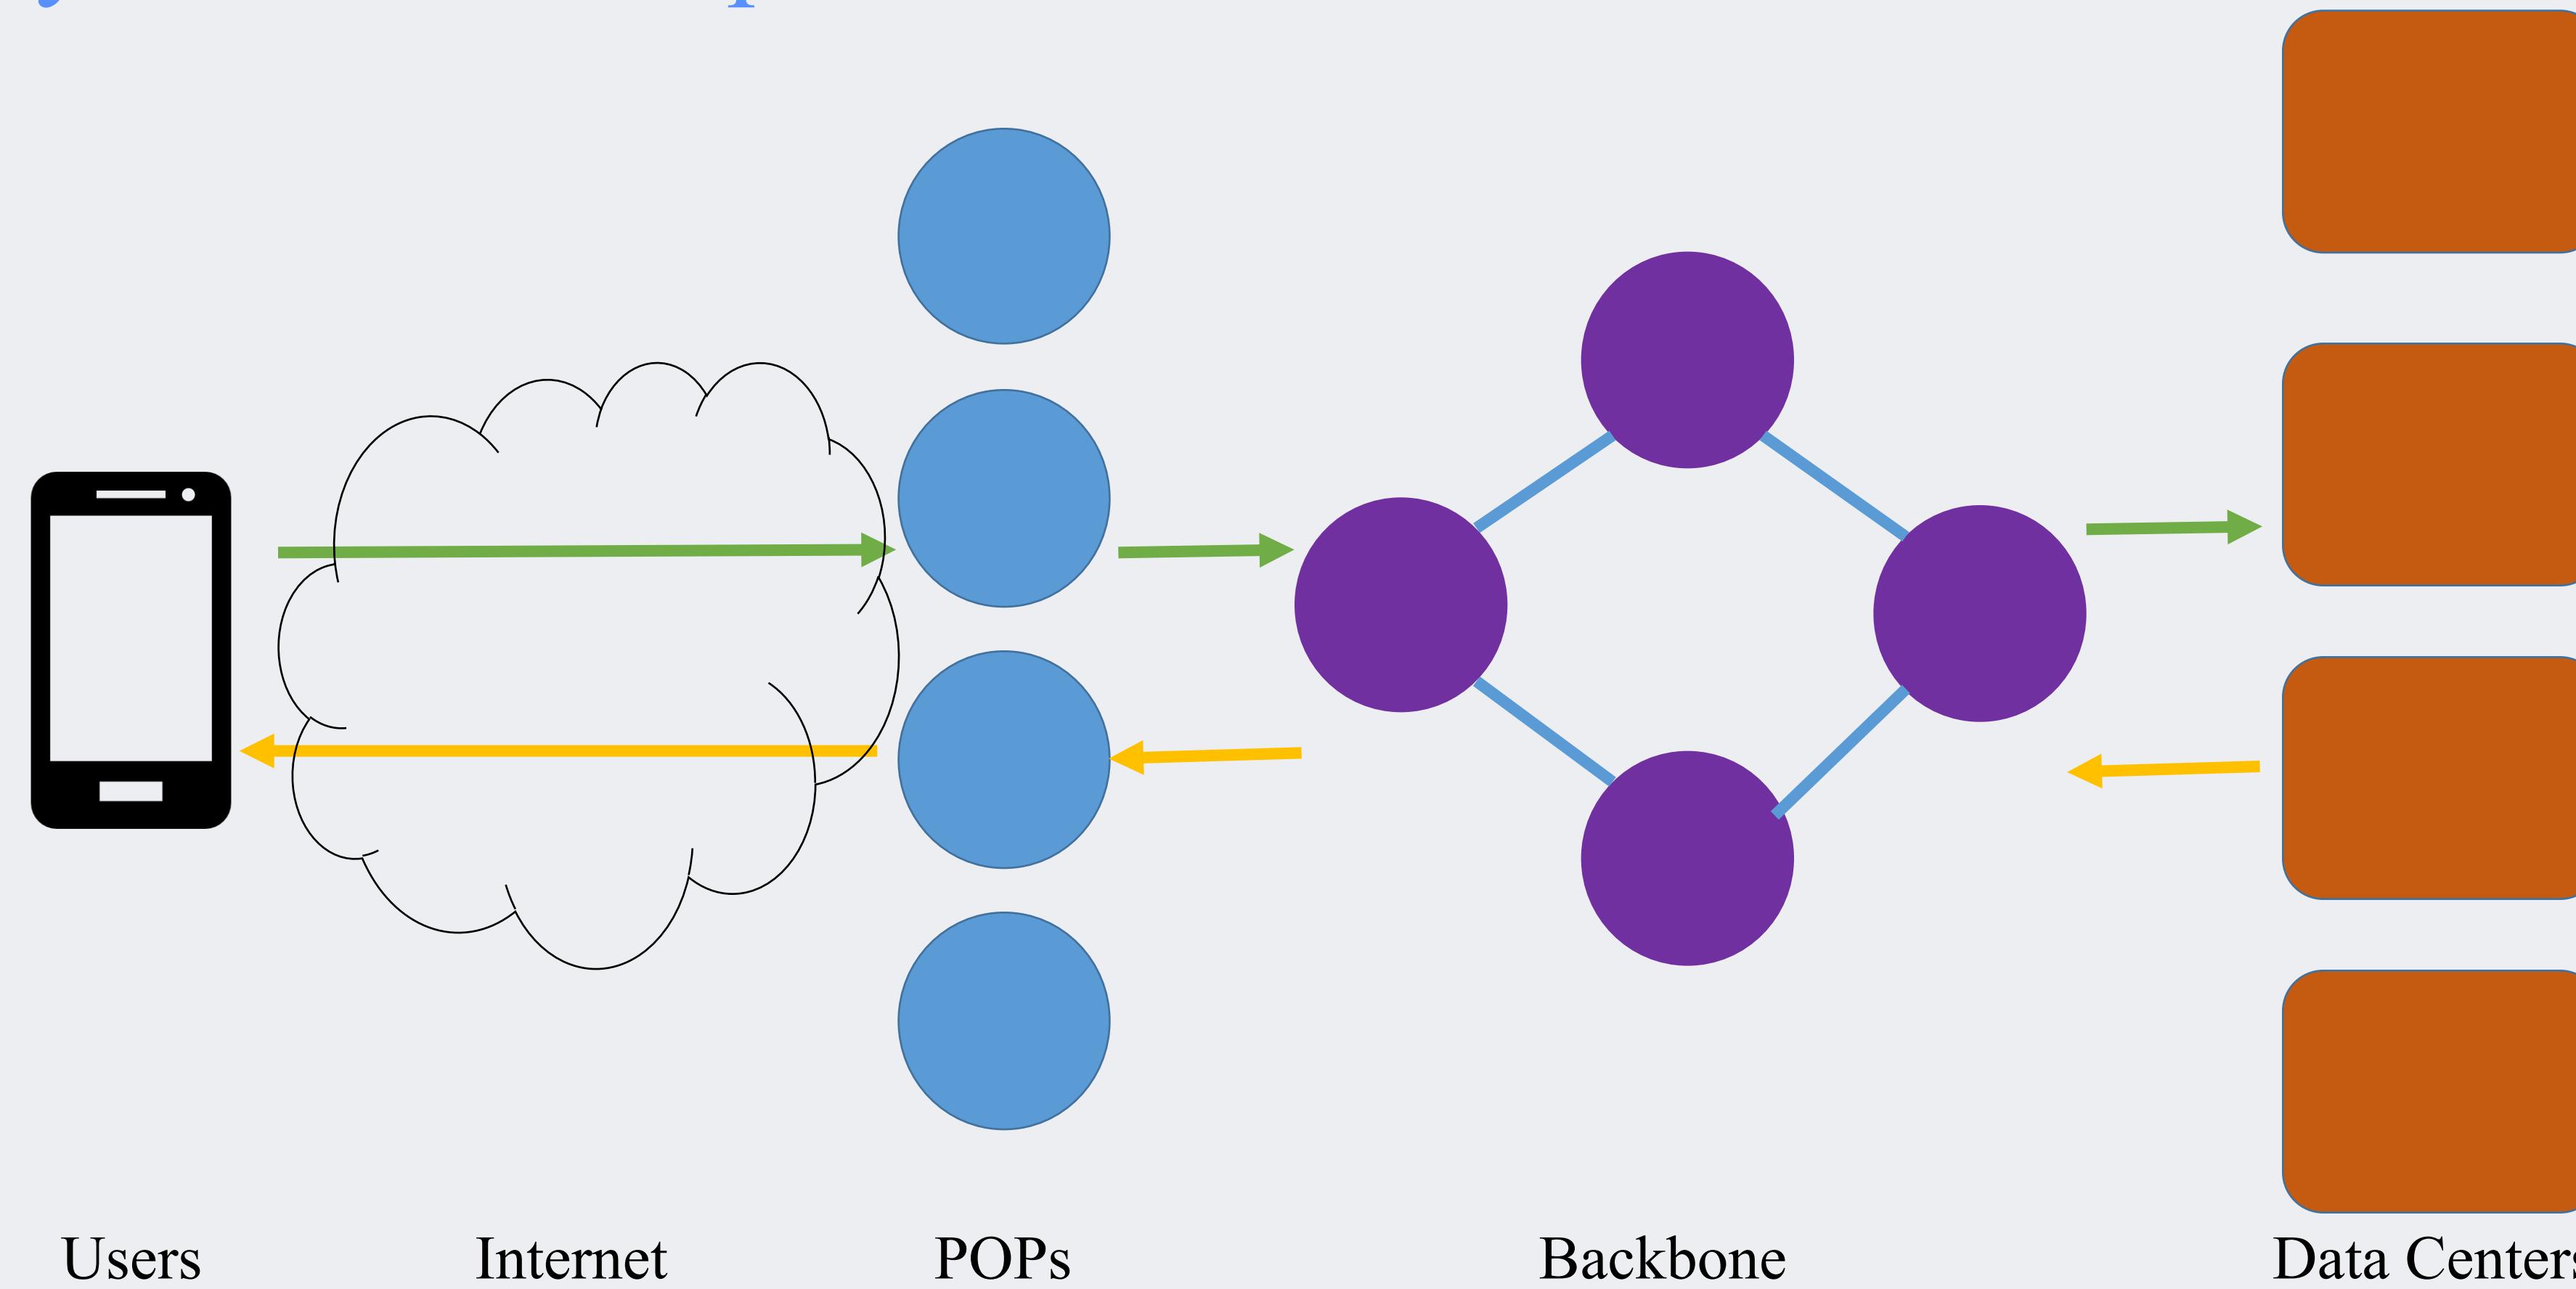
Contribution



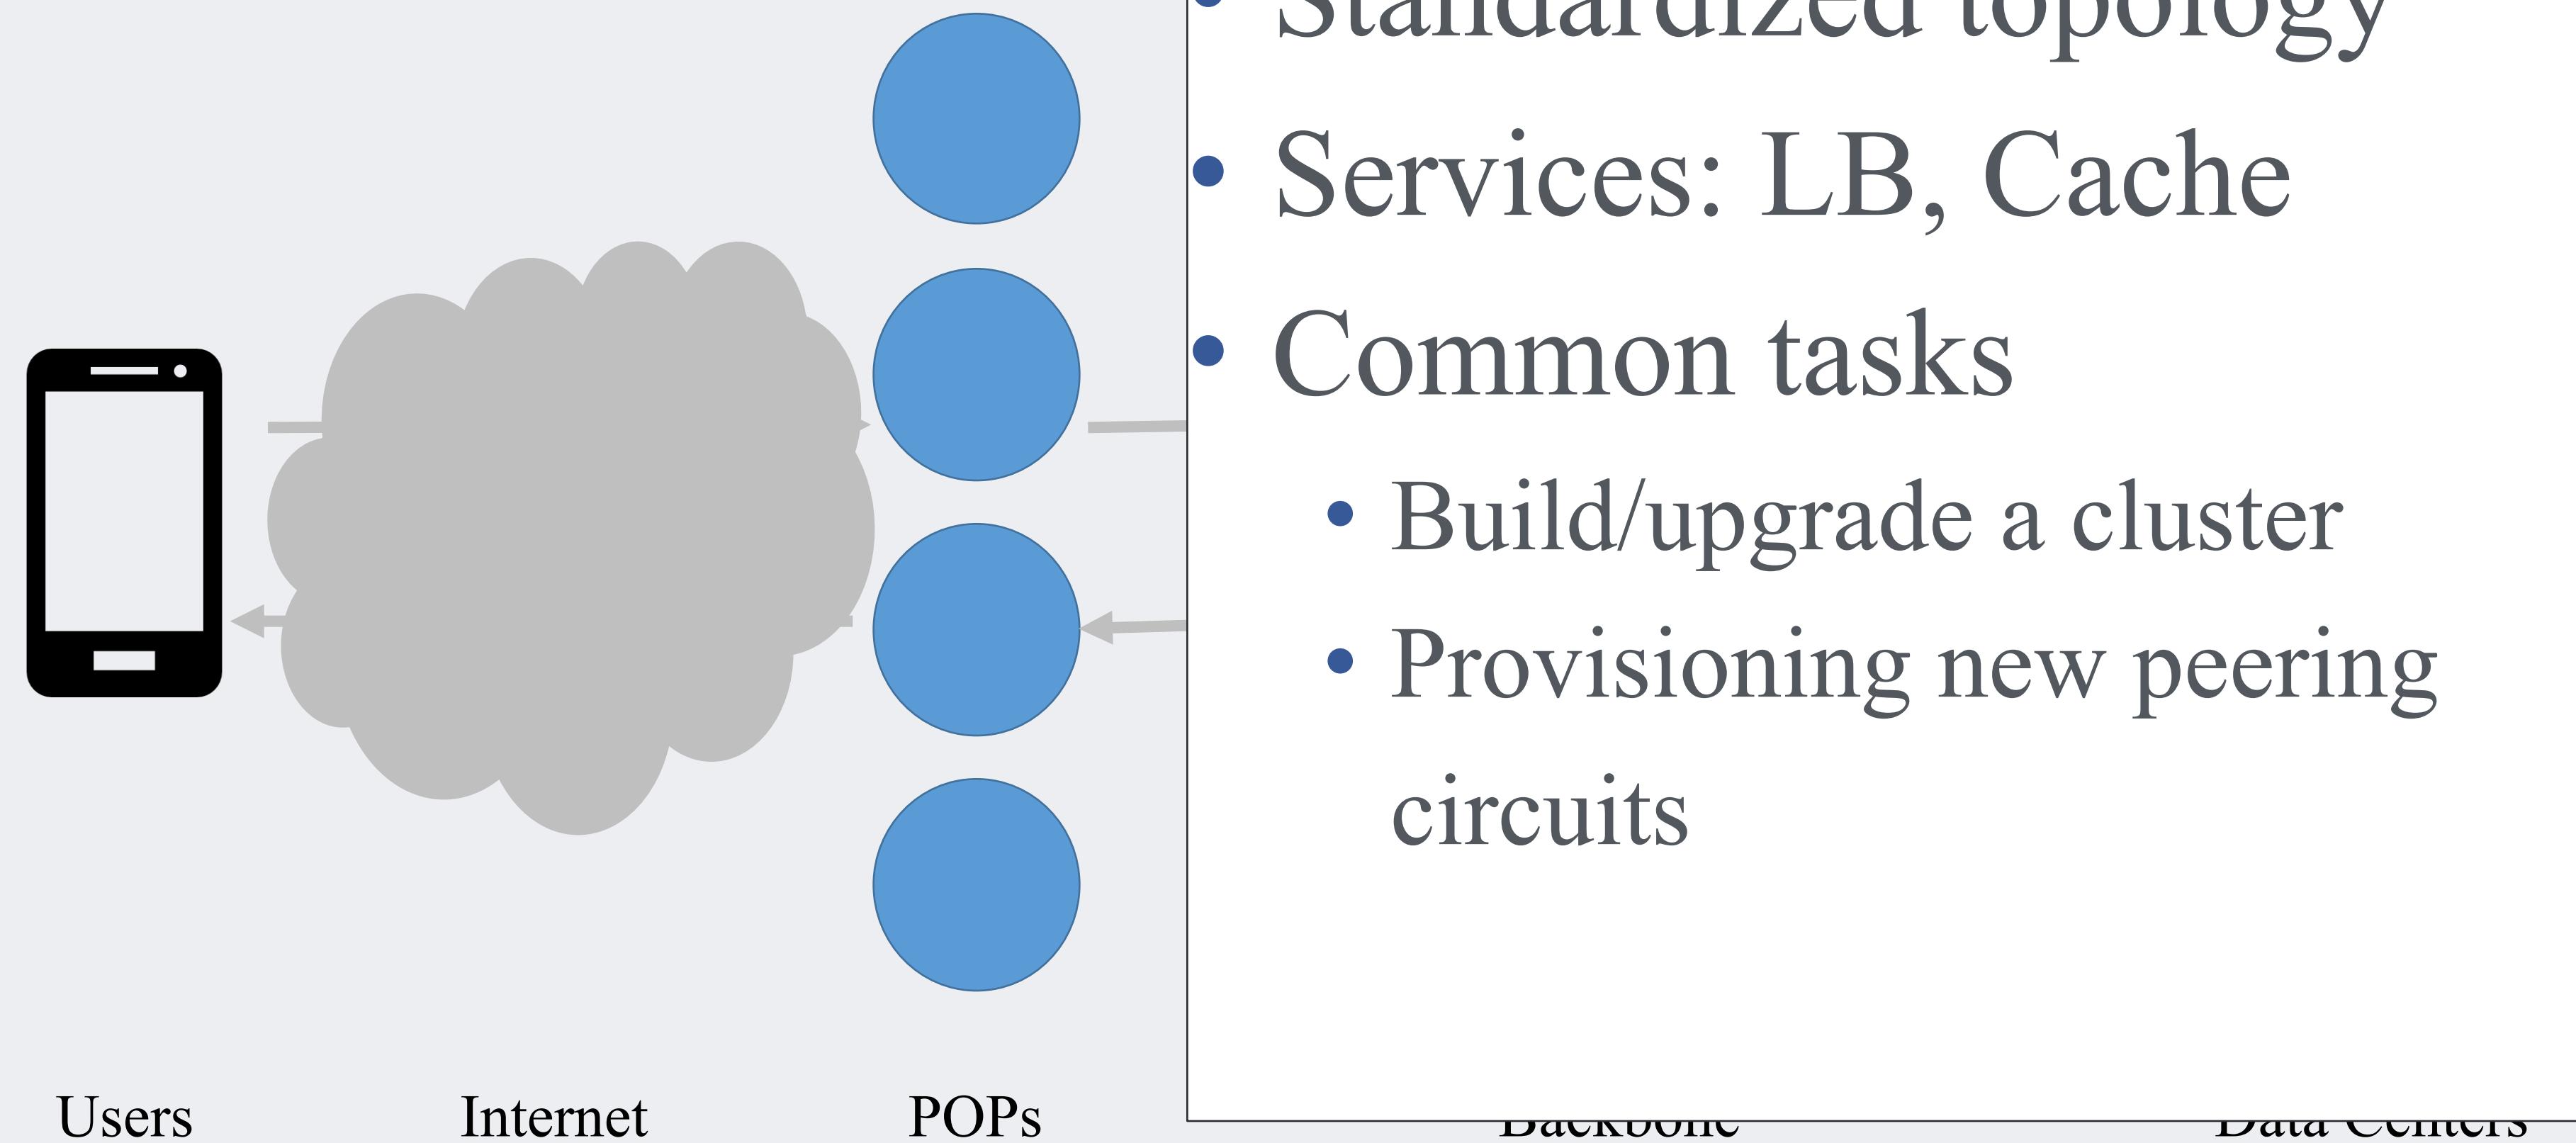
- Shed light on
 - Network management tasks
 - Robotron's usage
 - Evolution of Roboron
 - Our experiences using Robotron

Overview of Facebook's Network

Lifecycle of user requests



Point of Presence (POP)



Backbone

- Irregular, demand-driven topology
- Common tasks:
 - Add/migrate circuits
 - Add/remove routers

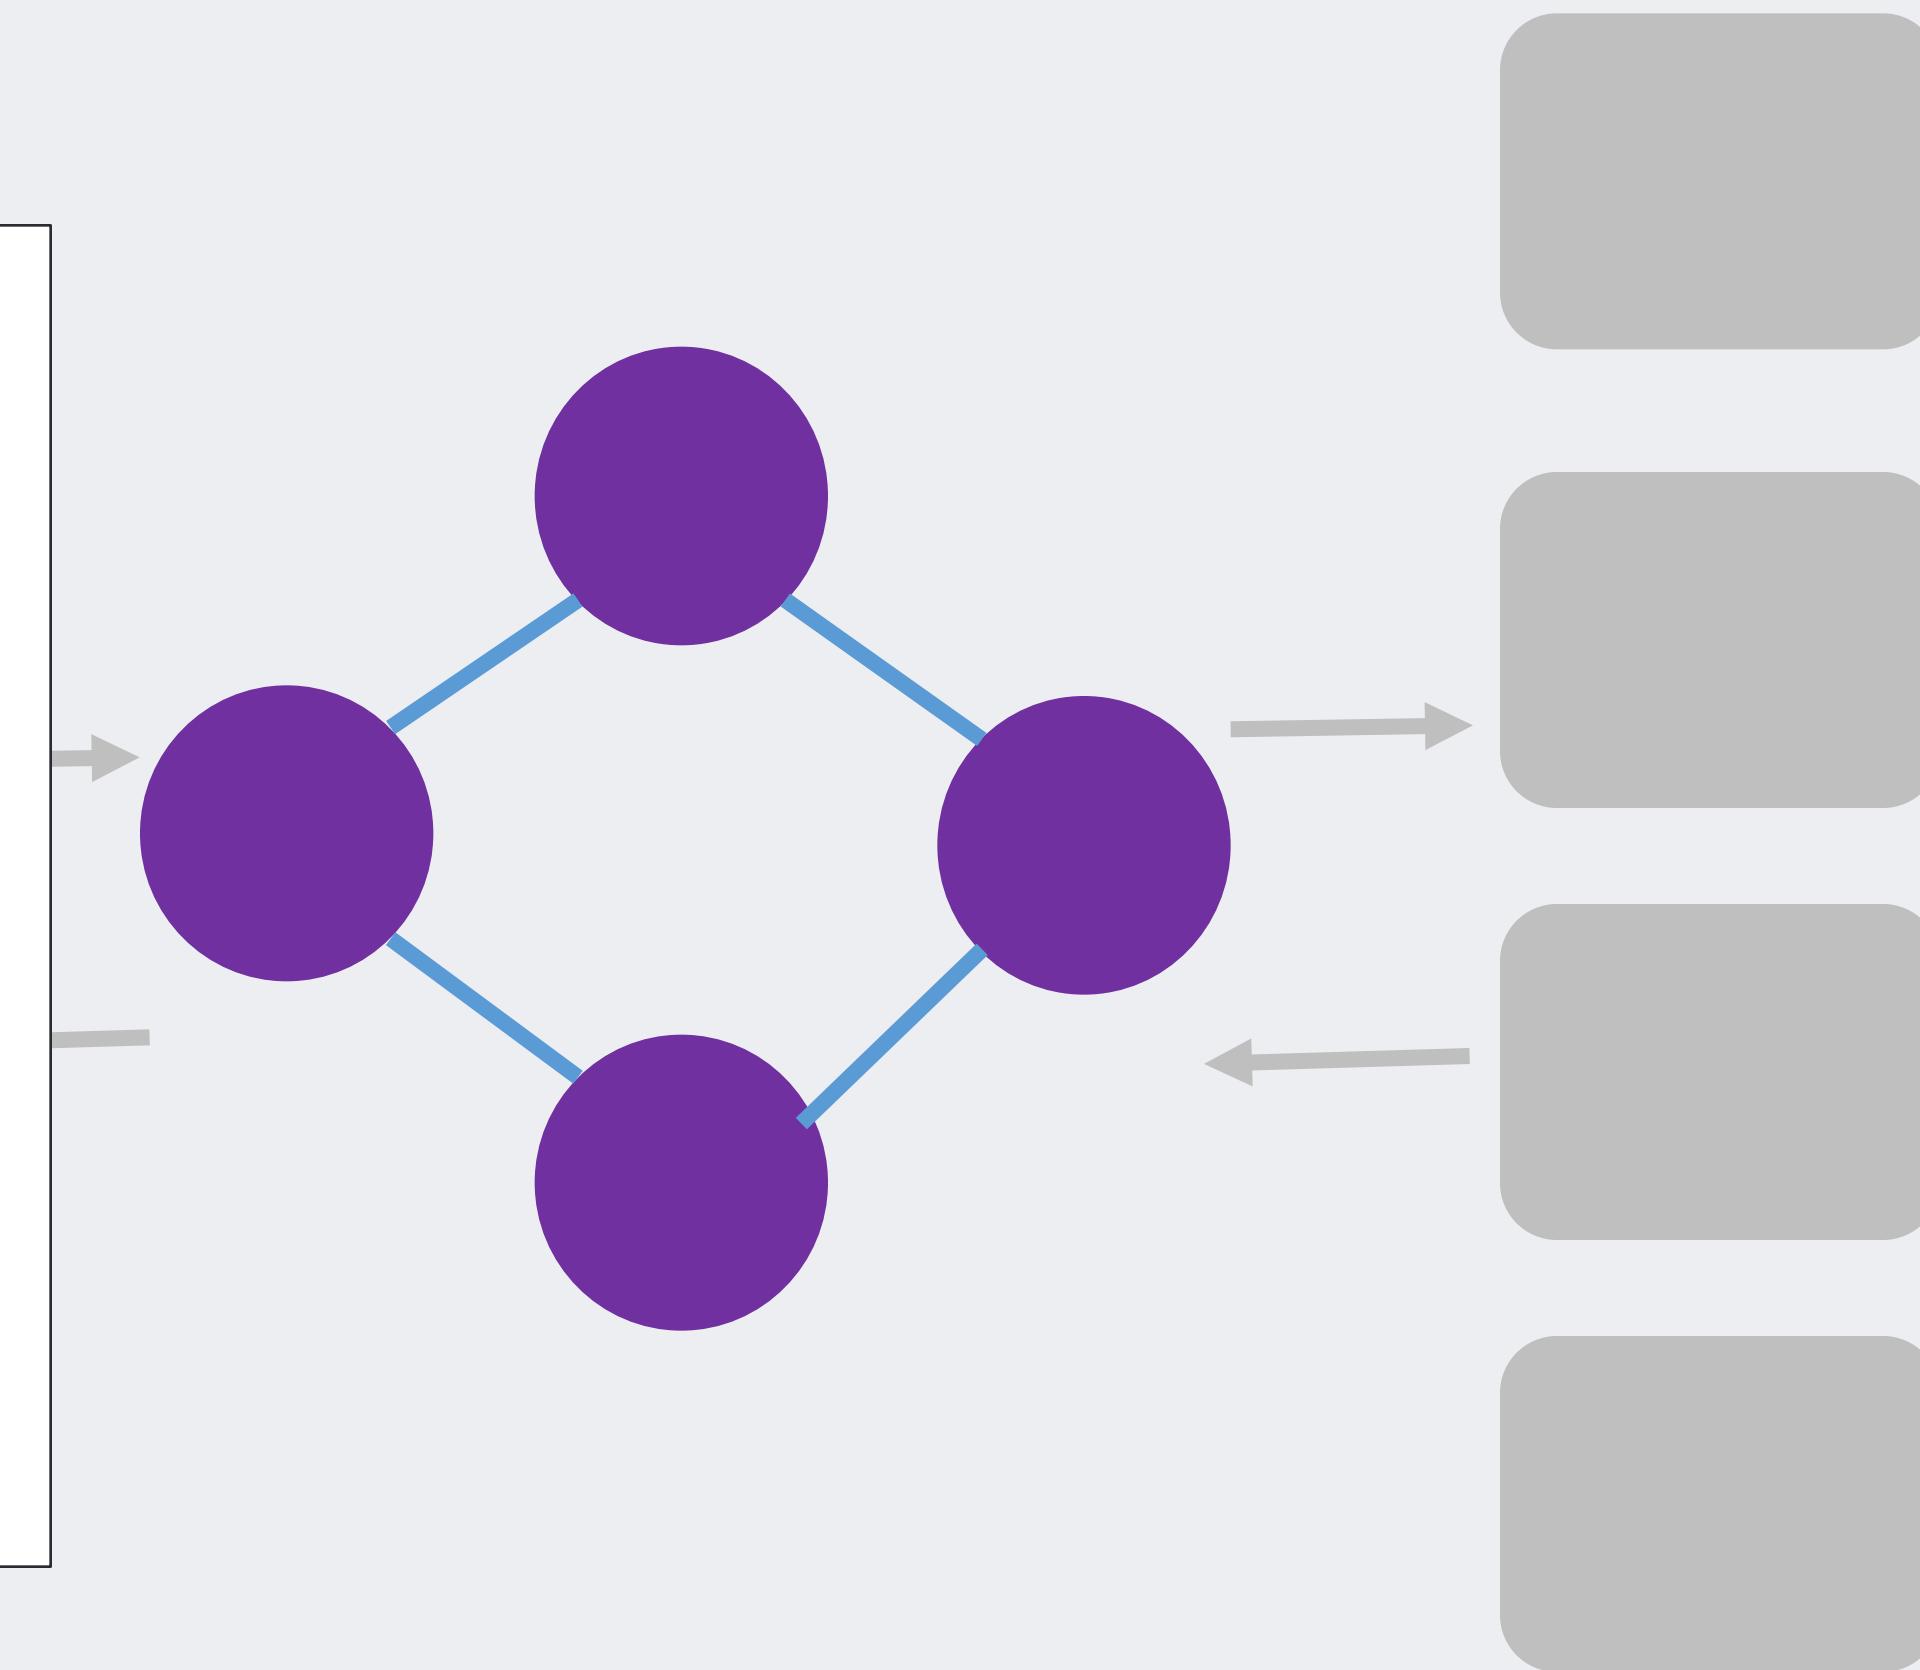
Users

Internet

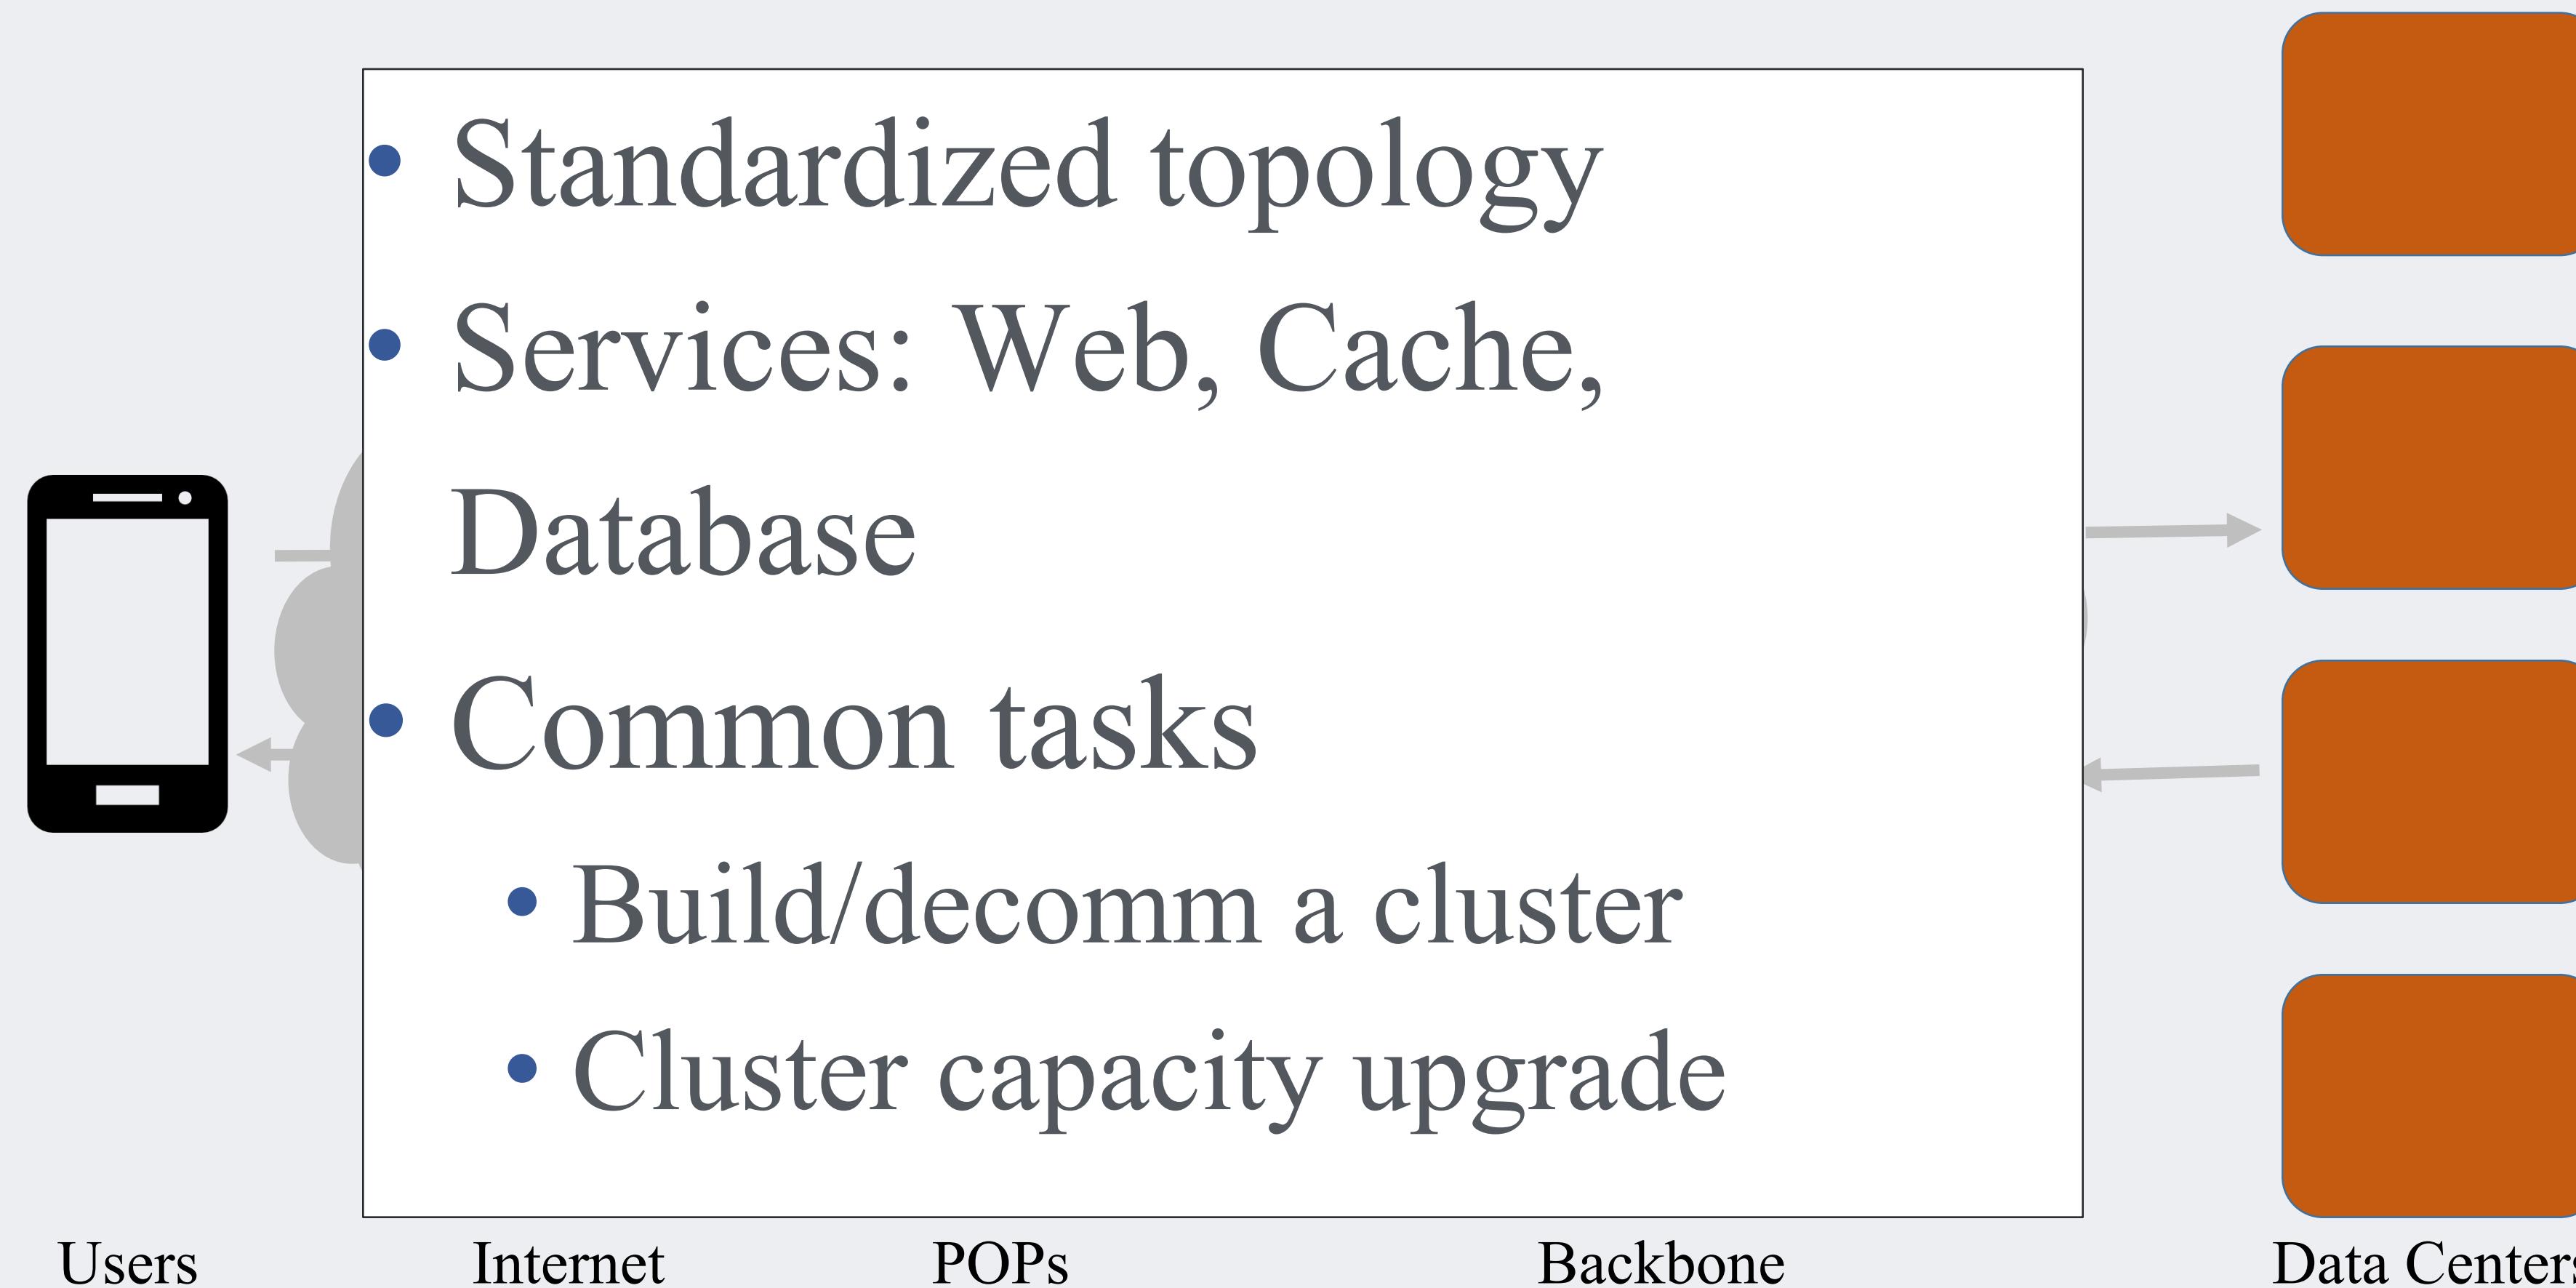
POPs

Backbone

Data Centers



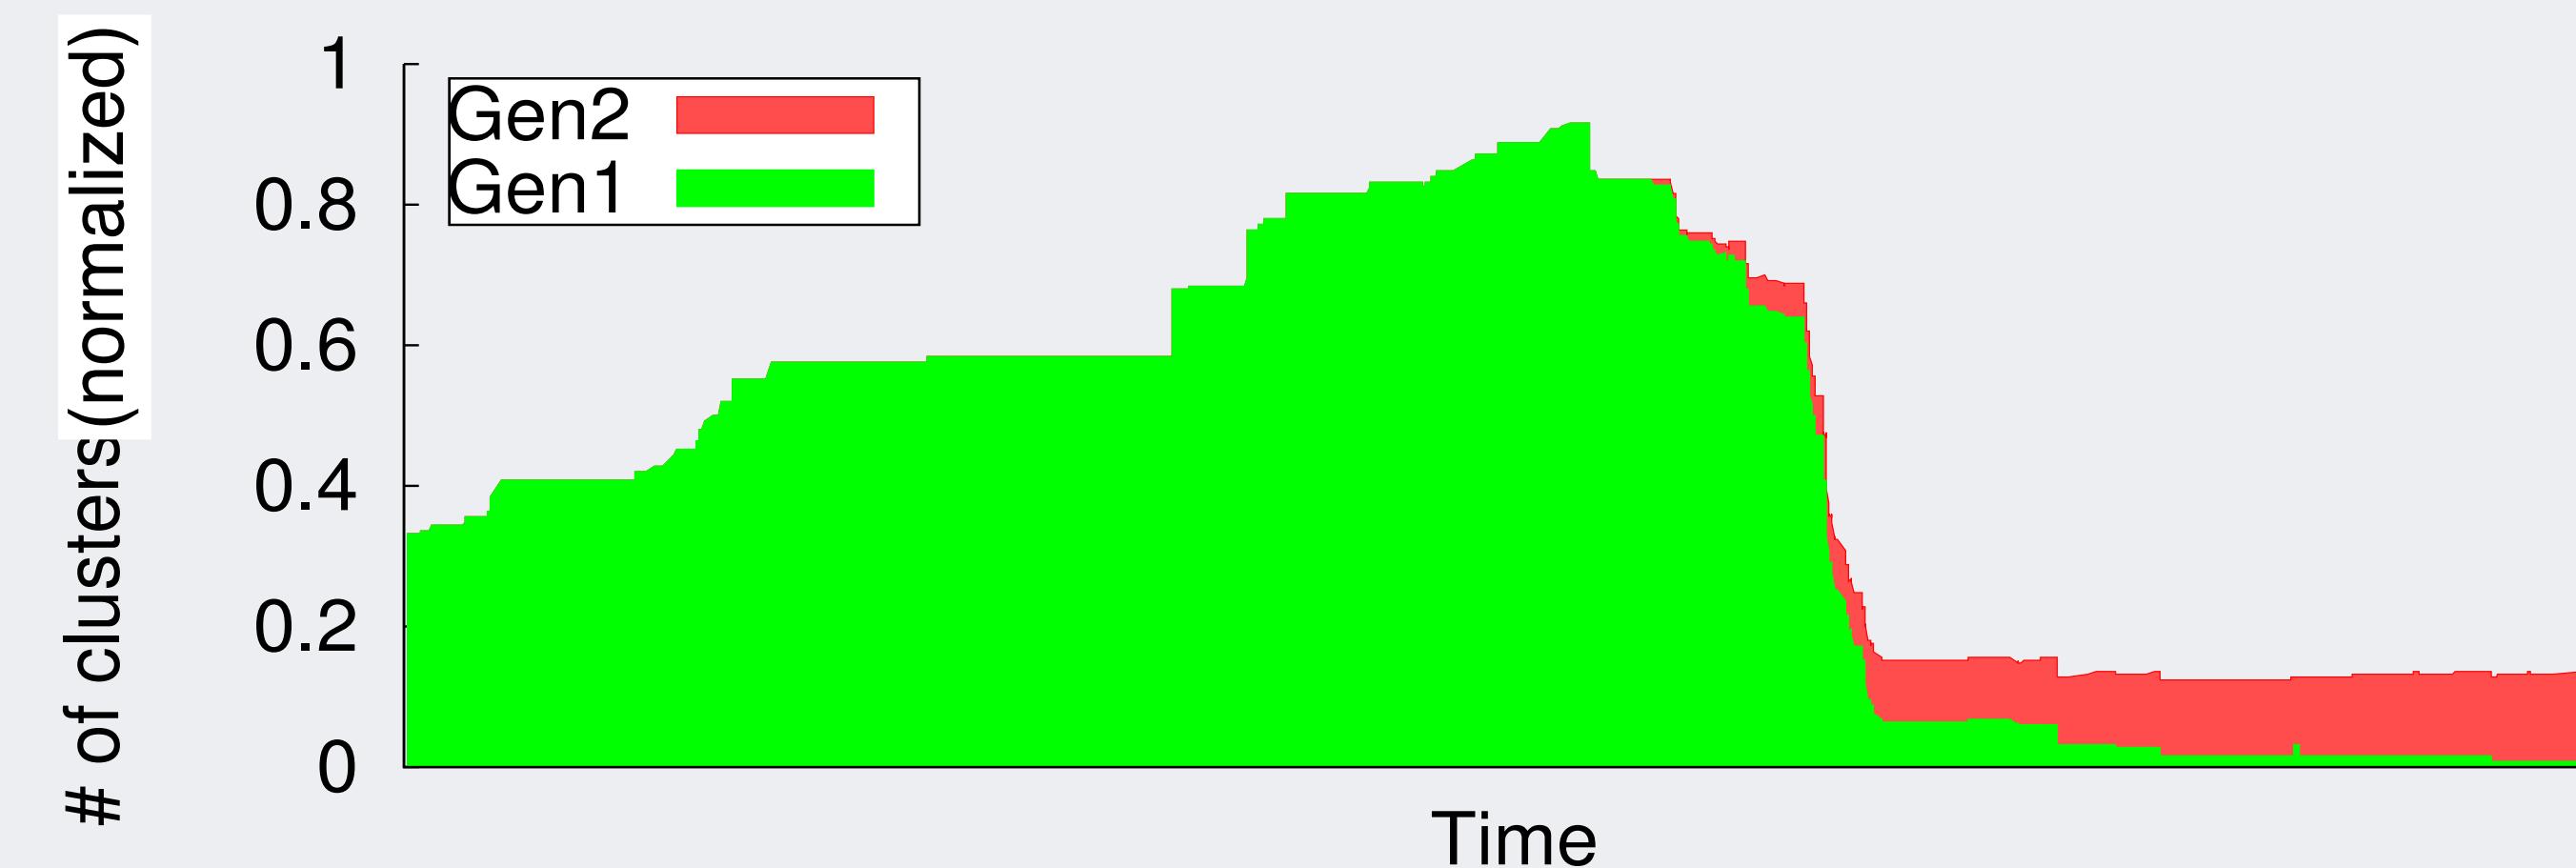
Datacenter



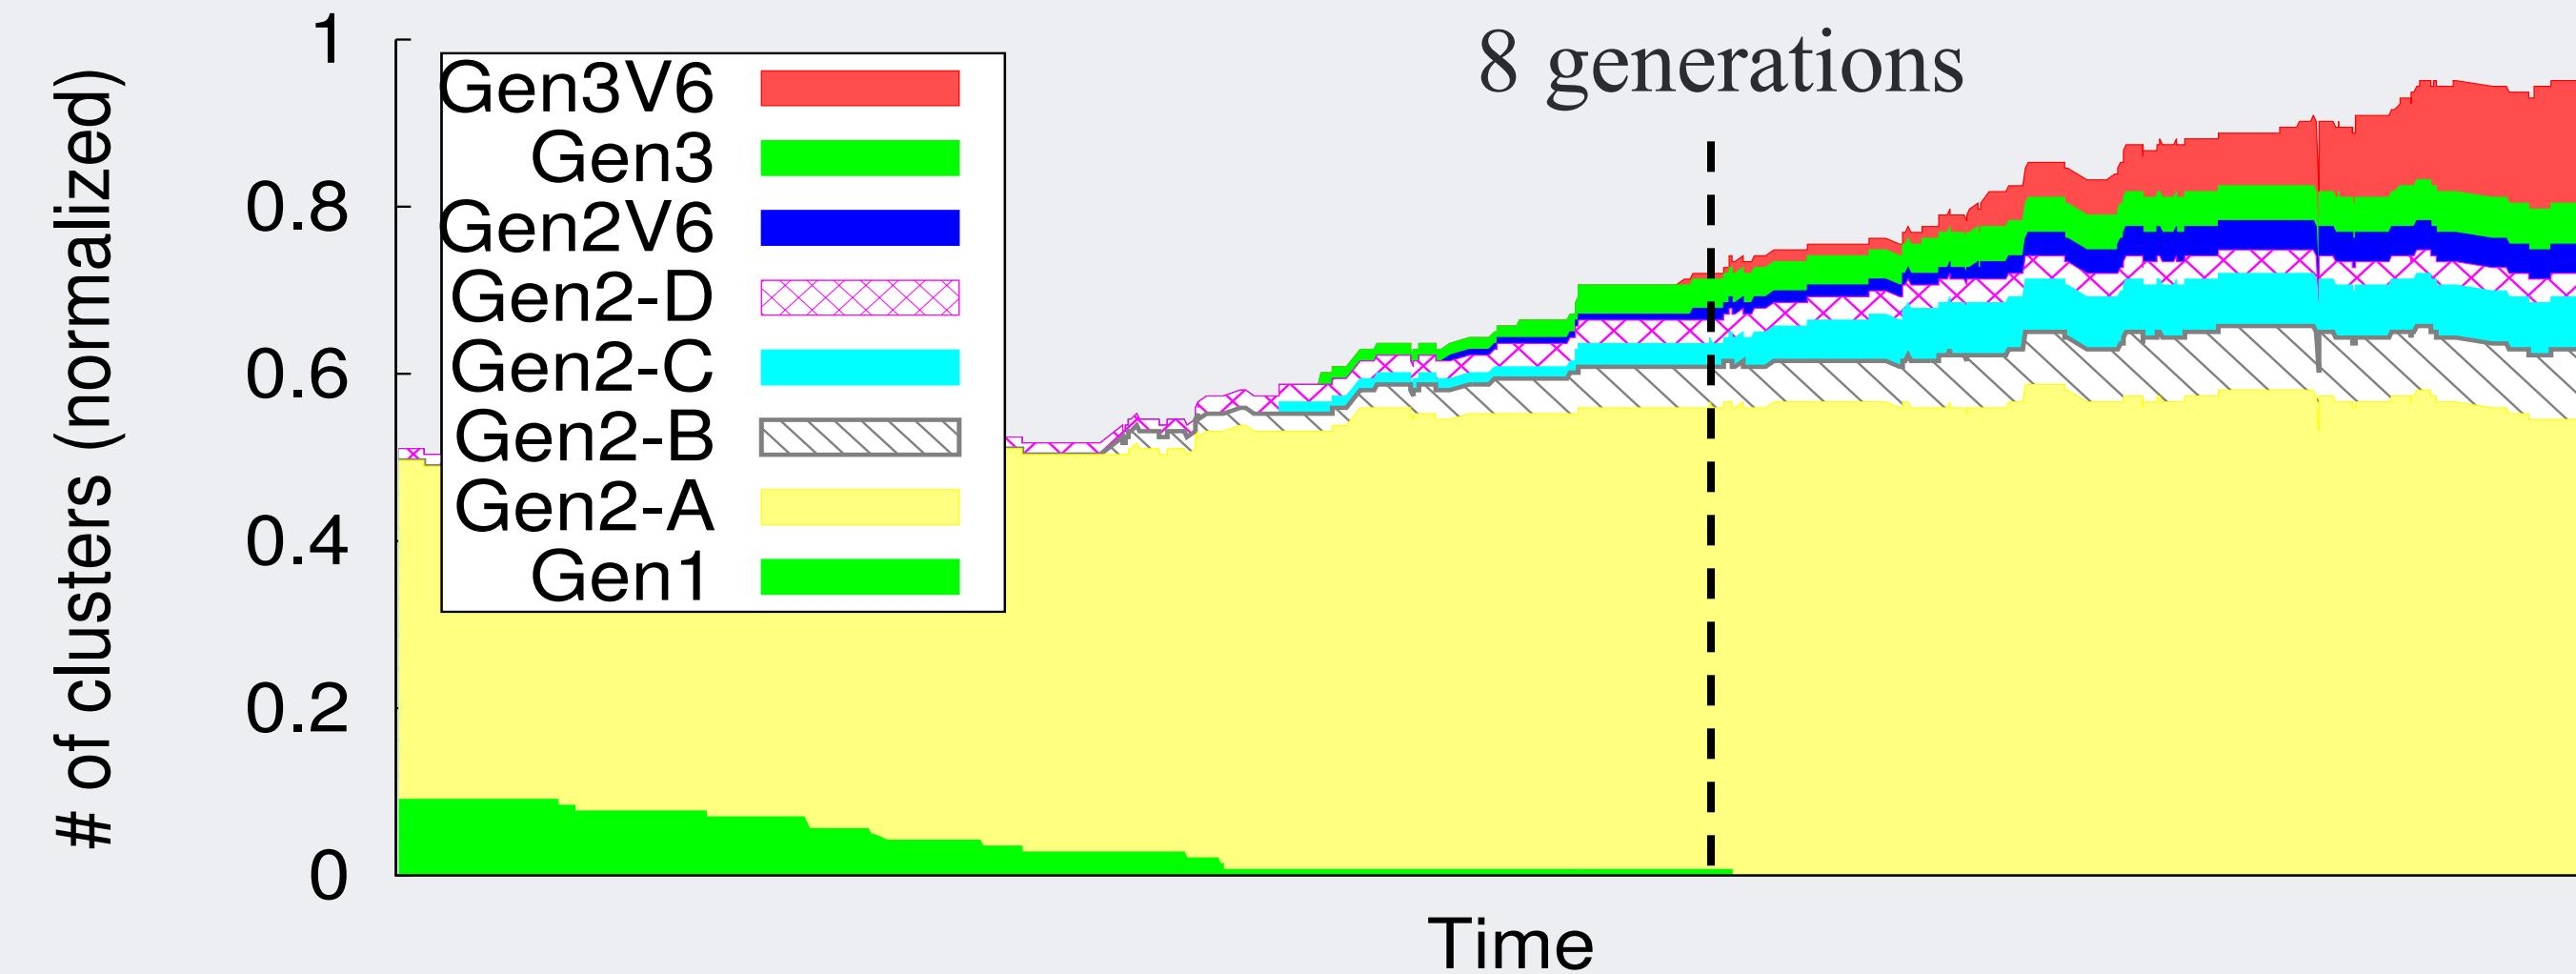
Overview of Facebook's Network

Multiple versions of FB cluster architectures co-exist

POP

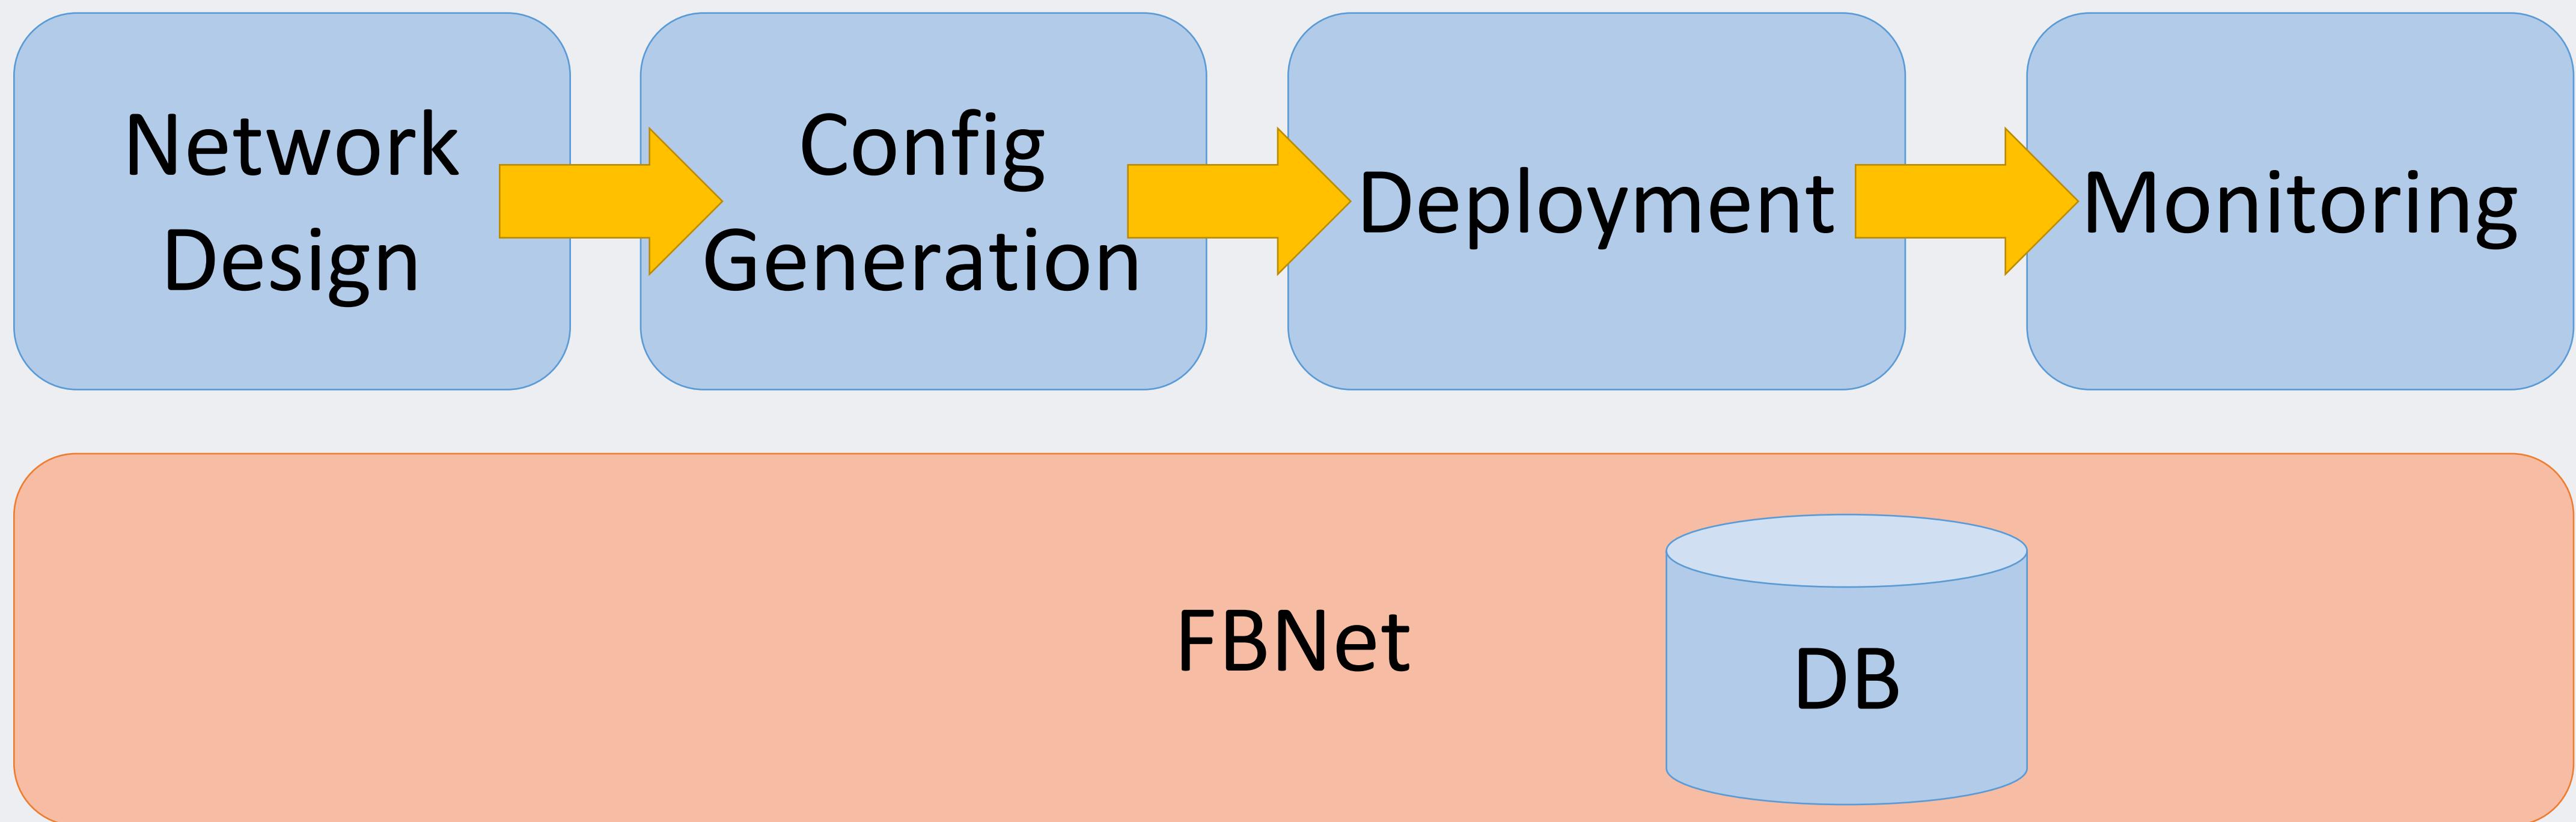


DC



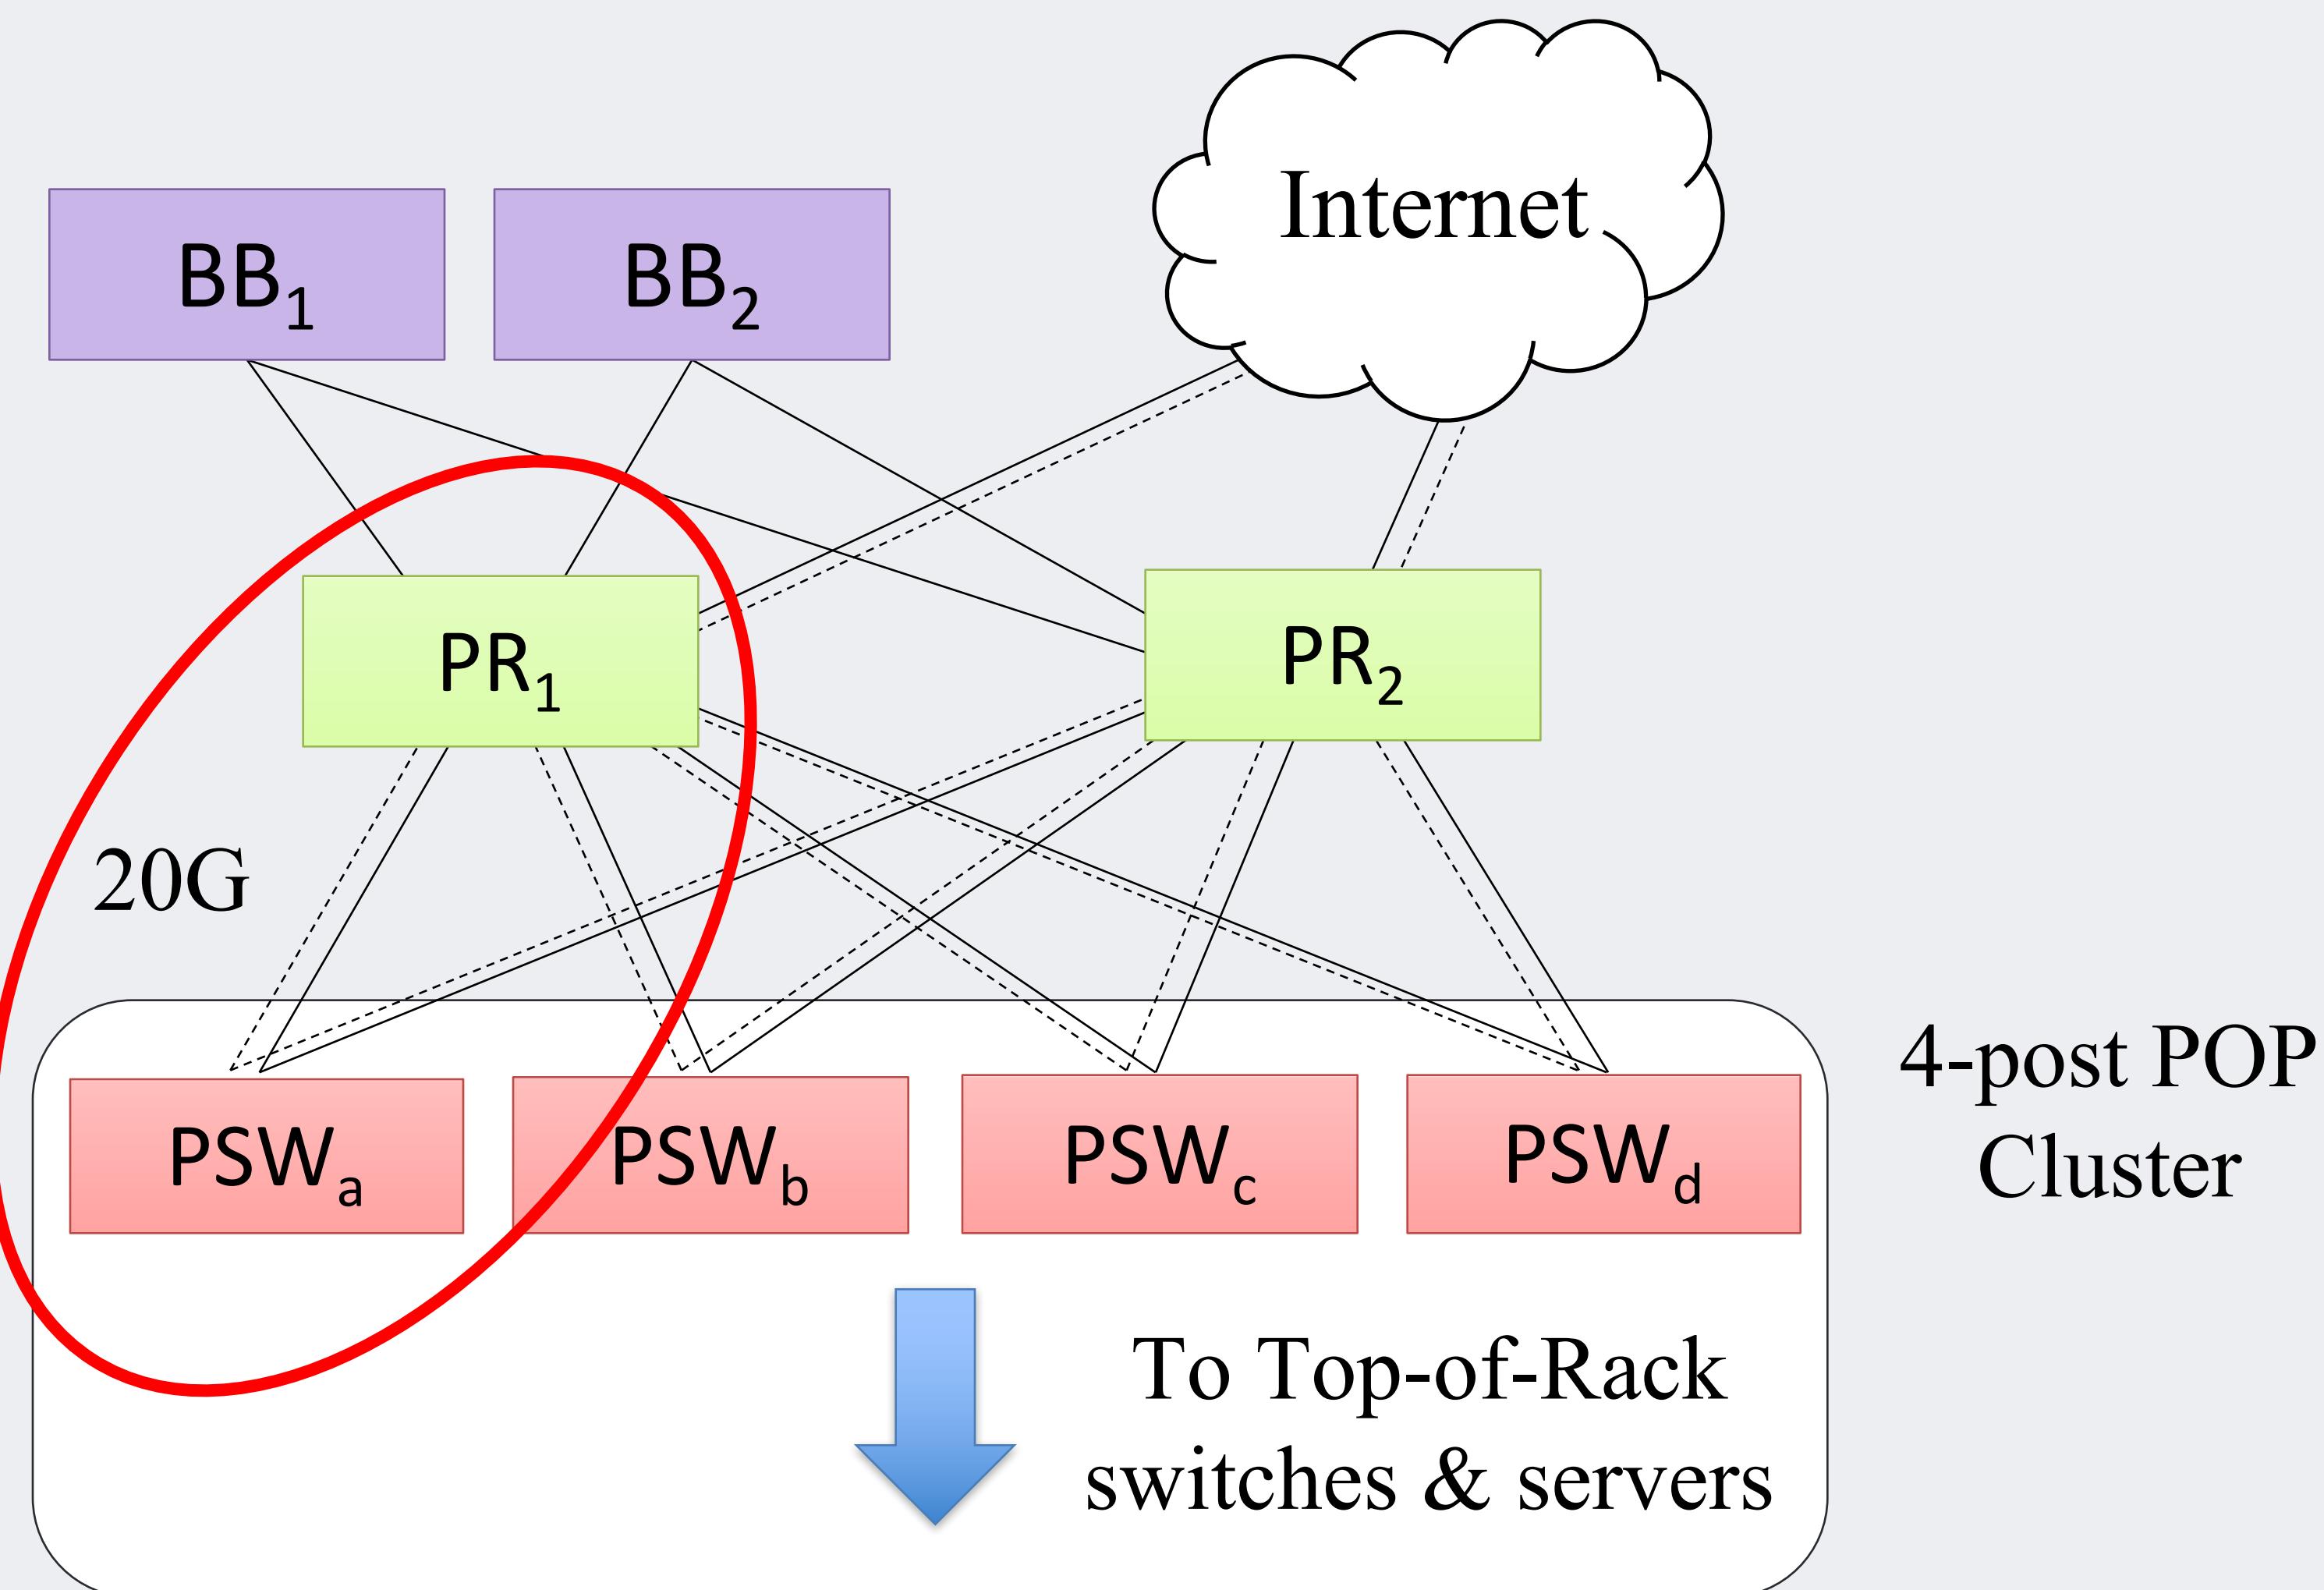
Robotron: “Top-Down” Network Management System@FB

Overview



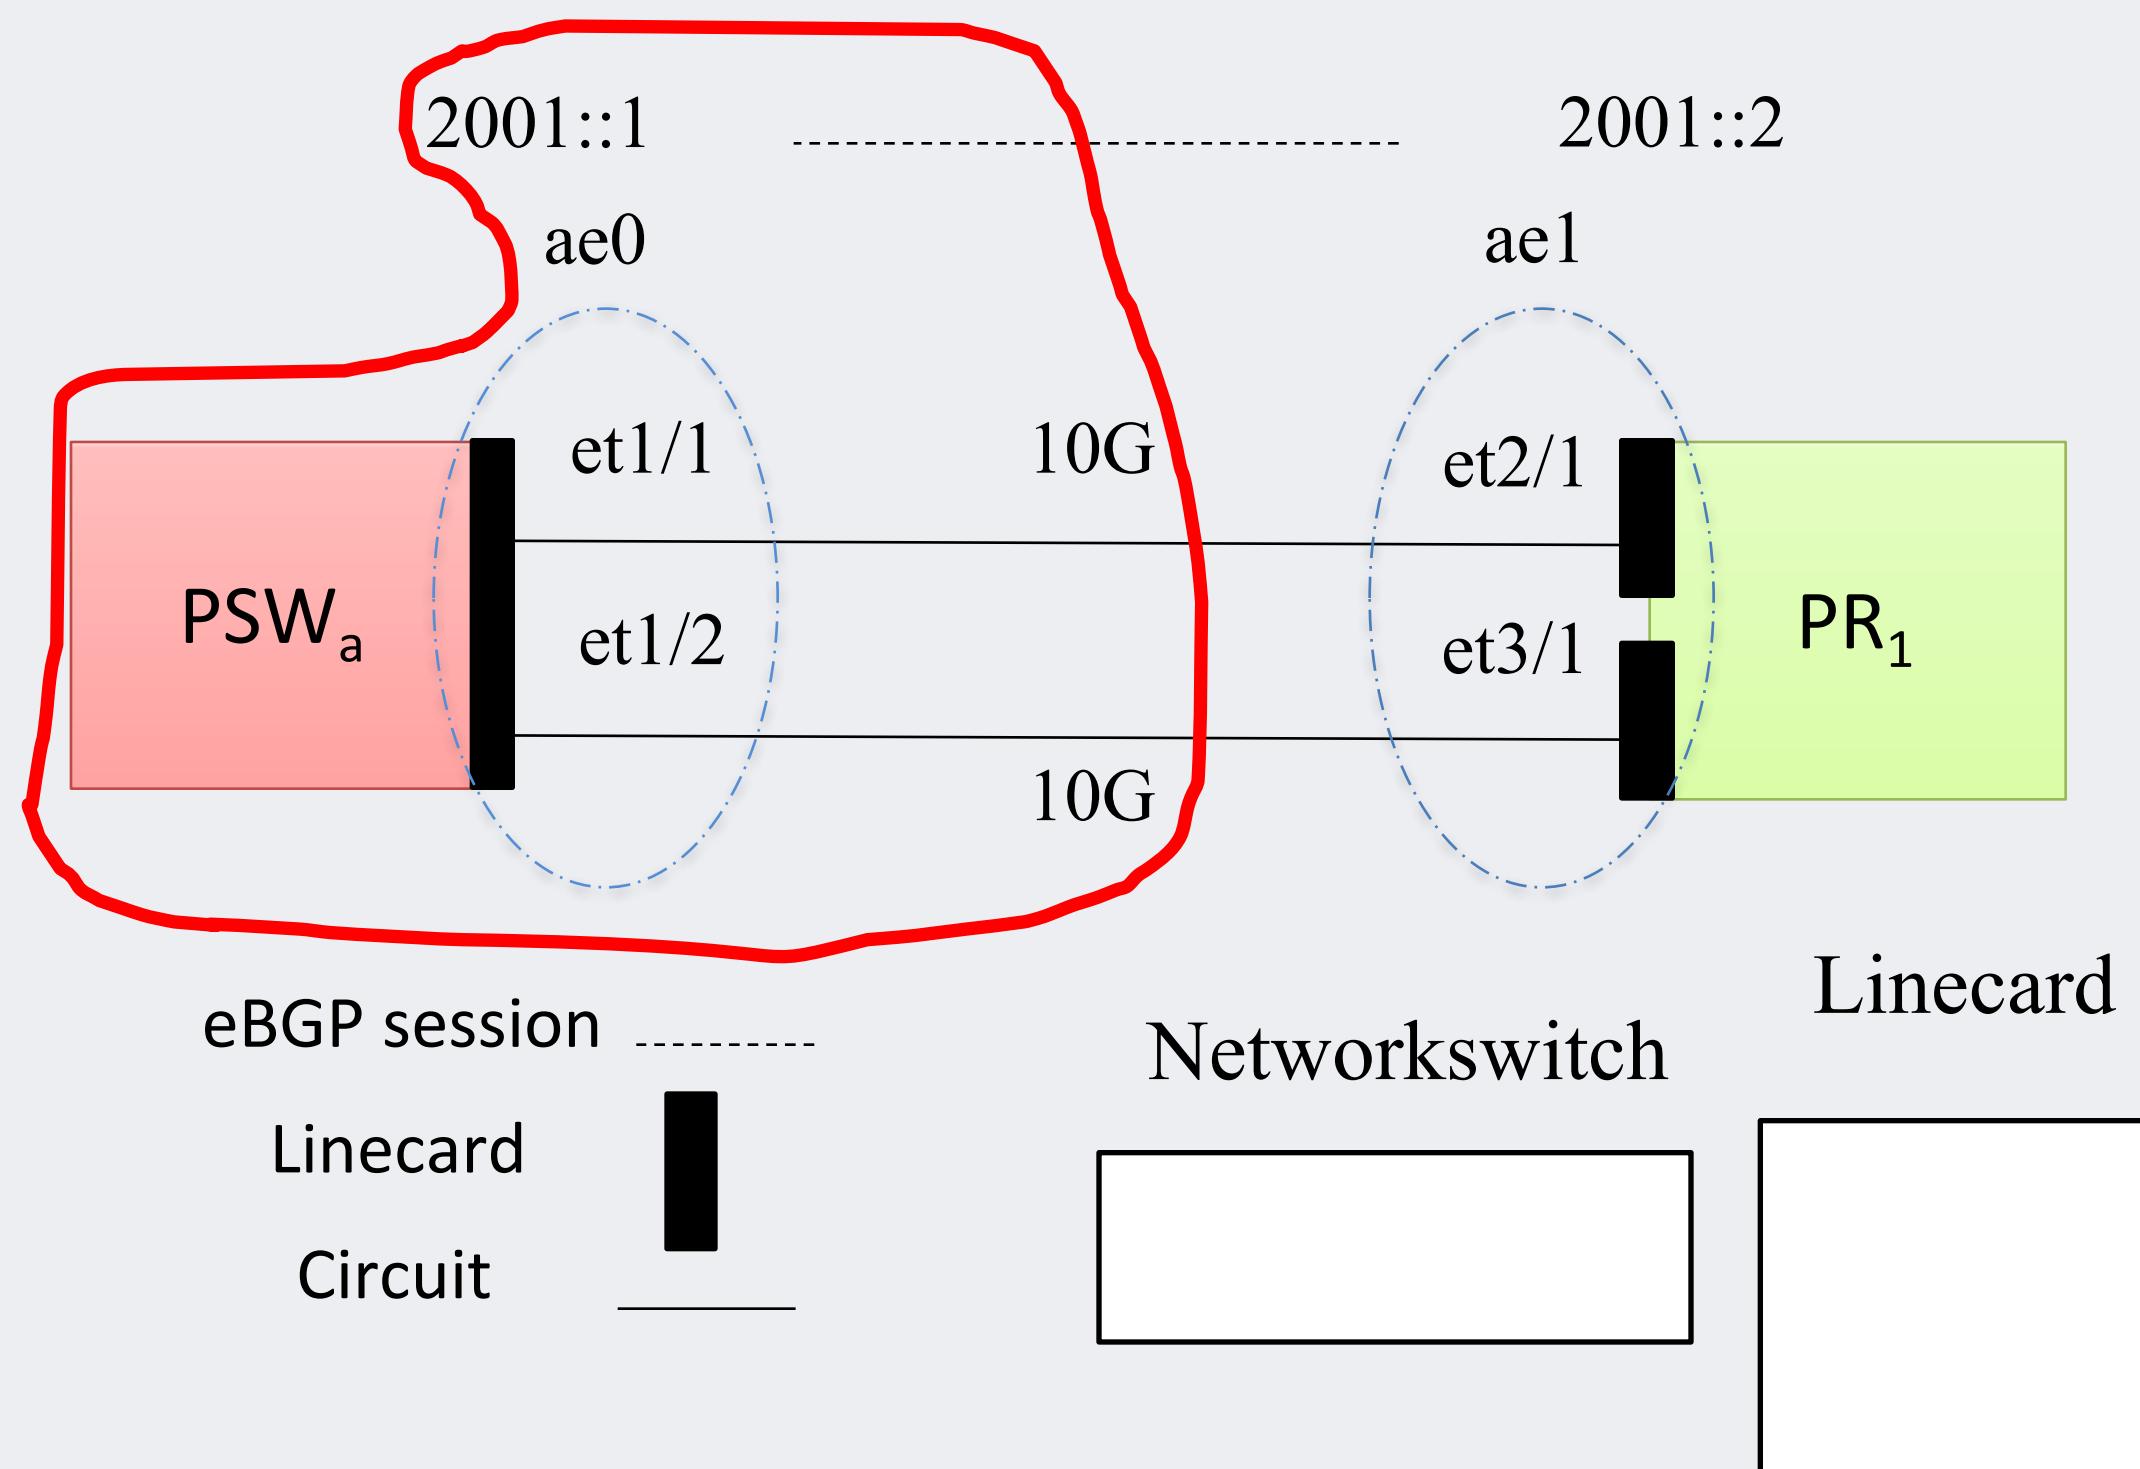
FBNet: Modeling the Network

Example 4-post POP cluster



FBNet: Modeling the Network

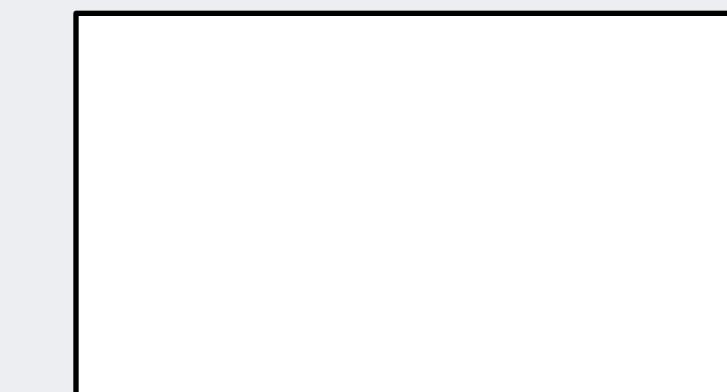
Object



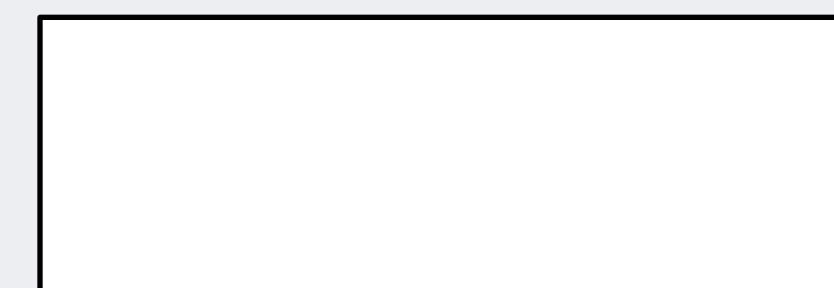
PhysicalInterface



PhysicalInterface



Circuit



AggregatedInterface



BgpV6Session

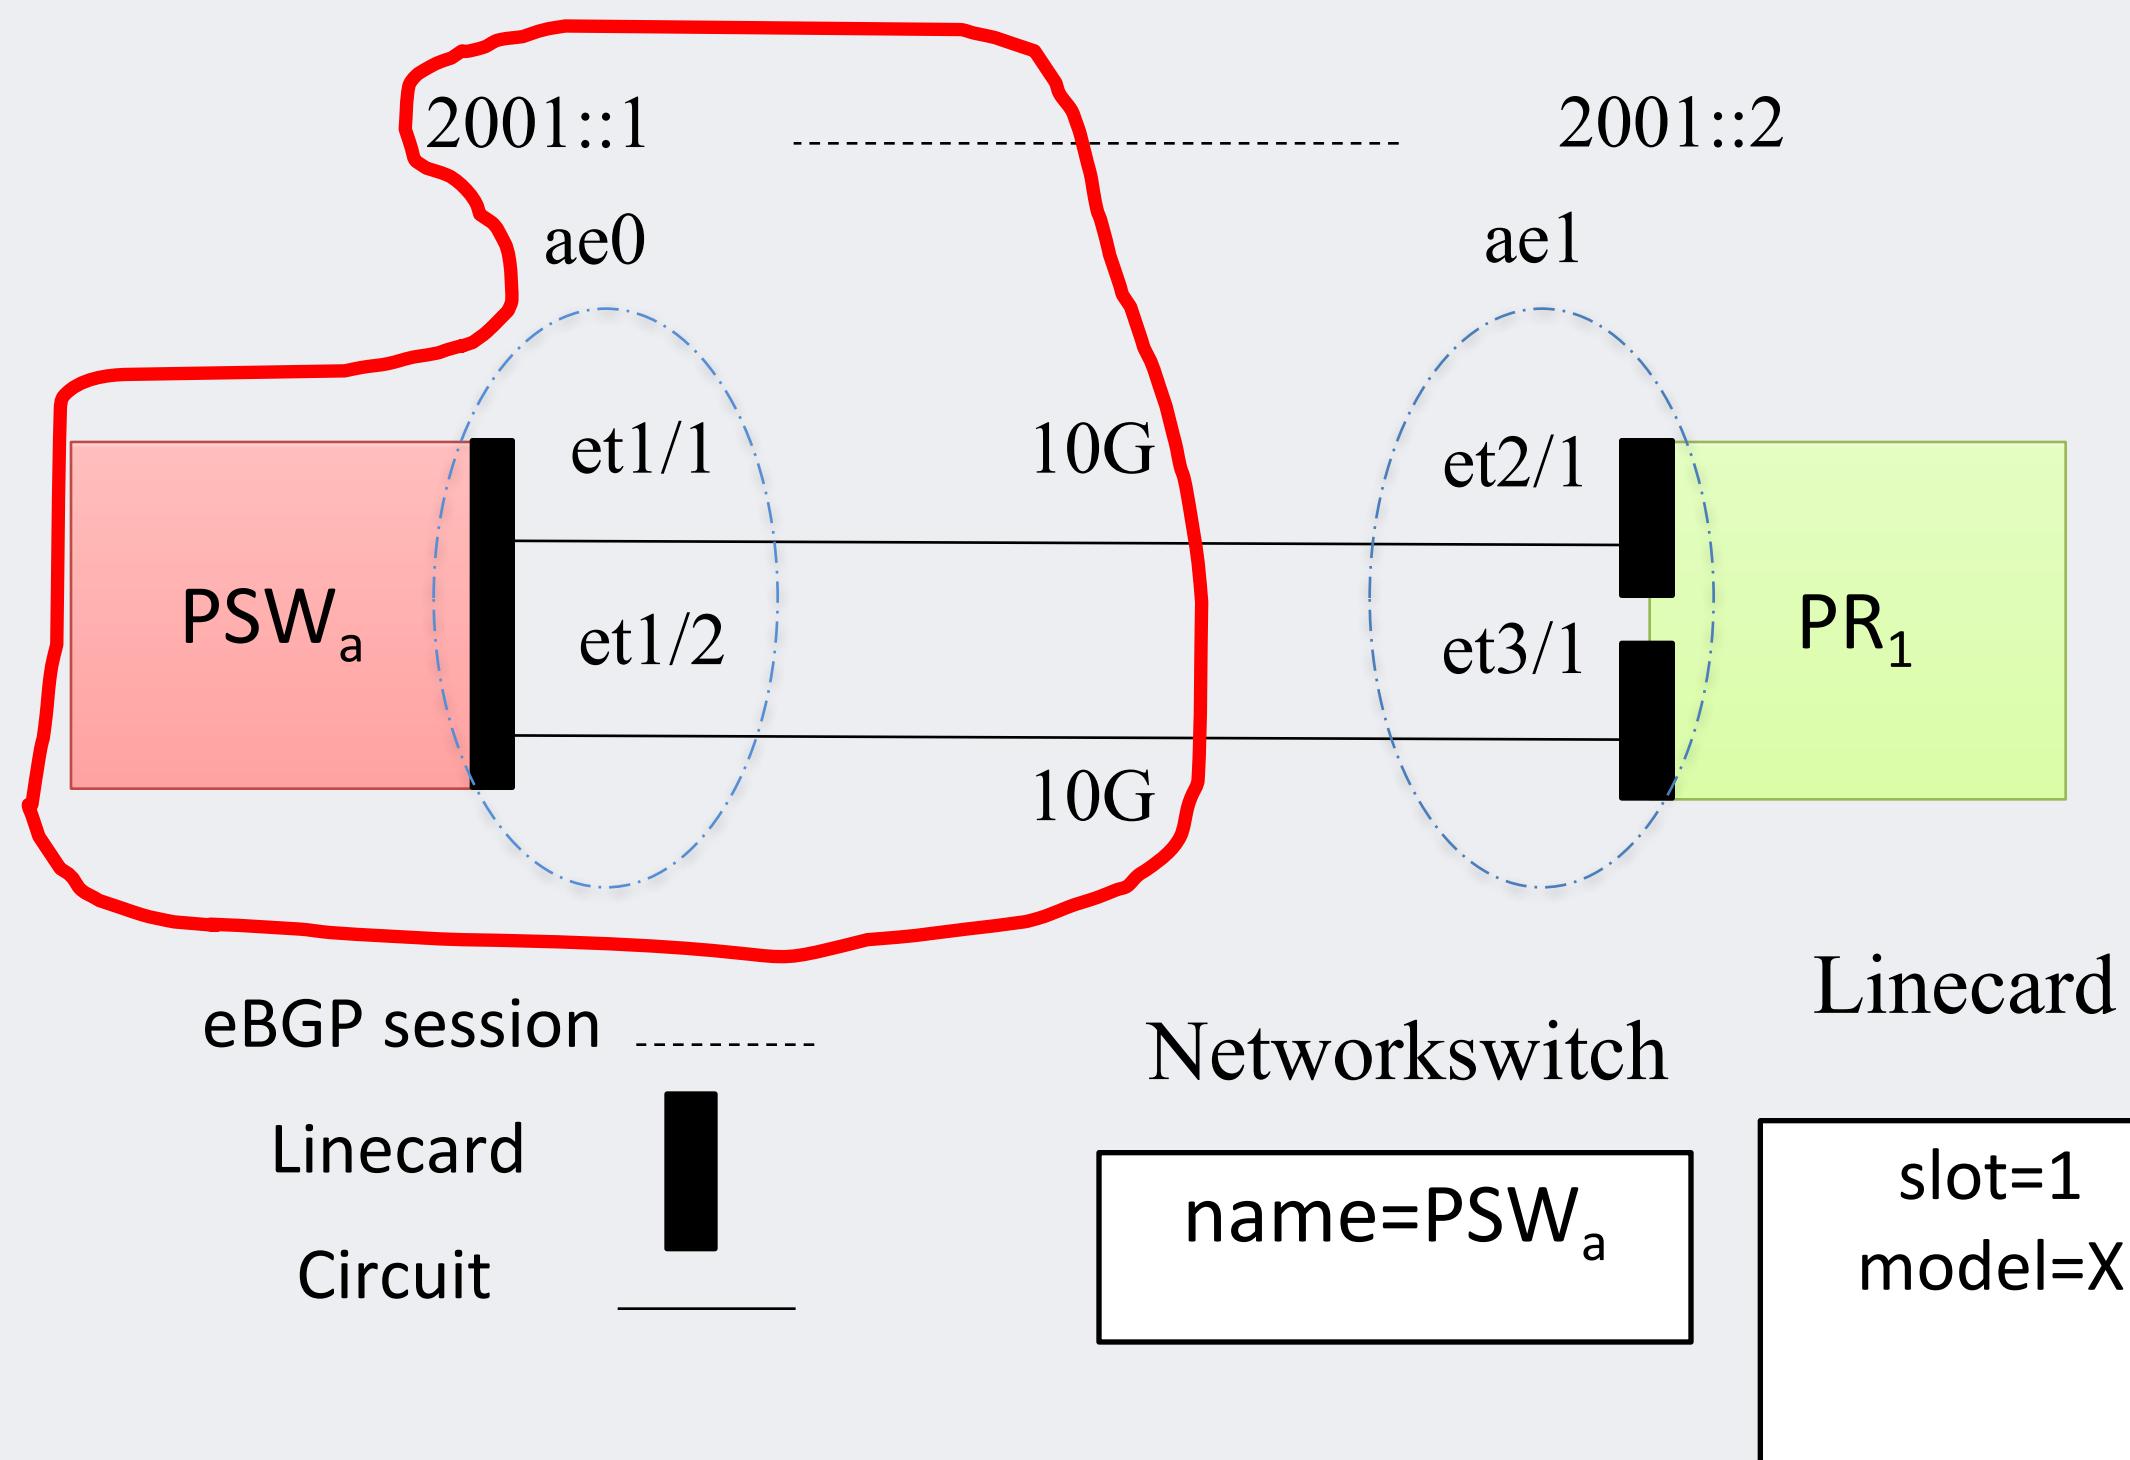


Circuit

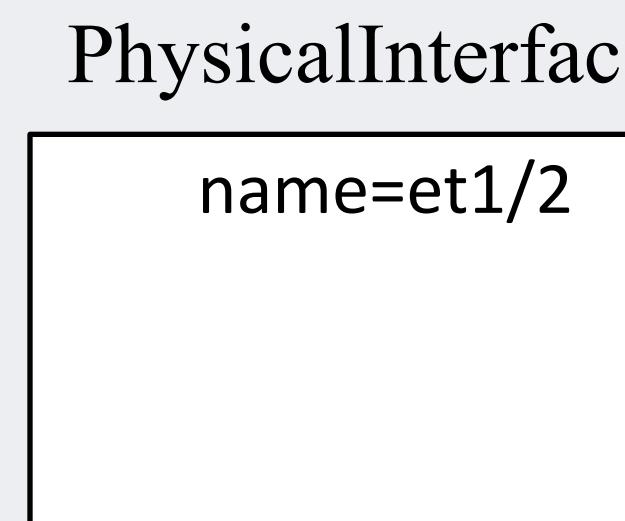
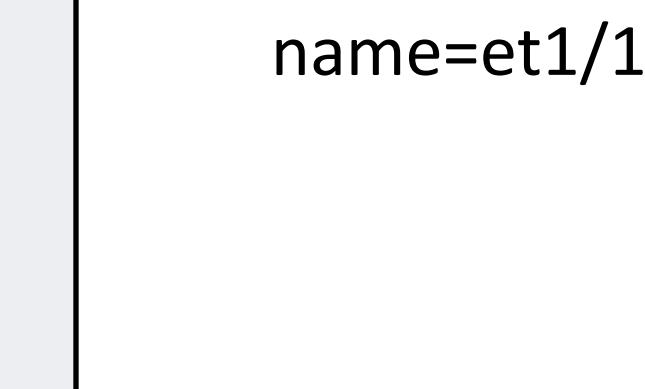


FBNet: Modeling the Network

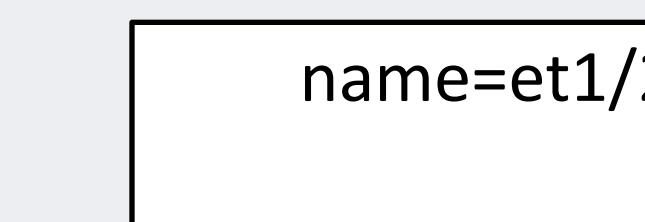
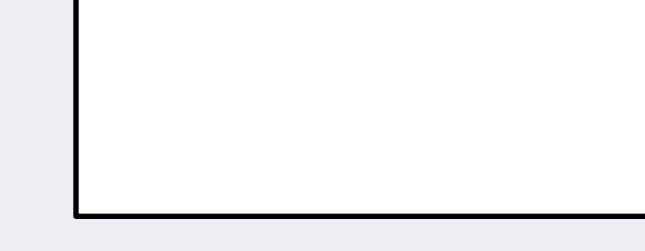
Value



PhysicalInterface

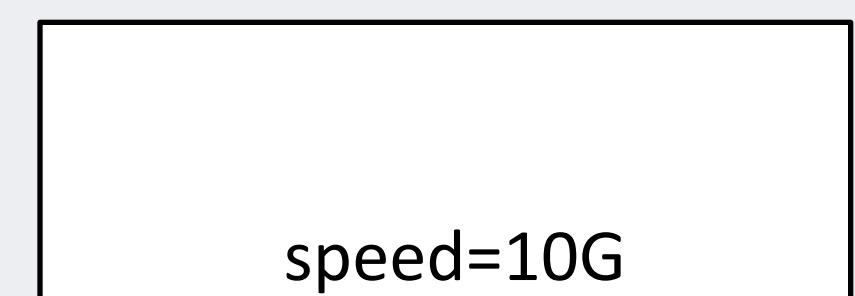


PhysicalInterface



V6Prefix

Circuit



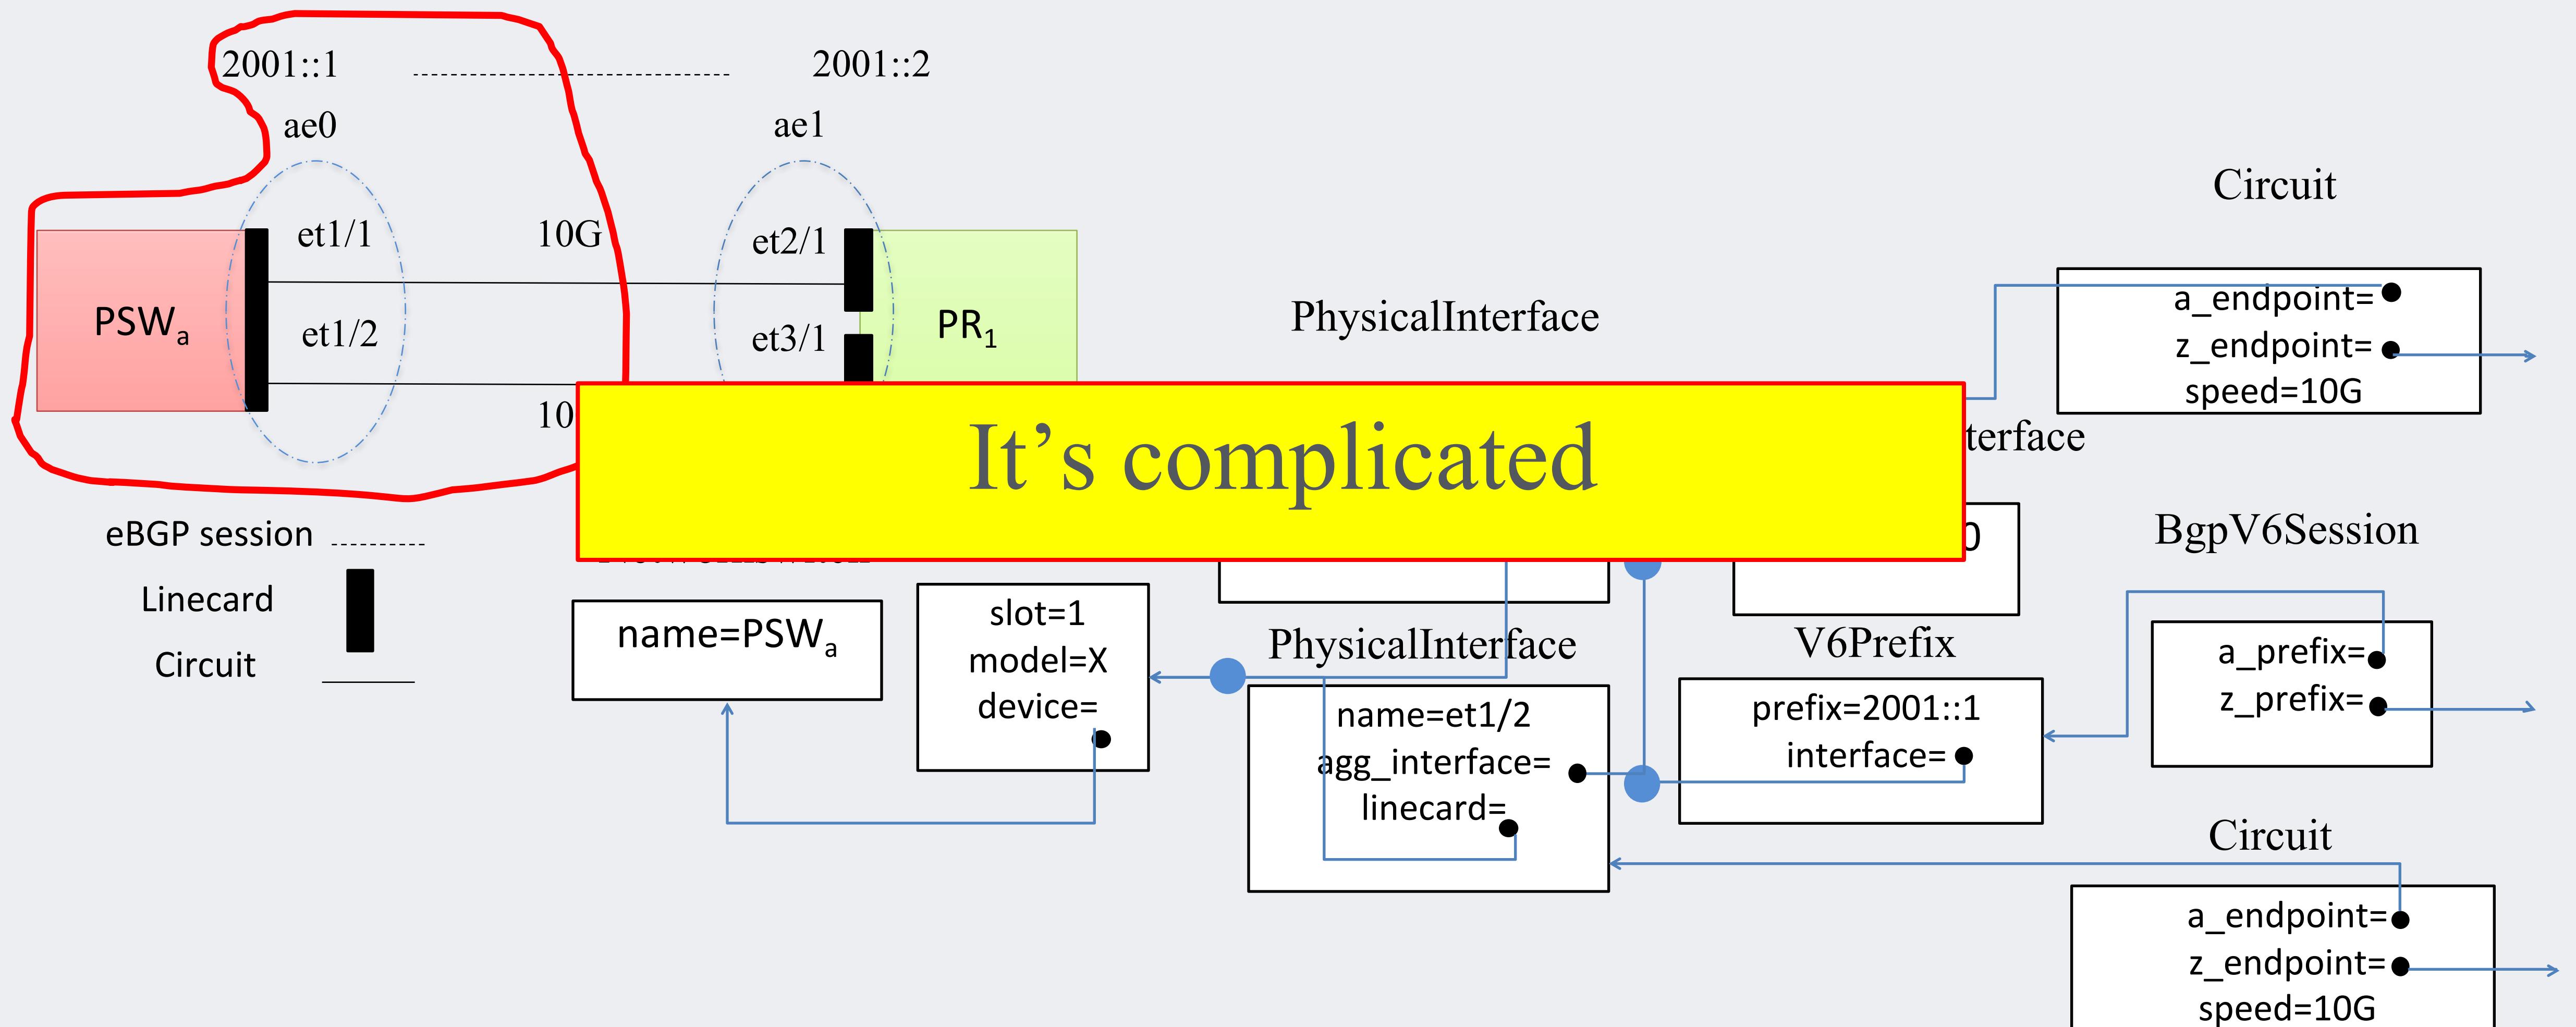
BgpV6Session



speed=10G

FBNet: Modeling the Network

Relationship



FBNet Model Snippet

```
class PhysicalInterface(Interface):
    linecard = models.ForeignKey(Linecard)
    agg_interface = models.ForeignKey(
        AggregatedInterface)
```

FBNet Model Snippet

Related models

```
class PhysicalInterface(Interface):
    linecard = models.ForeignKey(Linocard)
    agg_interface = models.ForeignKey(
        AggregatedInterface)
```

FBNet Model Snippet

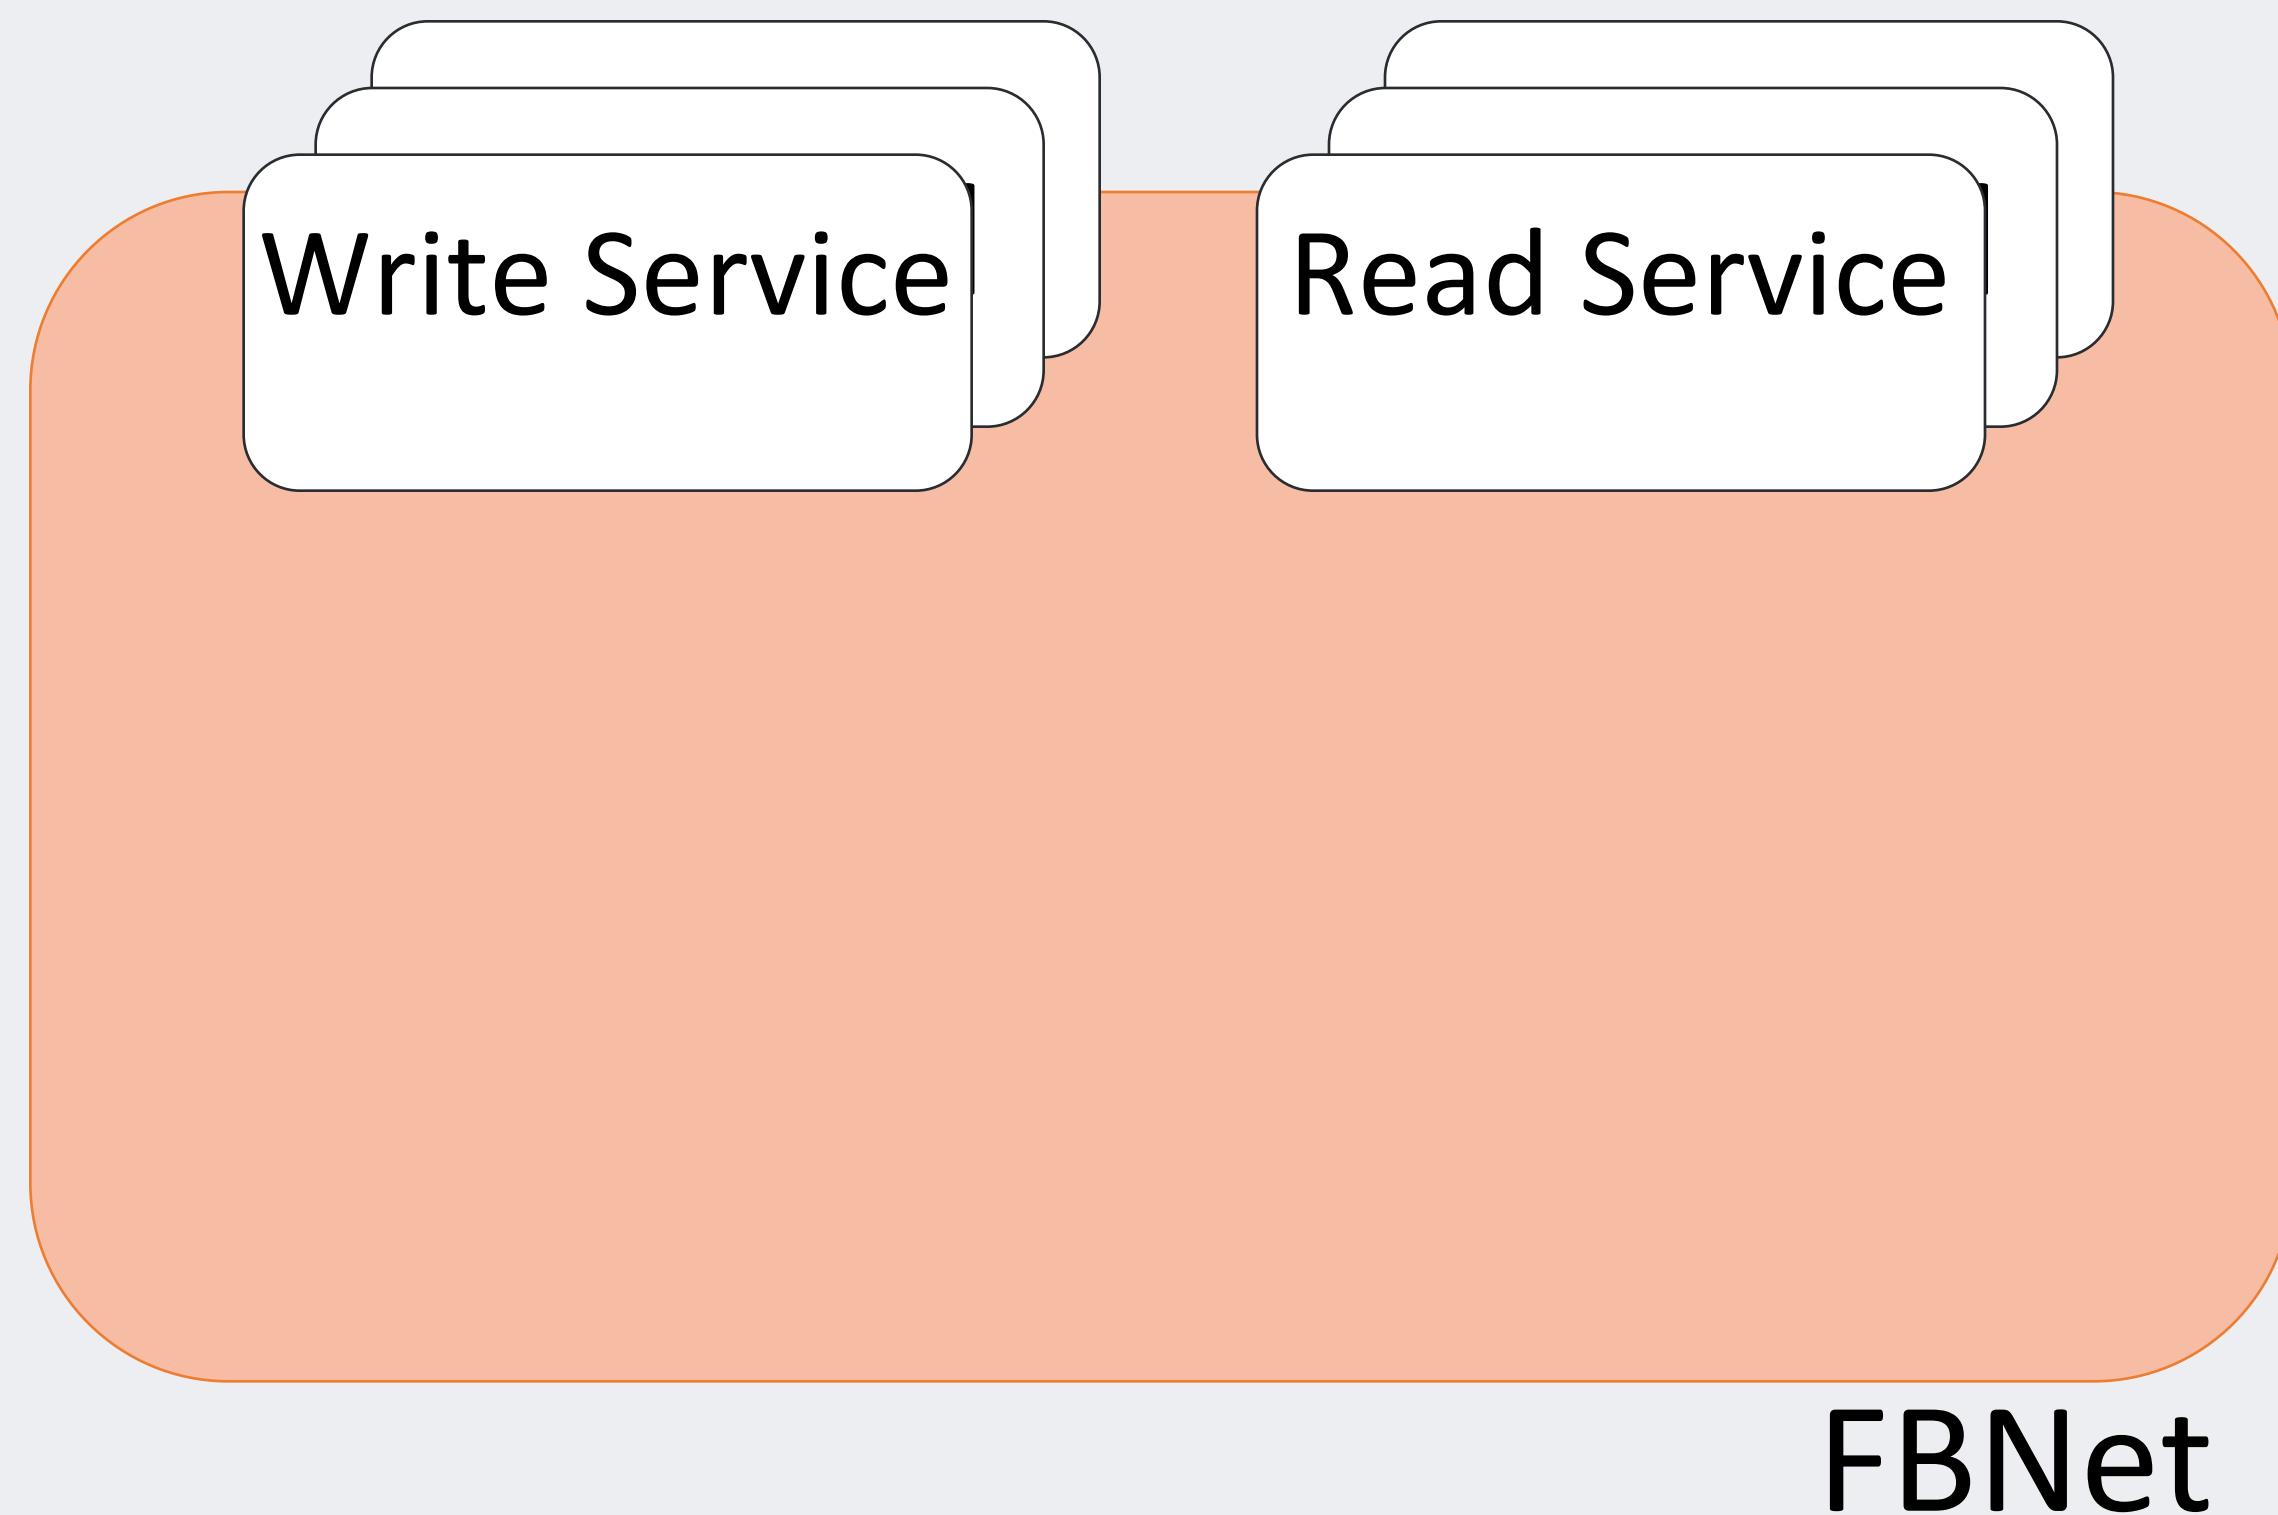
Model inheritance

```
class PhysicalInterface(Interface):
    linecard = models.ForeignKey(Linecard)
    agg_interface = models.ForeignKey(
        AggregatedInterface)
```

FBNet: Architecture

API Layer

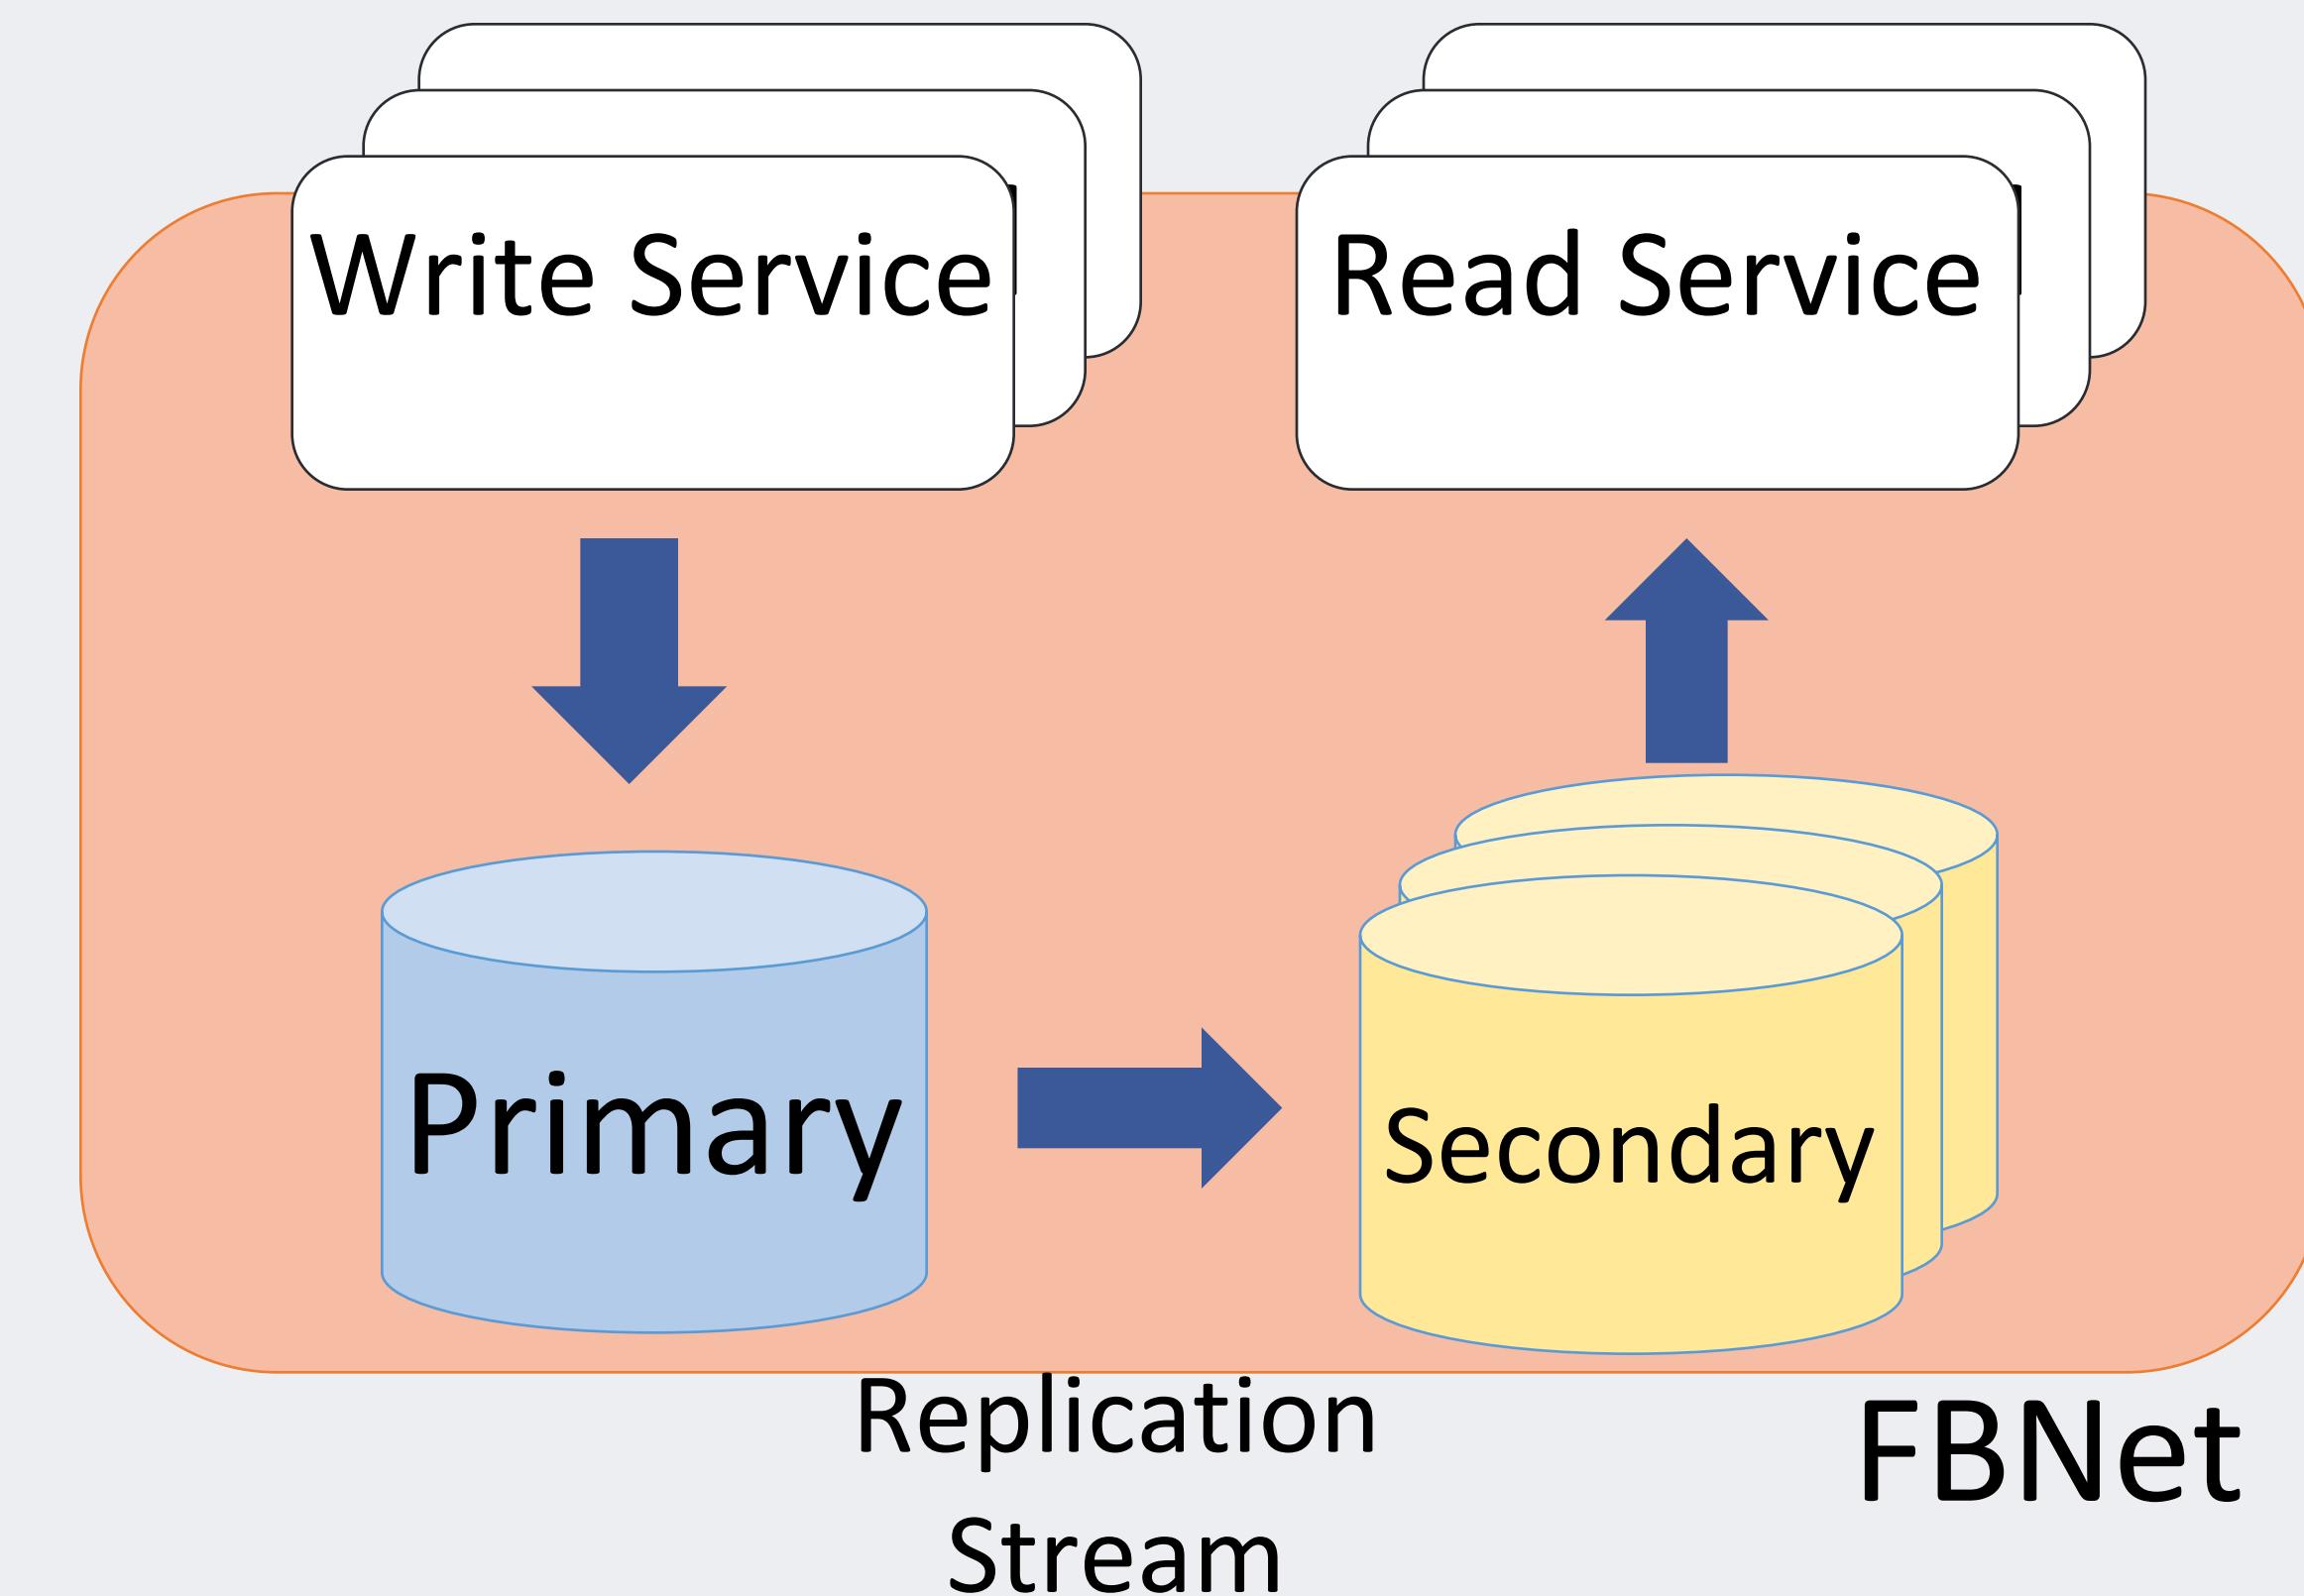
- RPC services
 - Read: fine-grained per-model query
 - Write: task-based
- High Availability: Multiple replicas per DC



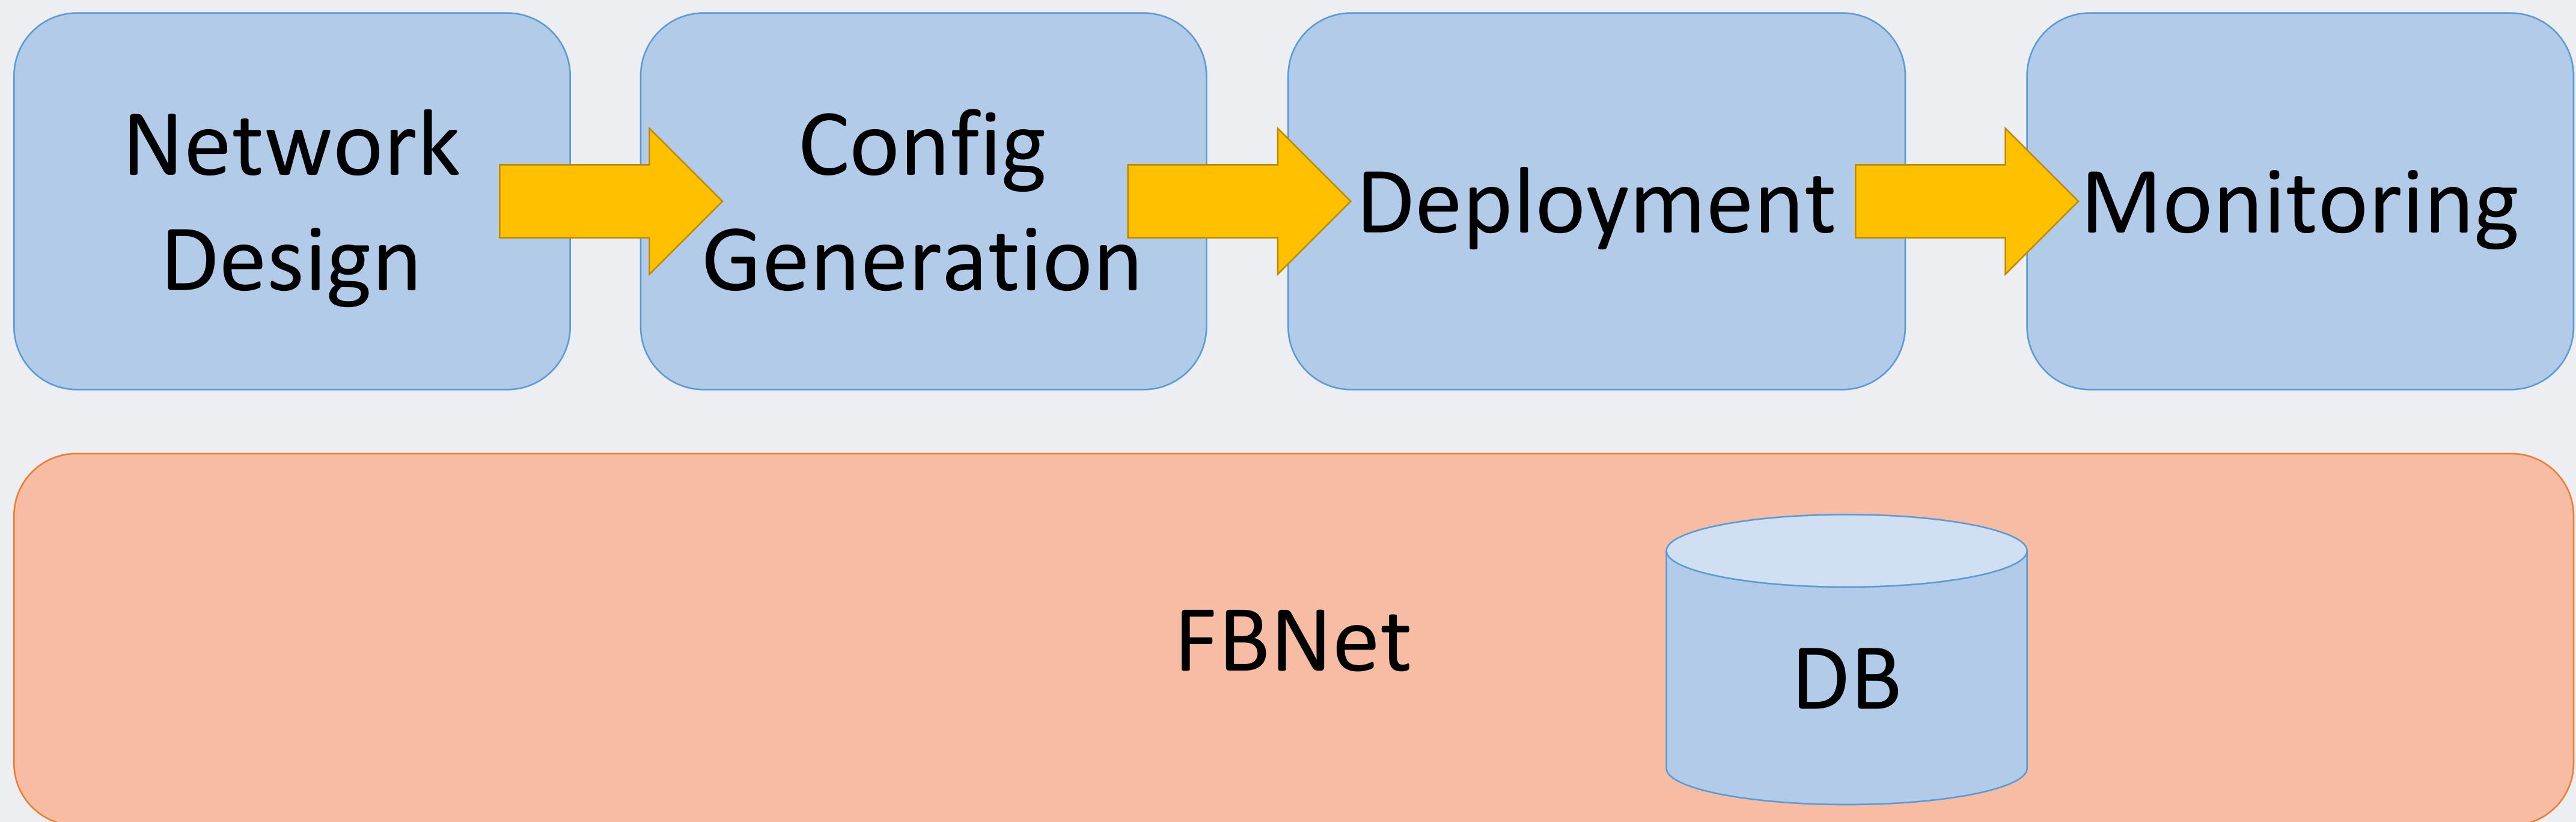
FBNet: Architecture

API Layer

- 1 primary, multiple secondary DBs
- Scalability: 1 slave per DC



Robotron's management life cycle



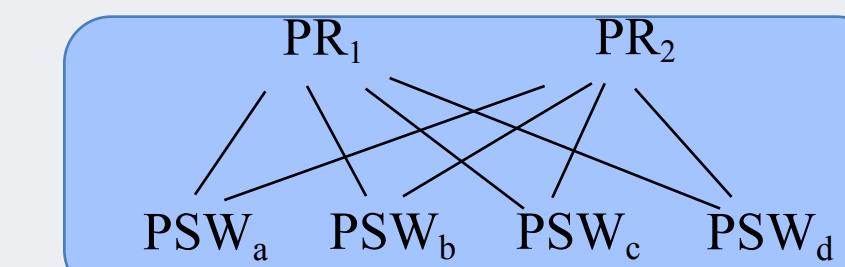
Network Design

Design intent → FBNet objects

Template for a POP cluster

```
Cluster(  
    devices={  
        PR: DeviceSpec(  
            hardware="Router_Vendor1"  
            num_devices=2)  
        PSW: DeviceSpec(  
            hardware="Switch_Vendor2"  
            num_devices=4)  
    },  
    Link_groups=[  
        LinkGroup(  
            a_device=PR,  
            z_device=PSW,  
            pifs_per_agg=2,  
            ip=V6)  
    ]  
)
```

FBNet objects



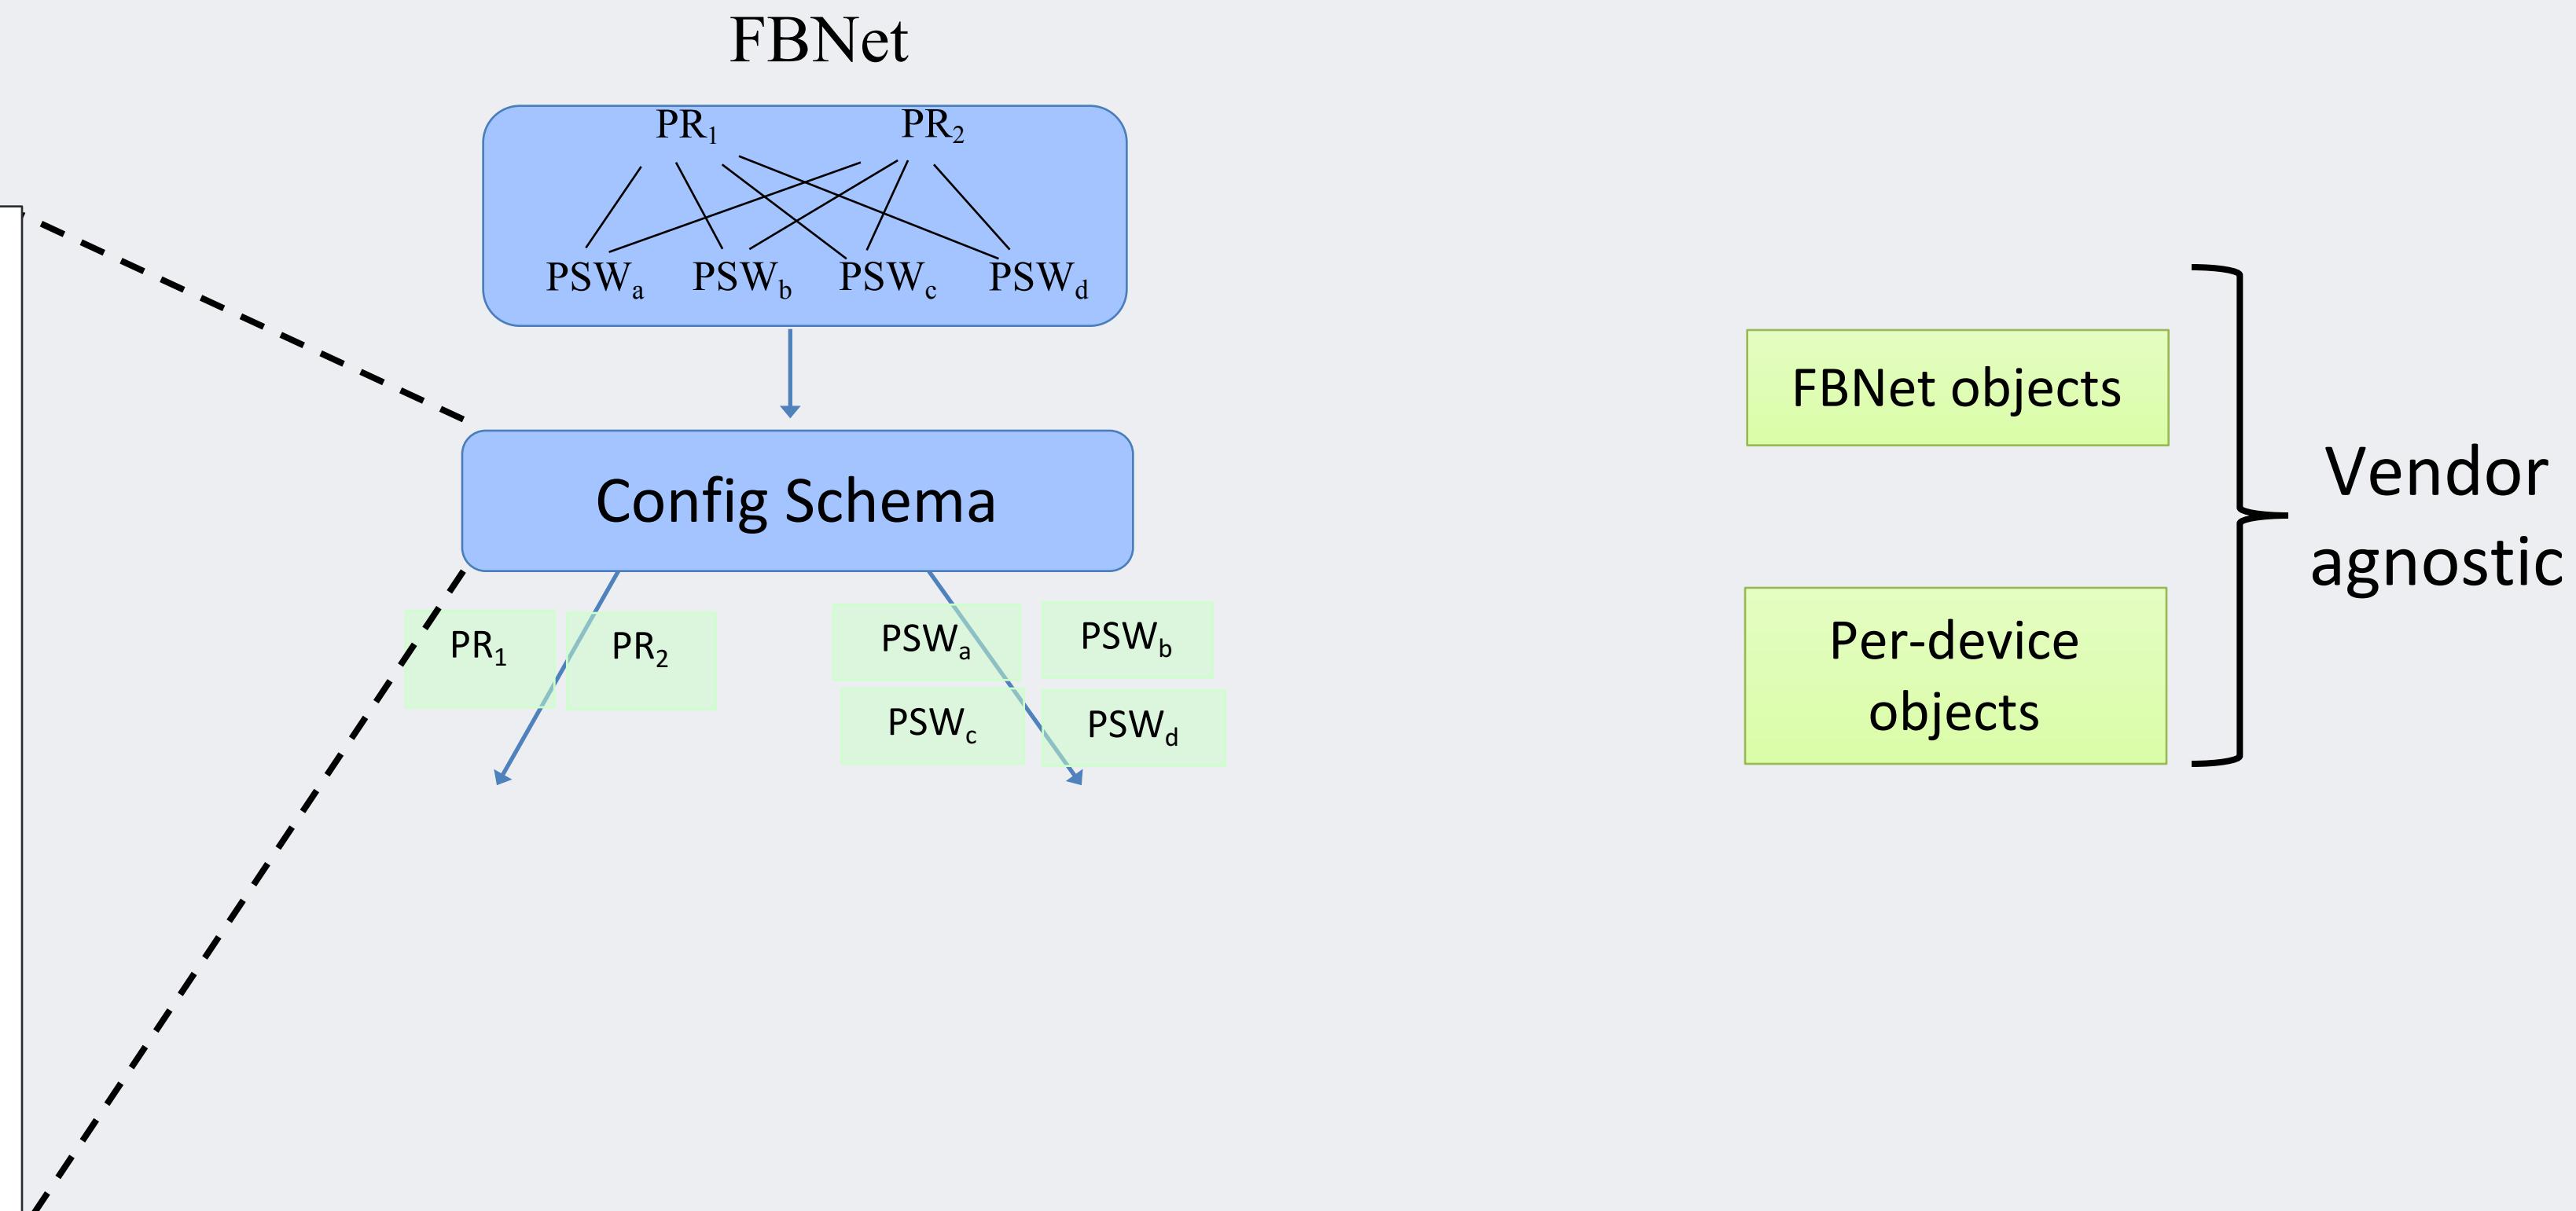
BackboneRouters: 2
NetworkSwitches: 4
Circuits: 16
PhysicalInterfaces: 32
AggregatedInterfaces: 16
V6Prefixes: 16
BgpV6Sessions: 8

94 objects across 7 models

Config Generation

FBNet objects → Device configs

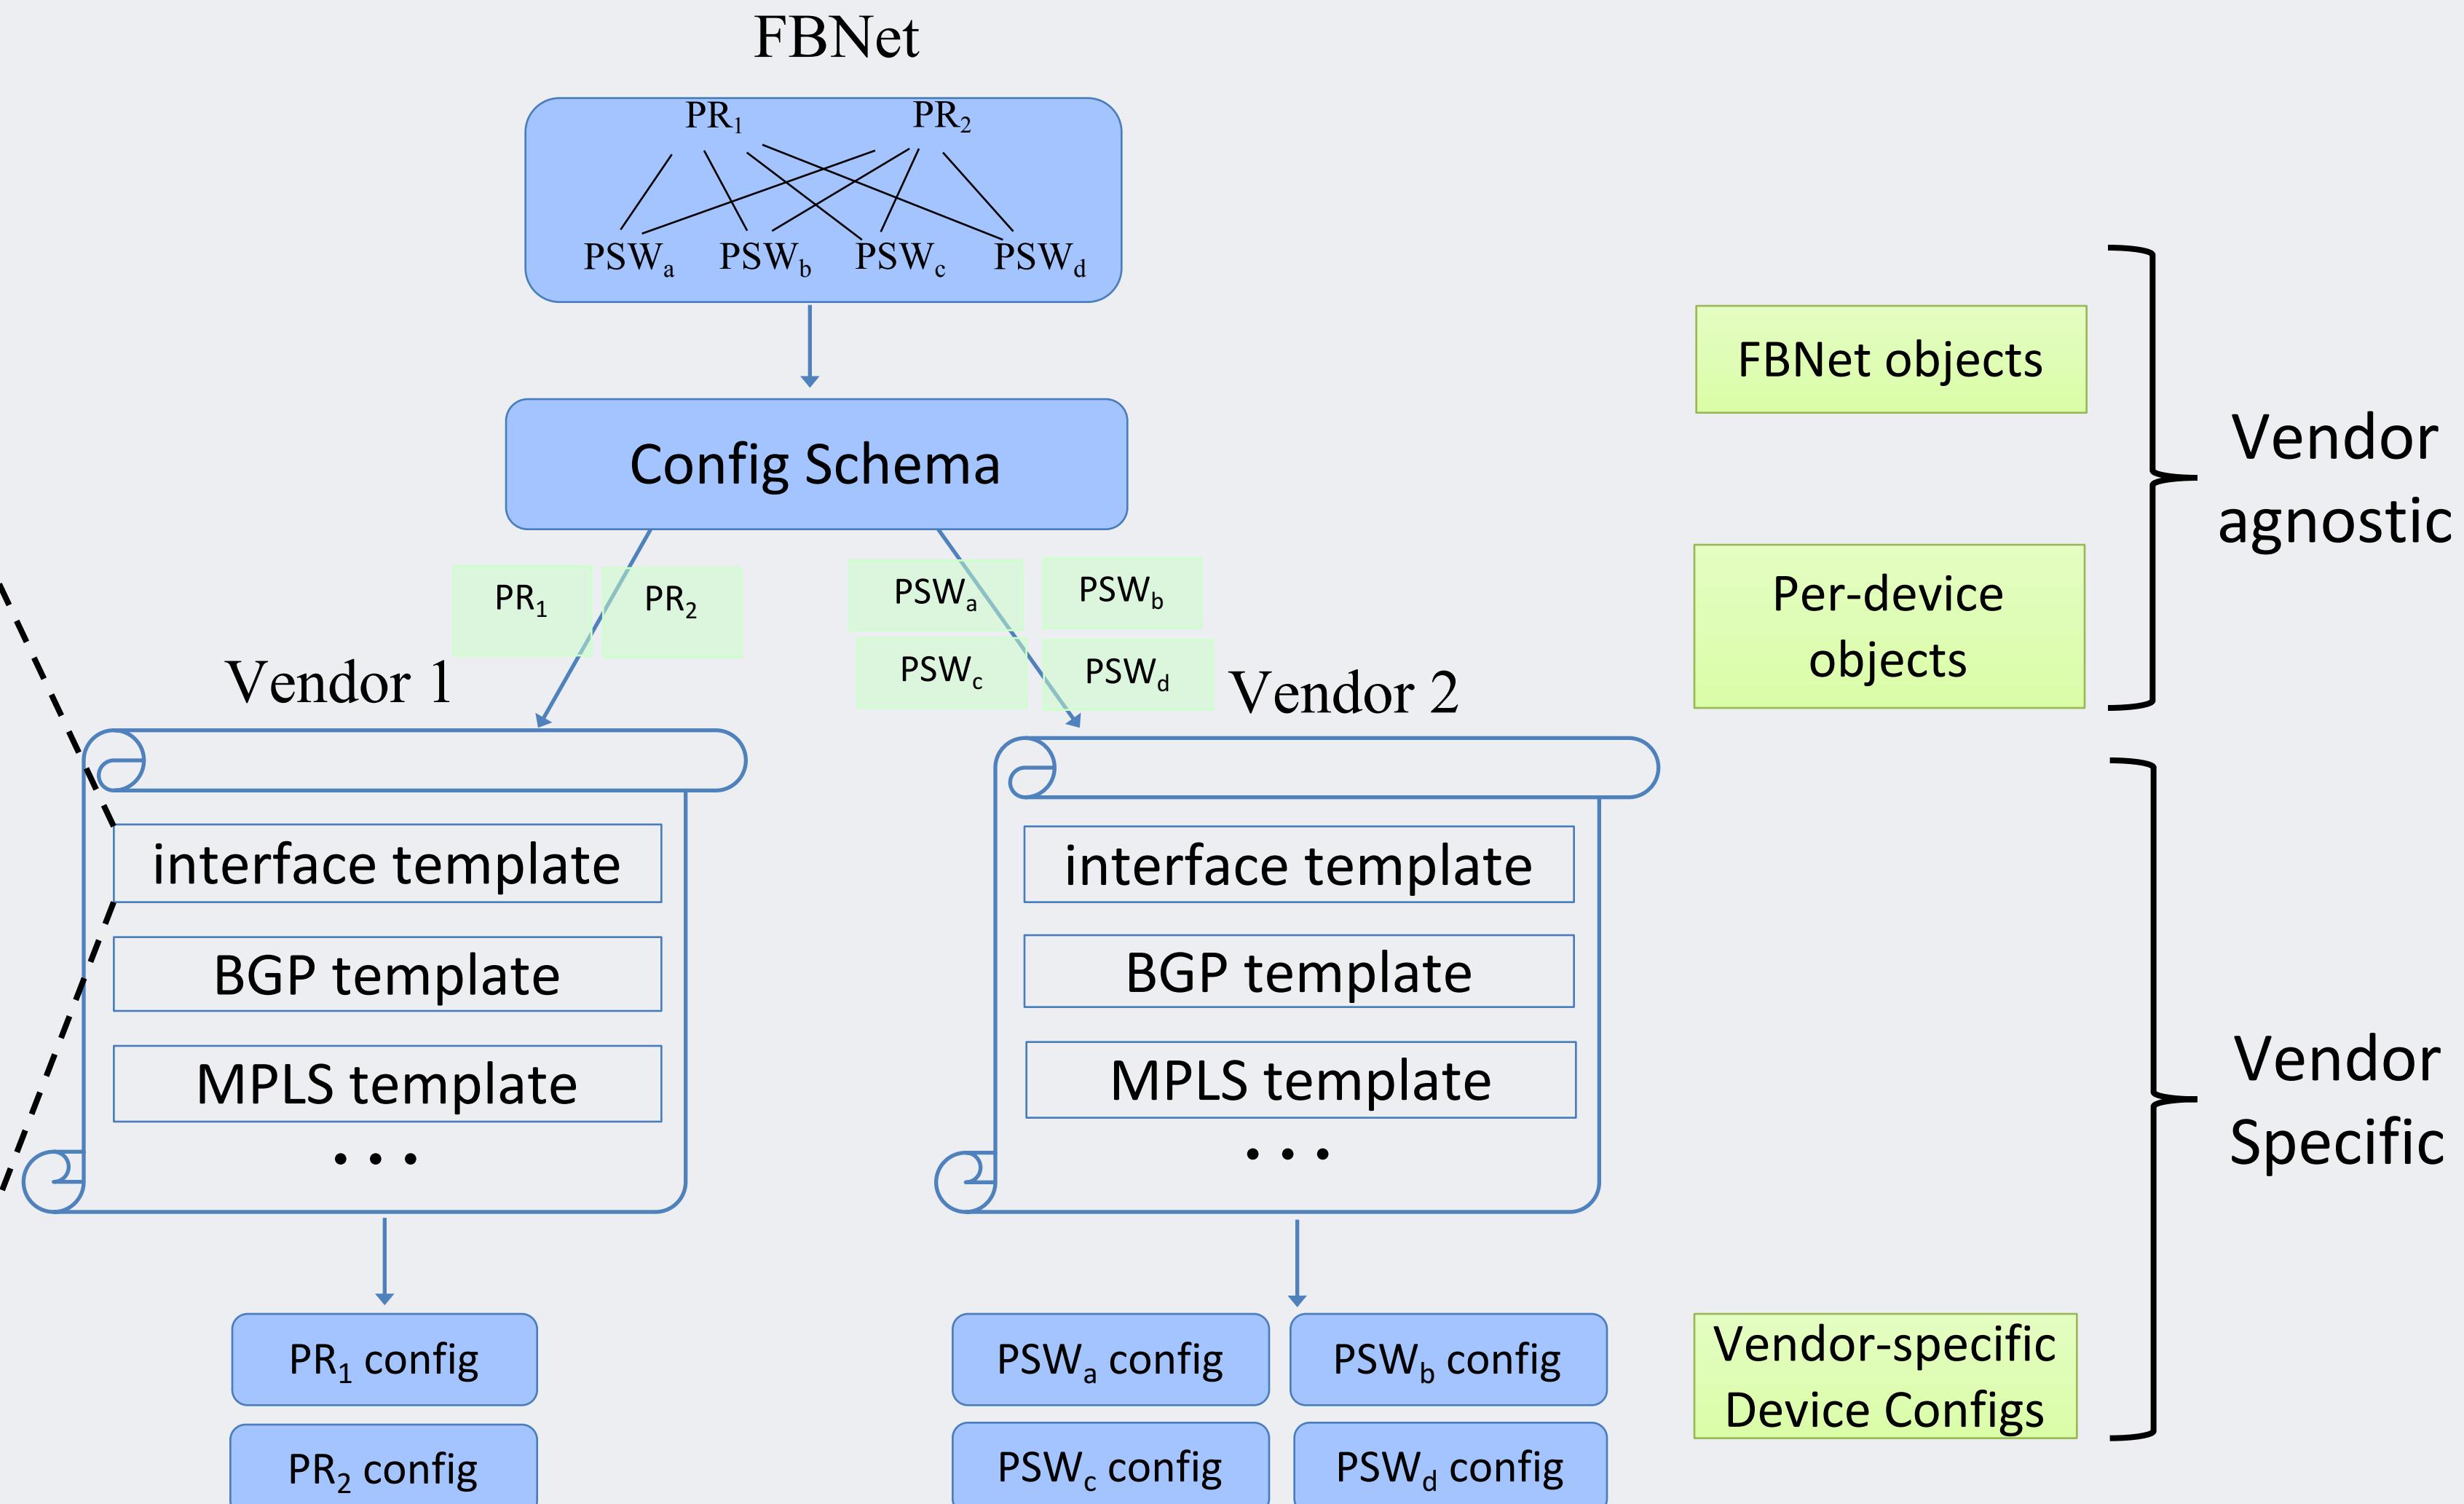
```
struct Device {  
    1: list<AggregatedInterface> aggs,  
}  
  
struct AggregatedInterface {  
    1: string name,  
    2: i32 number,  
    3: string v4_prefix,  
    4: string v6_prefix,  
    5: list<PhysicalInterface> pifs,  
}  
  
struct PhysicalInterface {  
    1: string name,  
}
```



Config Generation

FBNet objects → Device configs

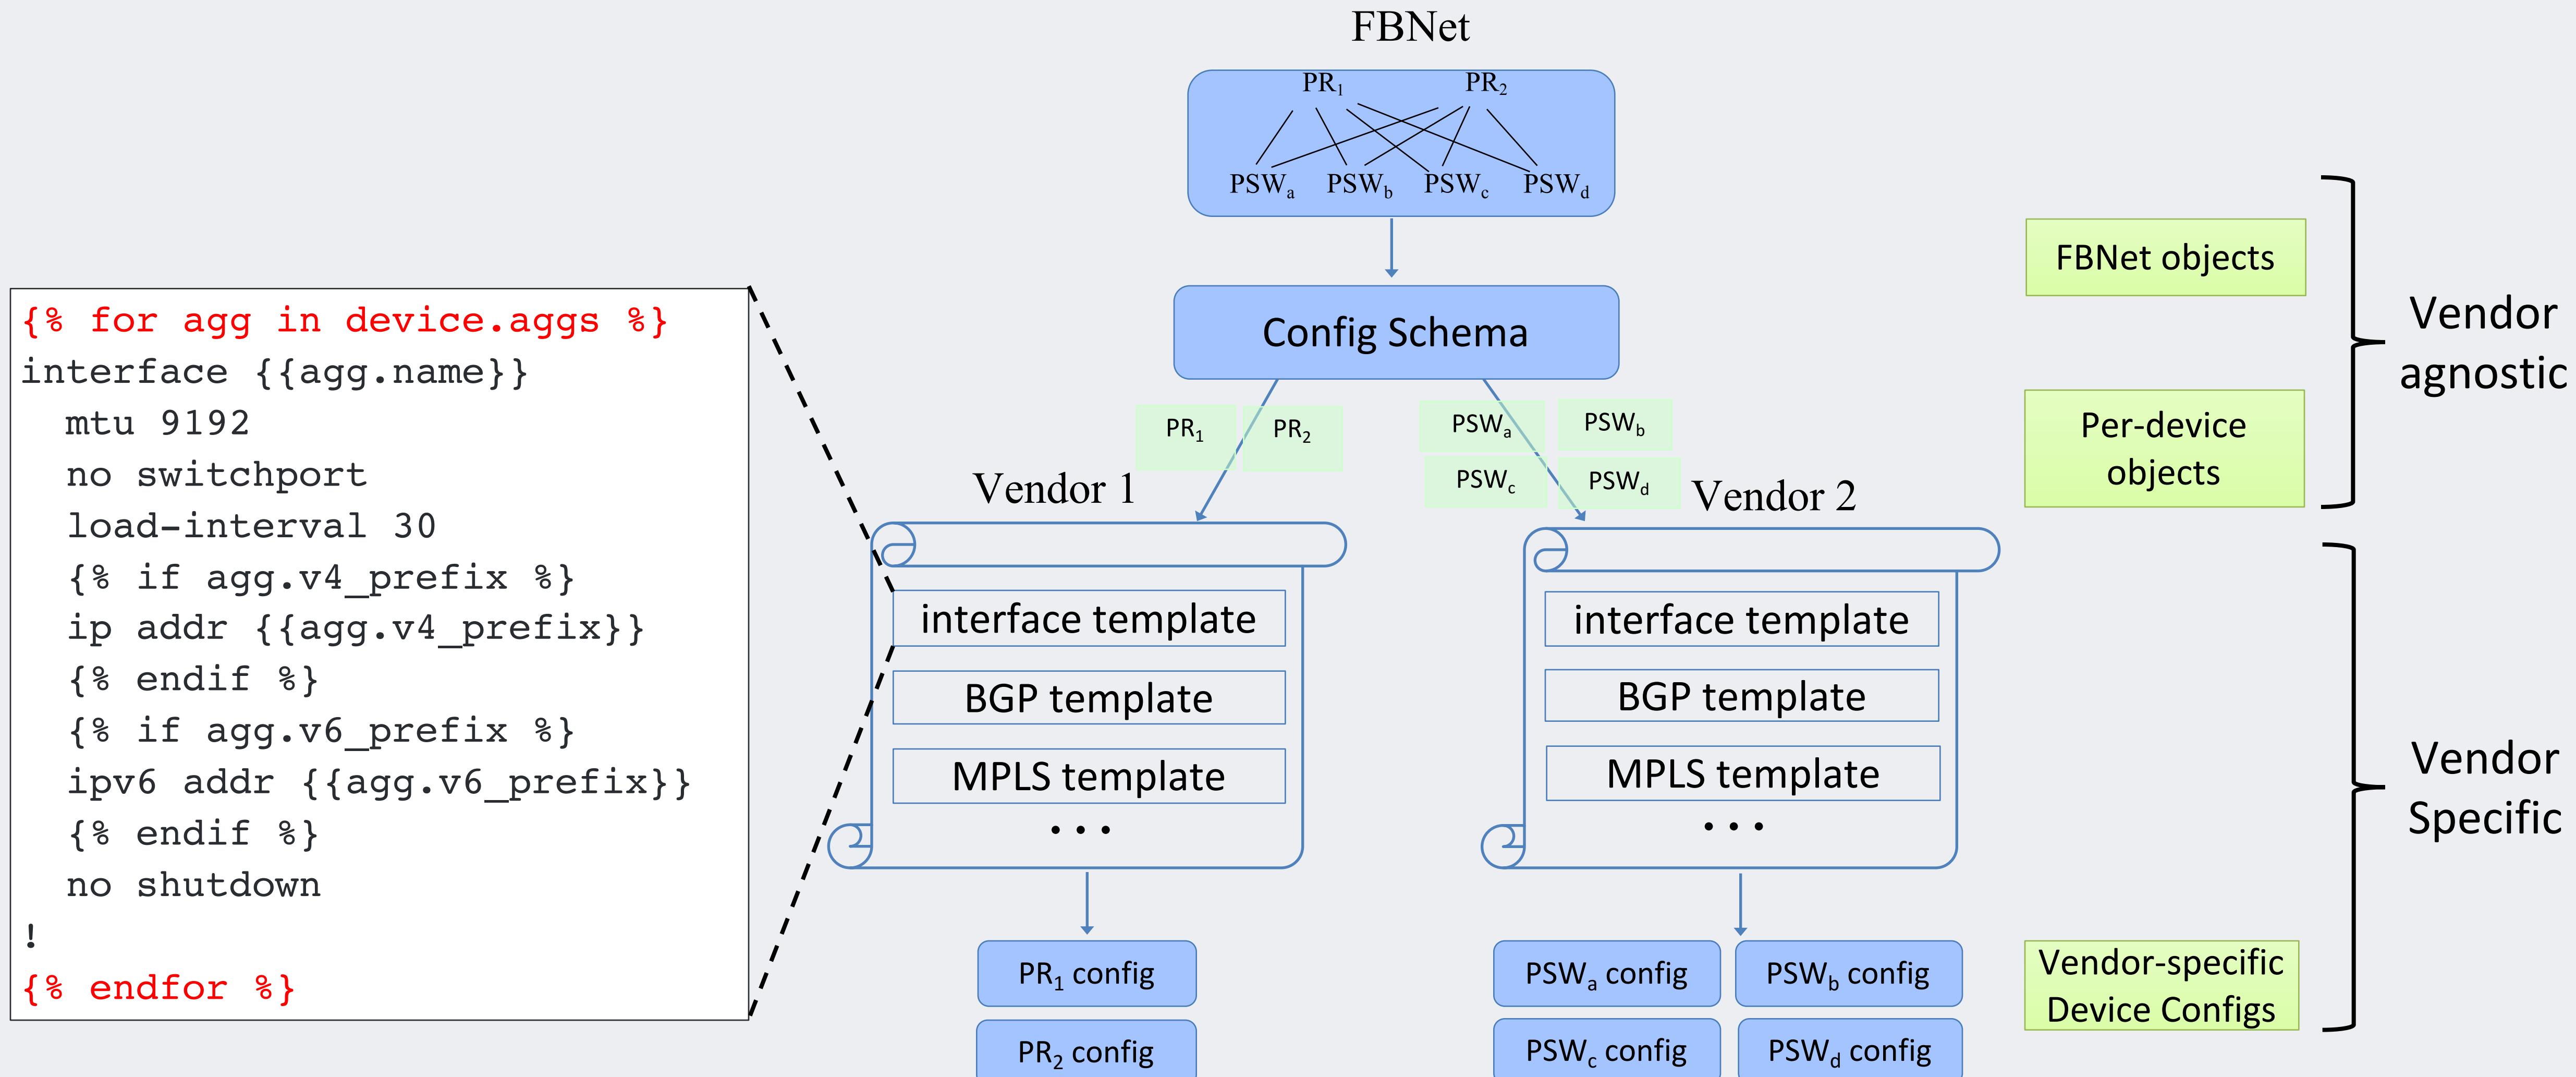
```
{% for agg in device.aggs %}  
interface {{agg.name}}  
  mtu 9192  
  no switchport  
  load-interval 30  
  {% if agg.v4_prefix %}  
    ip addr {{agg.v4_prefix}}  
  {% endif %}  
  {% if agg.v6_prefix %}  
    ipv6 addr {{agg.v6_prefix}}  
  {% endif %}  
  no shutdown  
!  
{% endfor %}
```



Config Generation

FBNet objects → Device configs

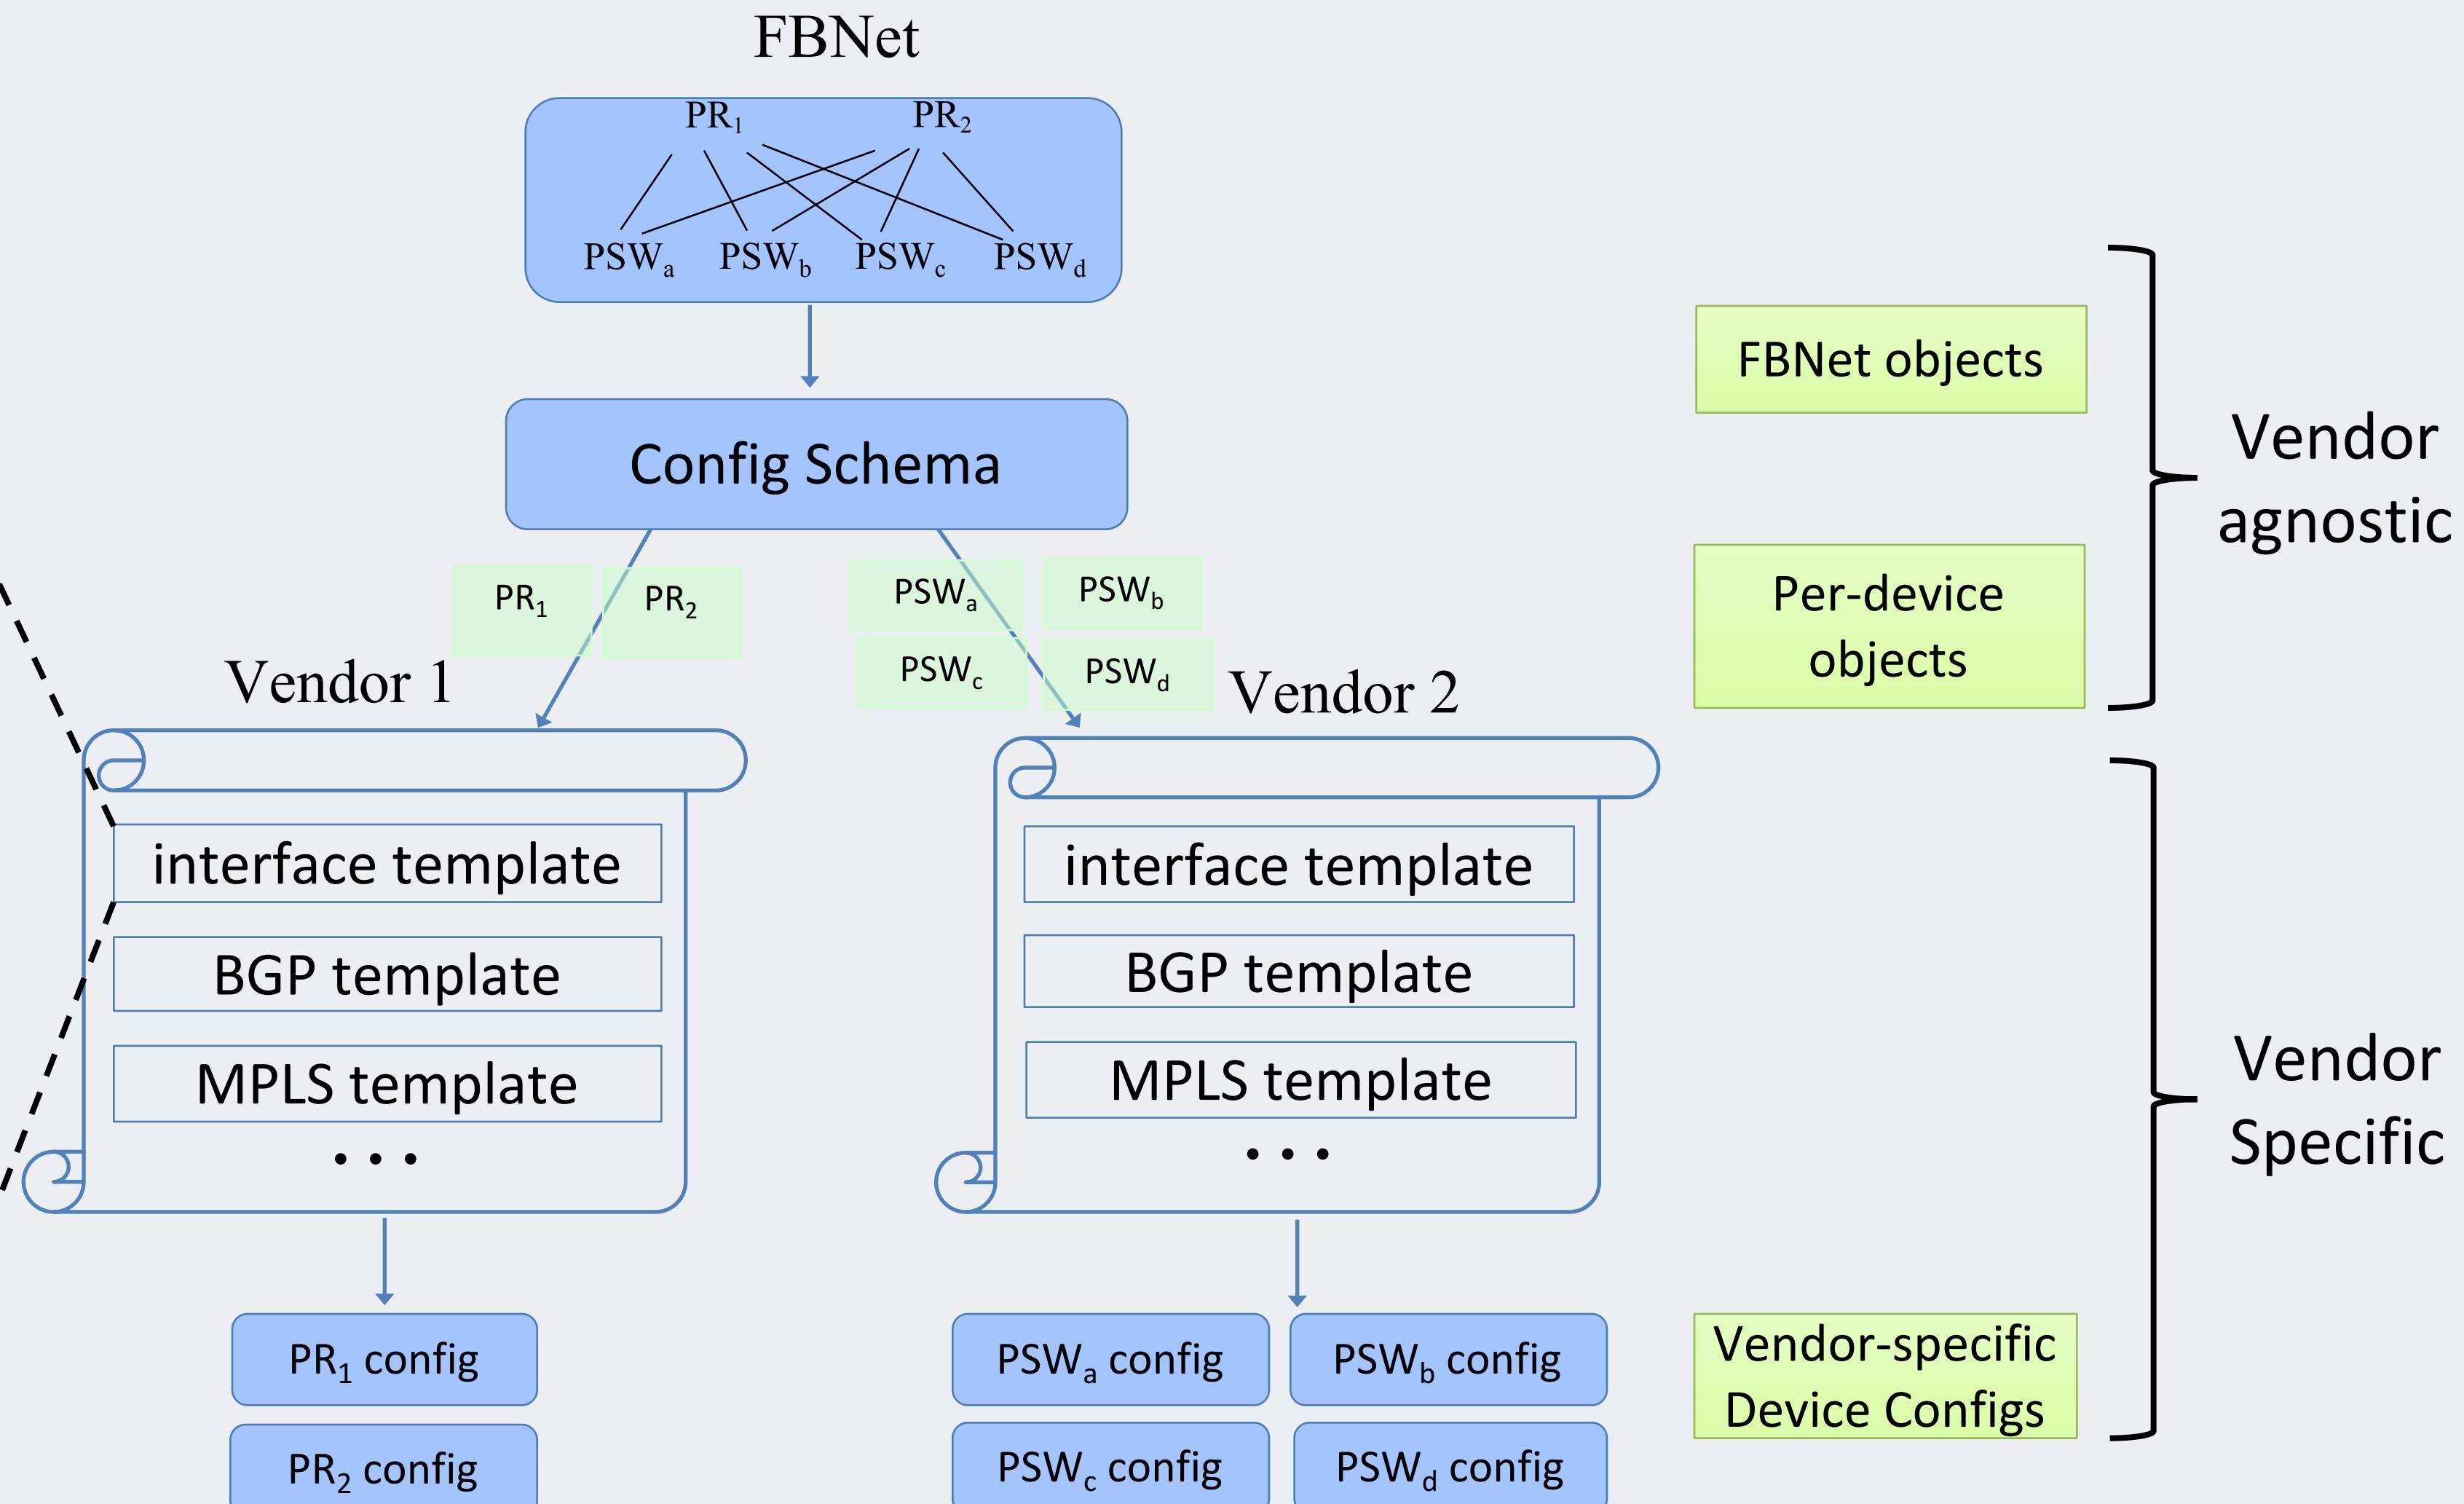
```
{% for agg in device.aggs %}  
interface {{agg.name}}  
mtu 9192  
no switchport  
load-interval 30  
{% if agg.v4_prefix %}  
ip addr {{agg.v4_prefix}}  
{% endif %}  
{% if agg.v6_prefix %}  
ipv6 addr {{agg.v6_prefix}}  
{% endif %}  
no shutdown  
!  
{% endfor %}
```



Config Generation

FBNet objects → Device configs

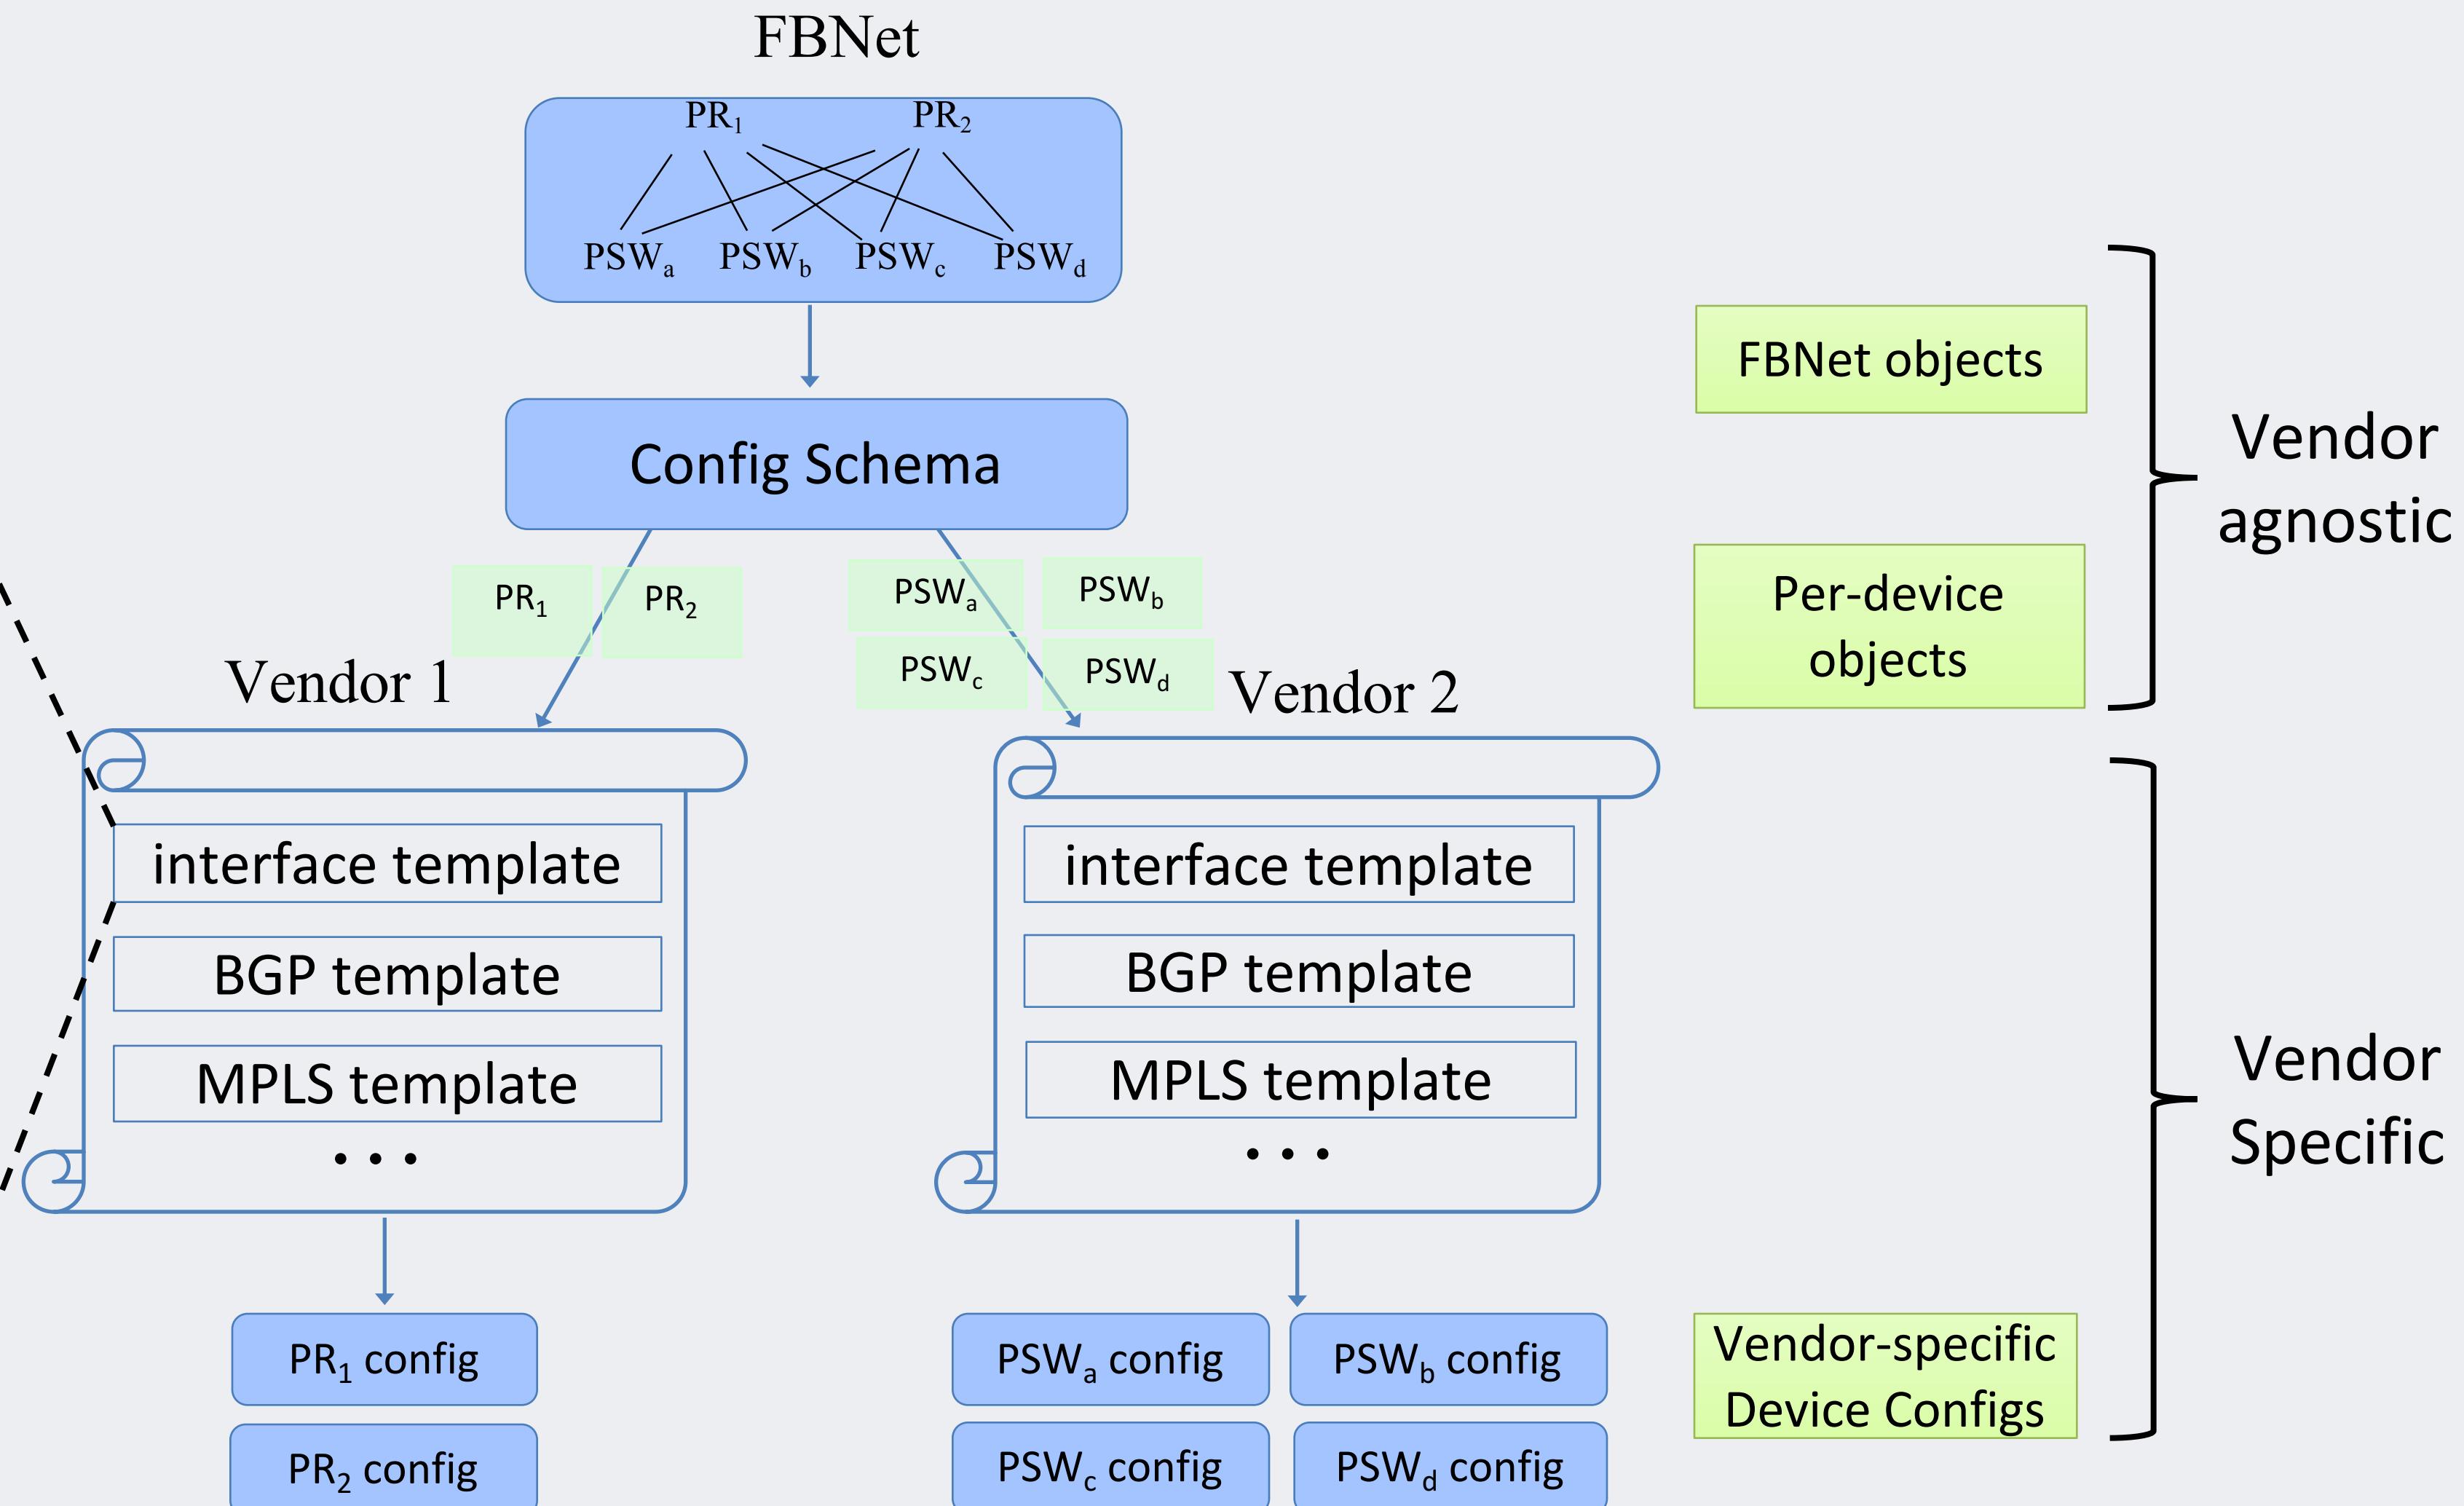
```
{% for agg in device.aggs %}  
interface {{agg.name}}  
  mtu 9192  
  no switchport  
  load-interval 30  
  {% if agg.v4_prefix %}  
    ip addr {{agg.v4_prefix}}  
  {% endif %}  
  {% if agg.v6_prefix %}  
    ipv6 addr {{agg.v6_prefix}}  
  {% endif %}  
  no shutdown  
!  
{% endfor %}
```



Config Generation

FBNet objects → Device configs

```
{% for agg in device.aggs %}  
interface {{agg.name}}  
  mtu 9192  
  no switchport  
  load-interval 30  
  {% if agg.v4_prefix %}  
    ip addr {{agg.v4_prefix}}  
  {% endif %}  
  {% if agg.v6_prefix %}  
    ipv6 addr {{agg.v6_prefix}}  
  {% endif %}  
  no shutdown  
!  
{% endfor %}
```

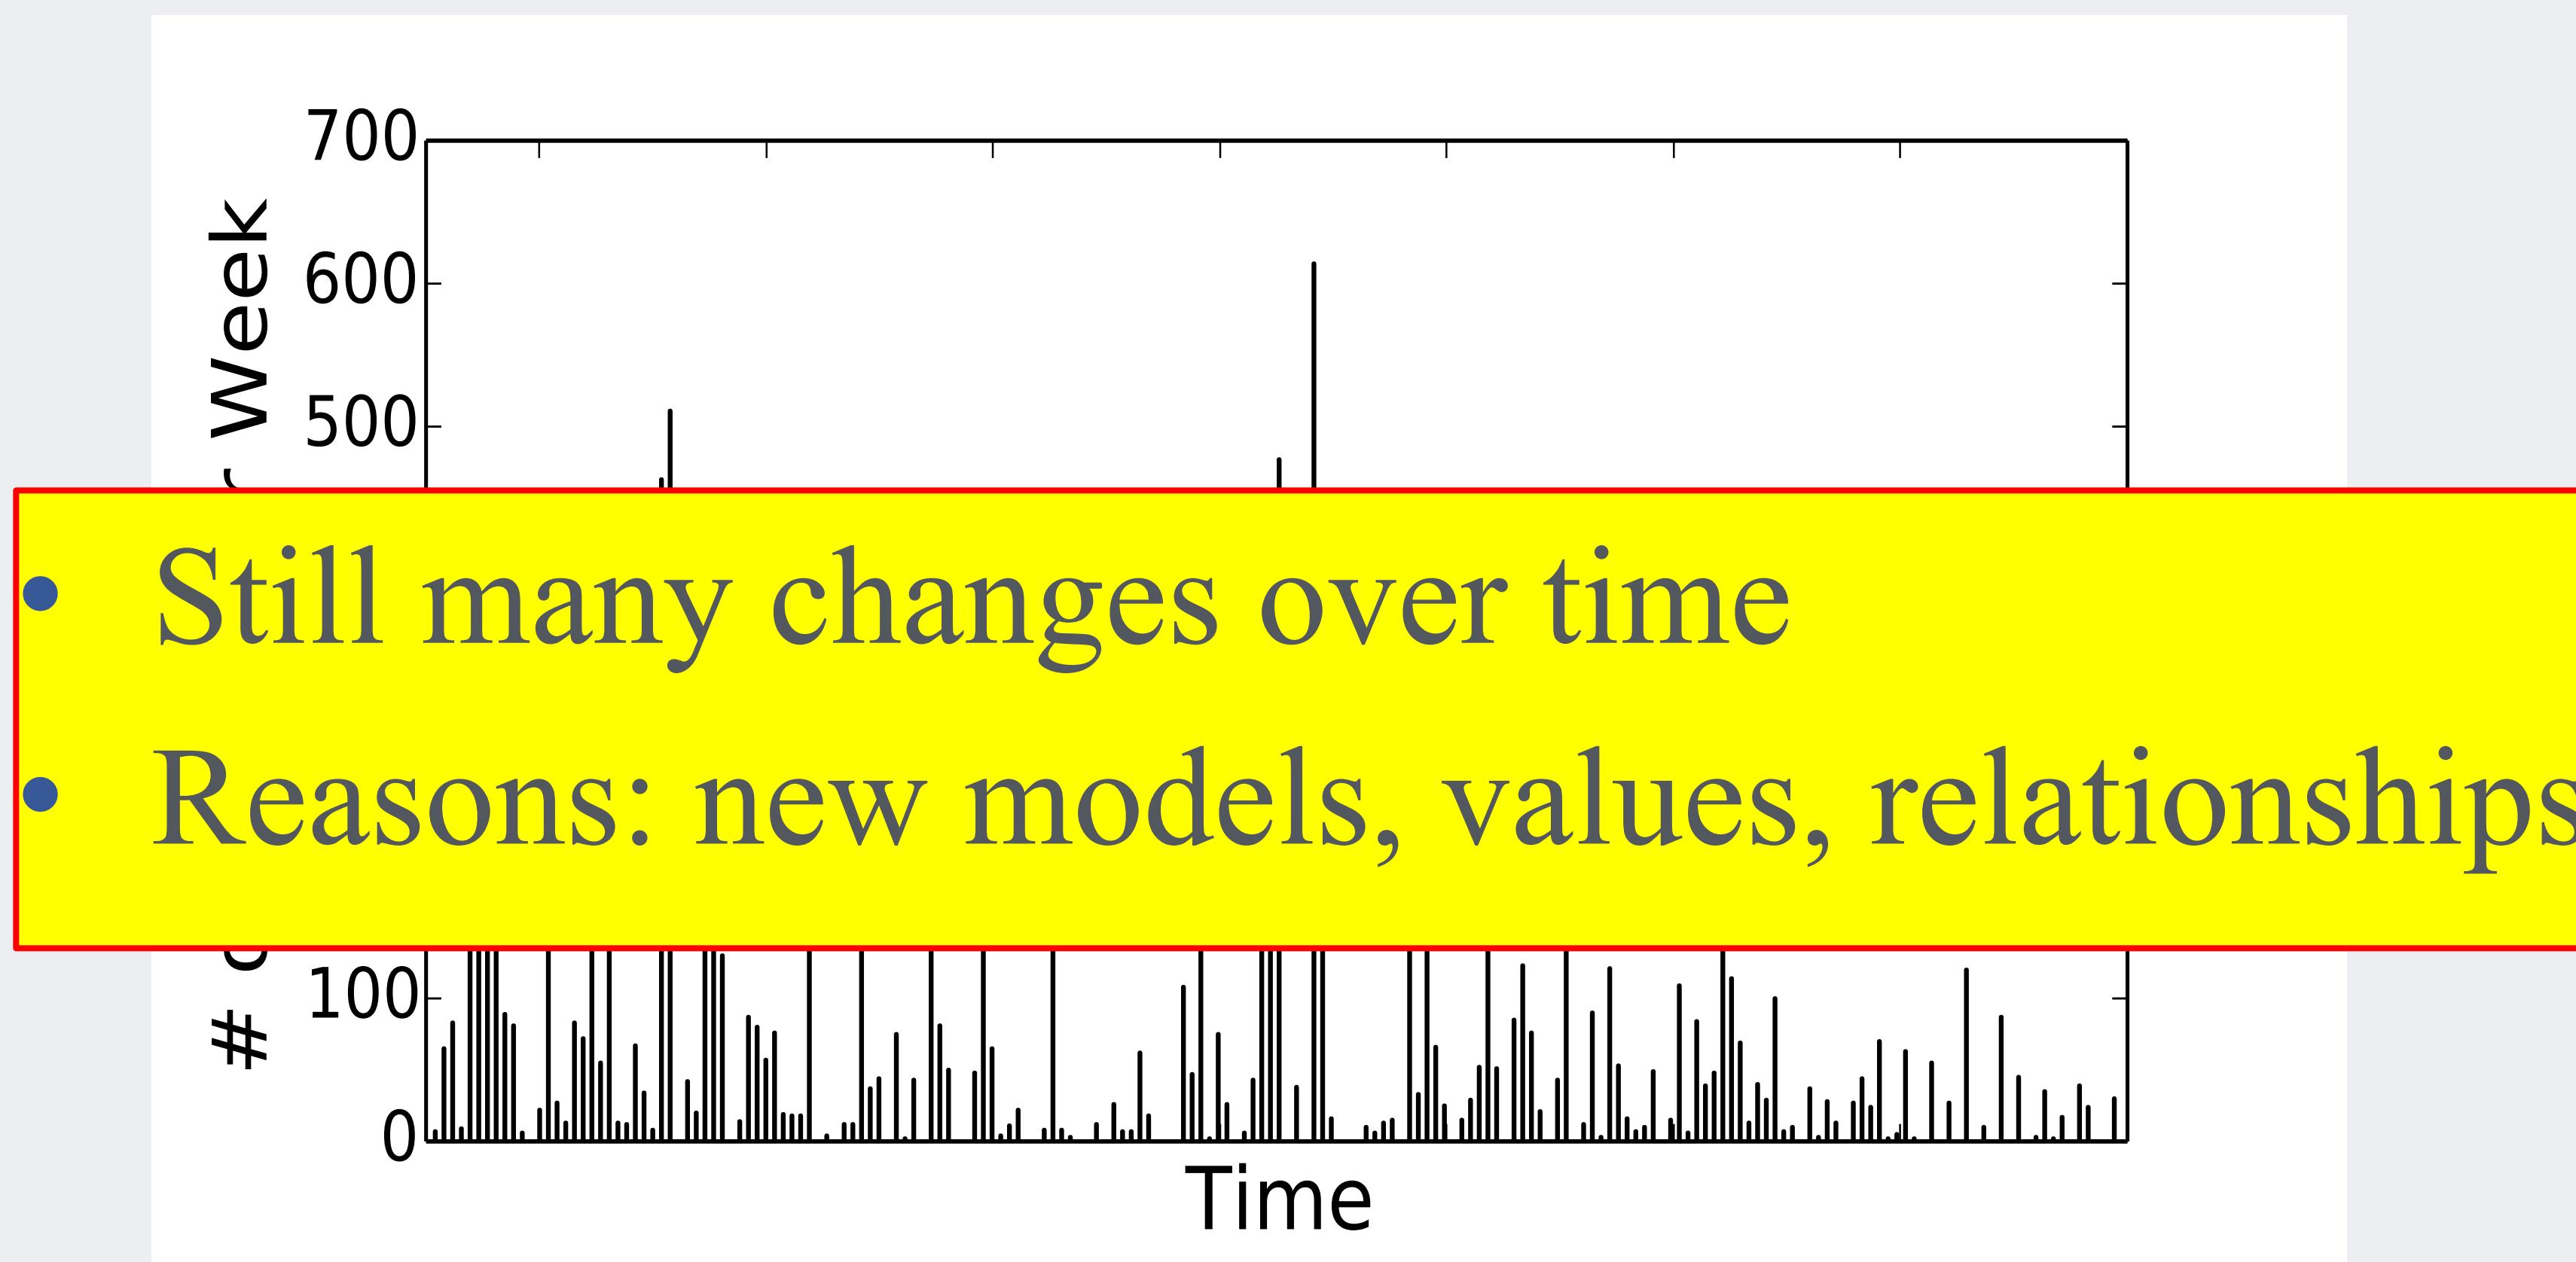


Usage Statistics

- # of FBNet model change?
- # changed FBNet objects per design change?
- Frequency and size of config change?

FBNet Model Changes

How much does FBNet model change over time?

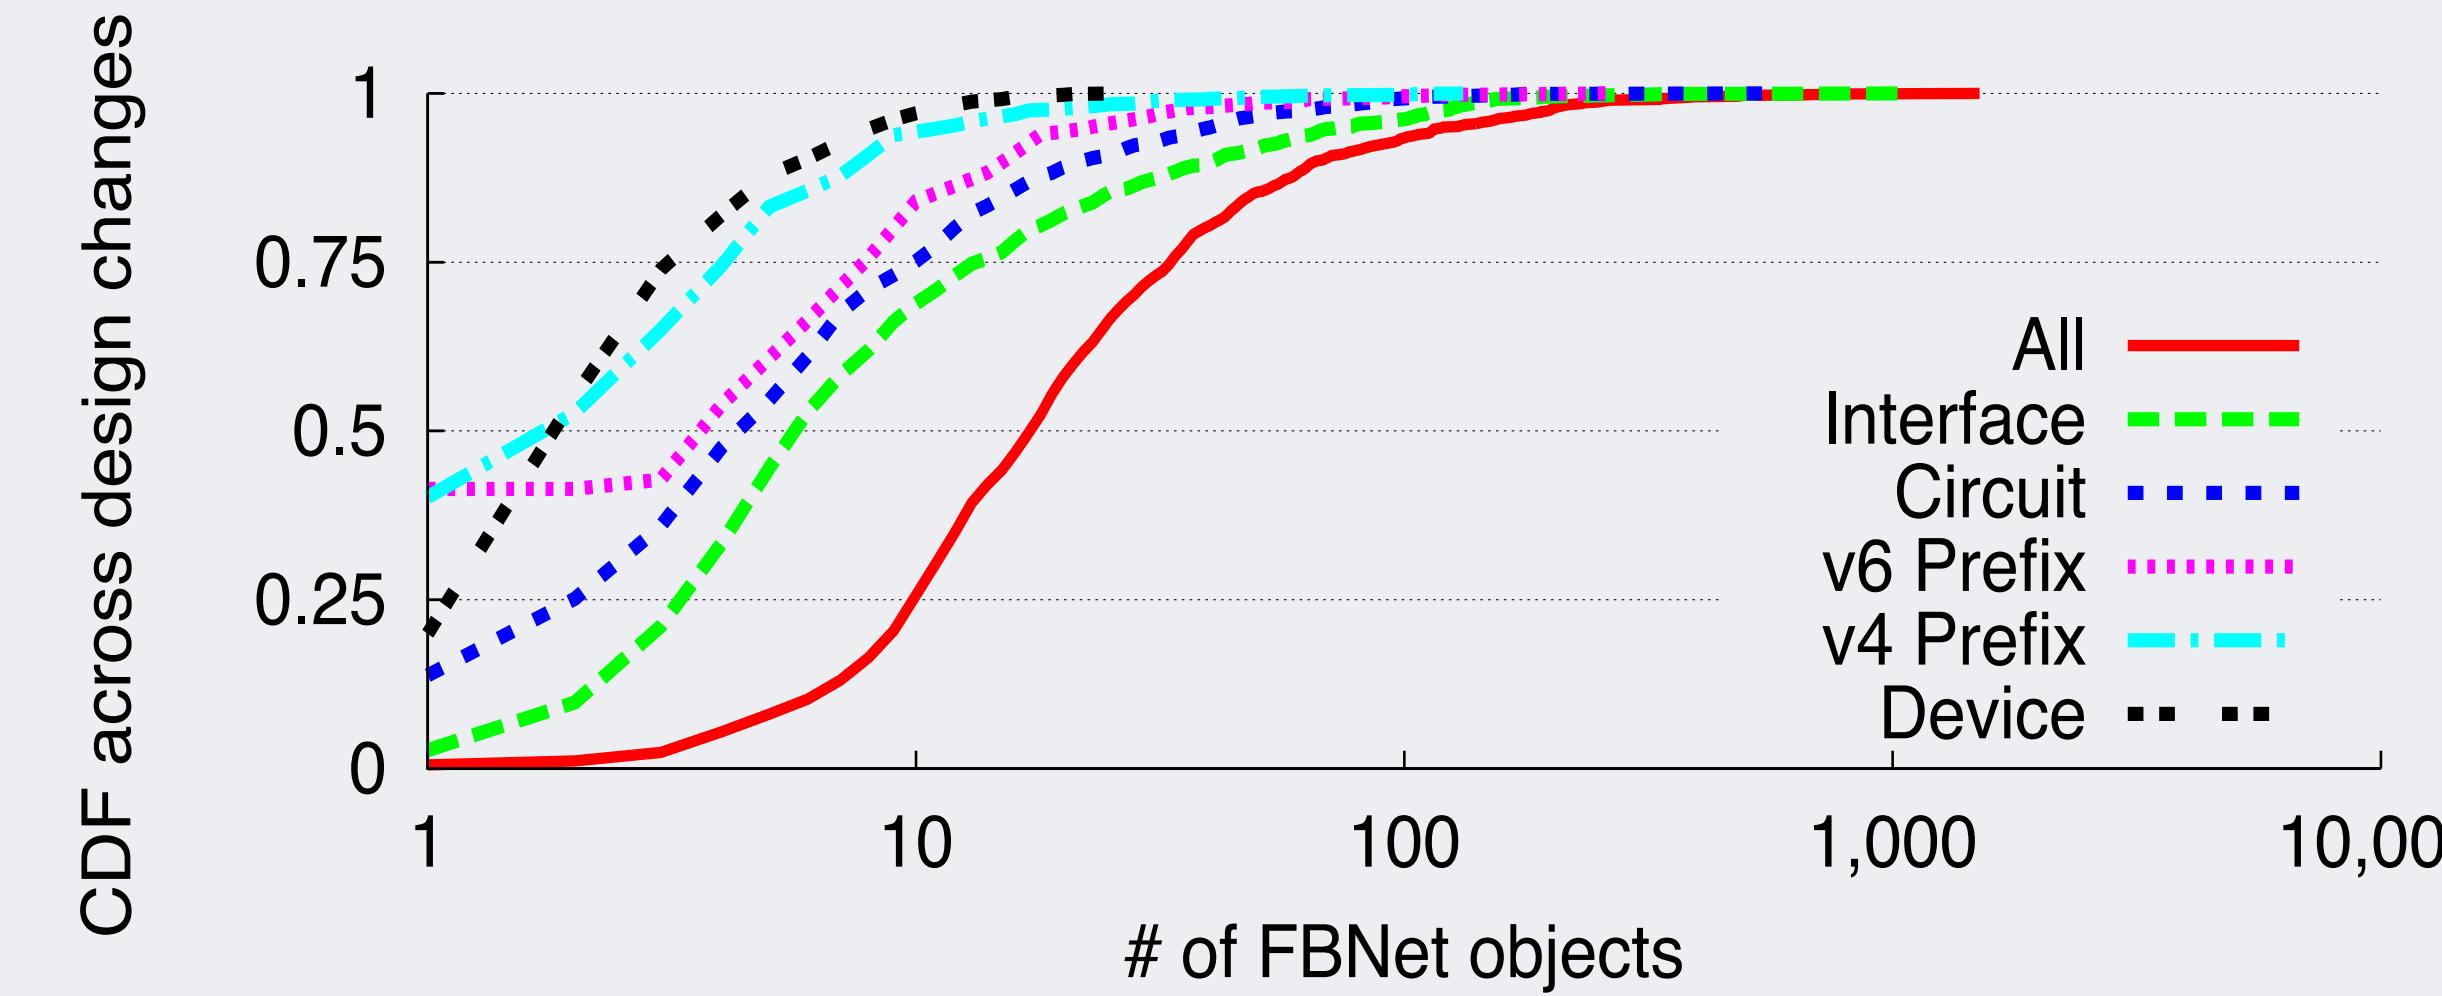
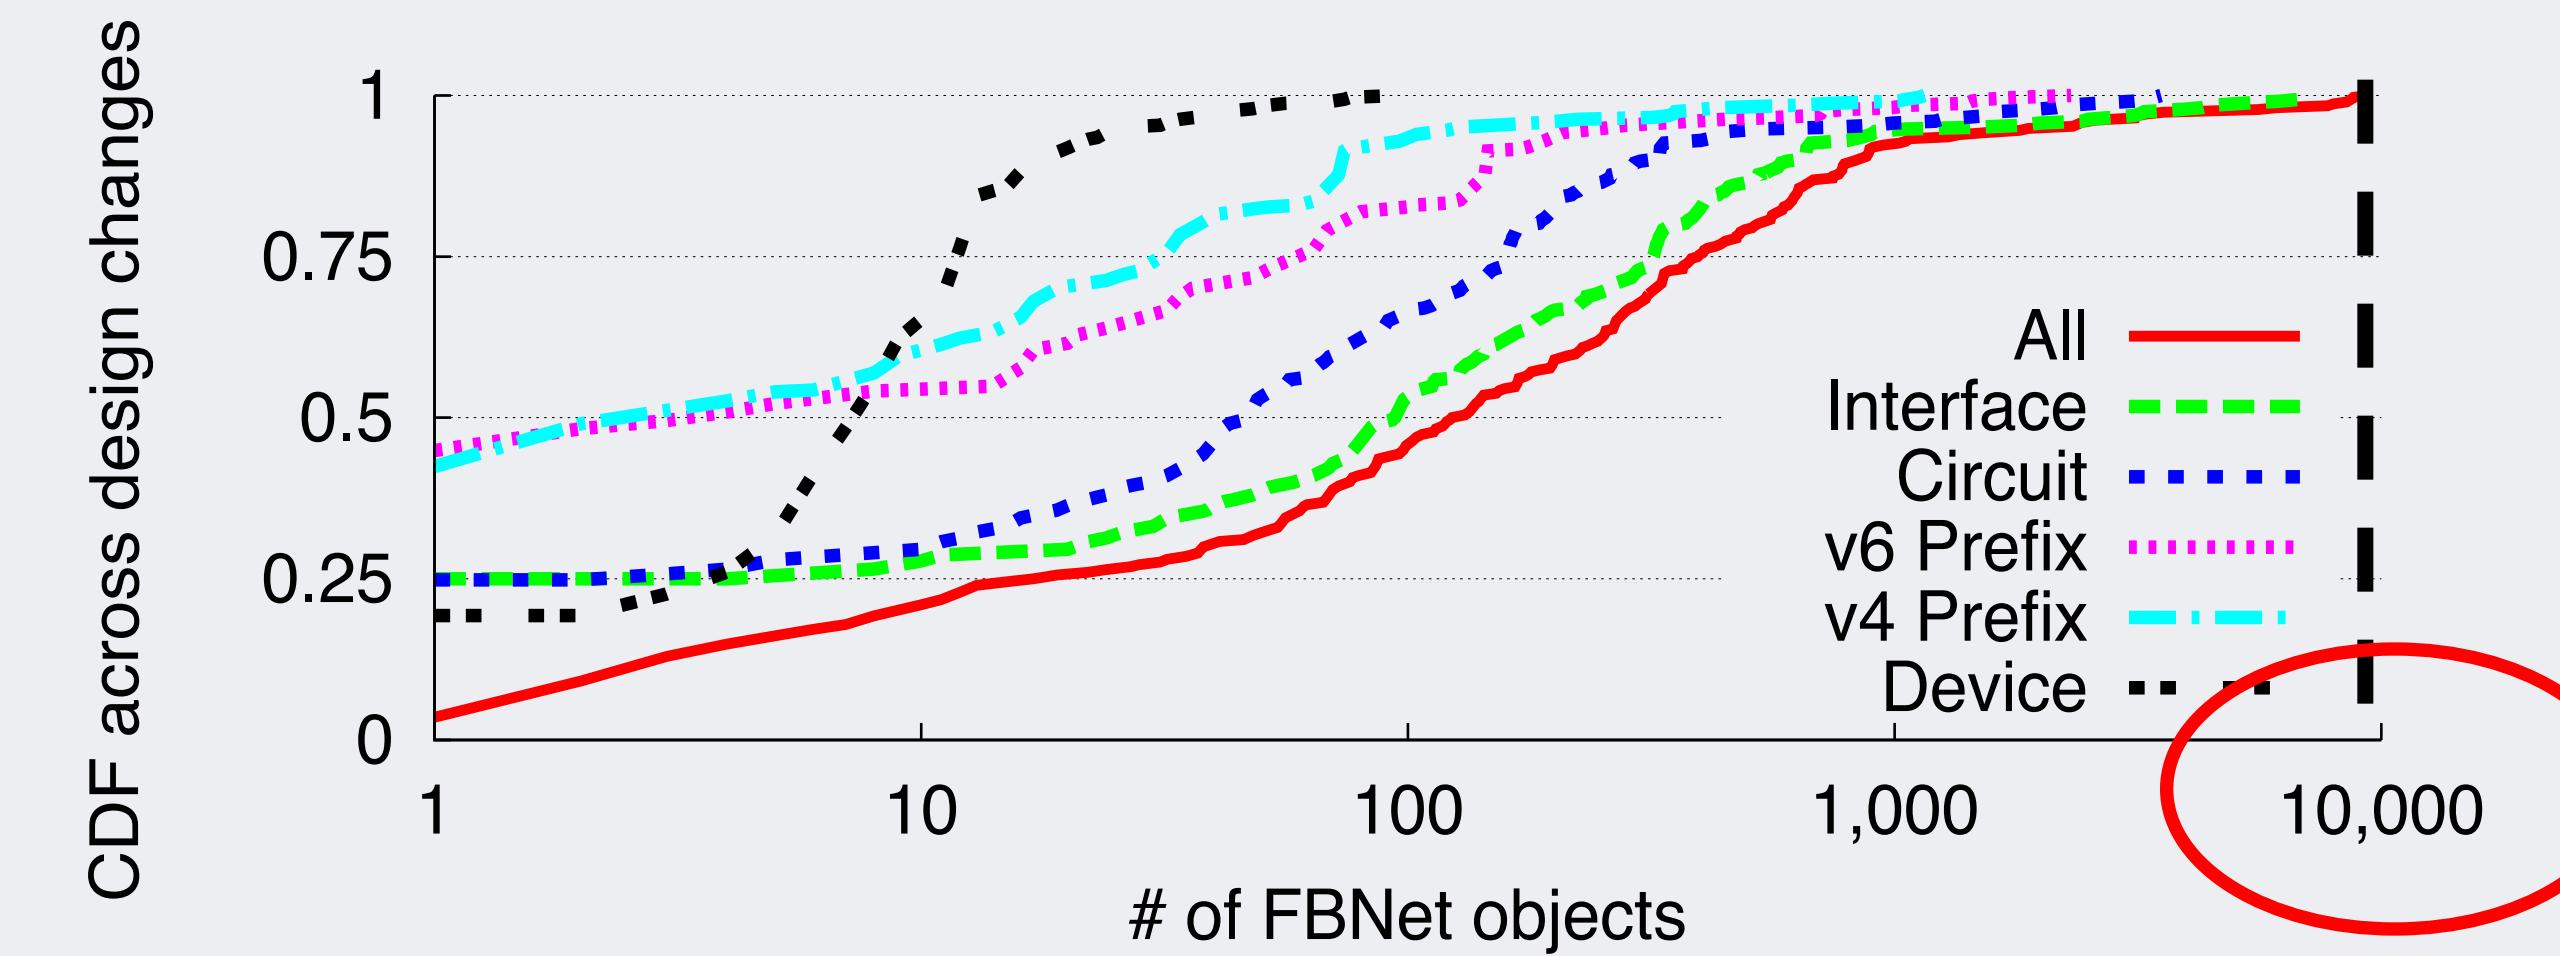


Design Changes

How many FBNet object are changed per design change?

POP/DC

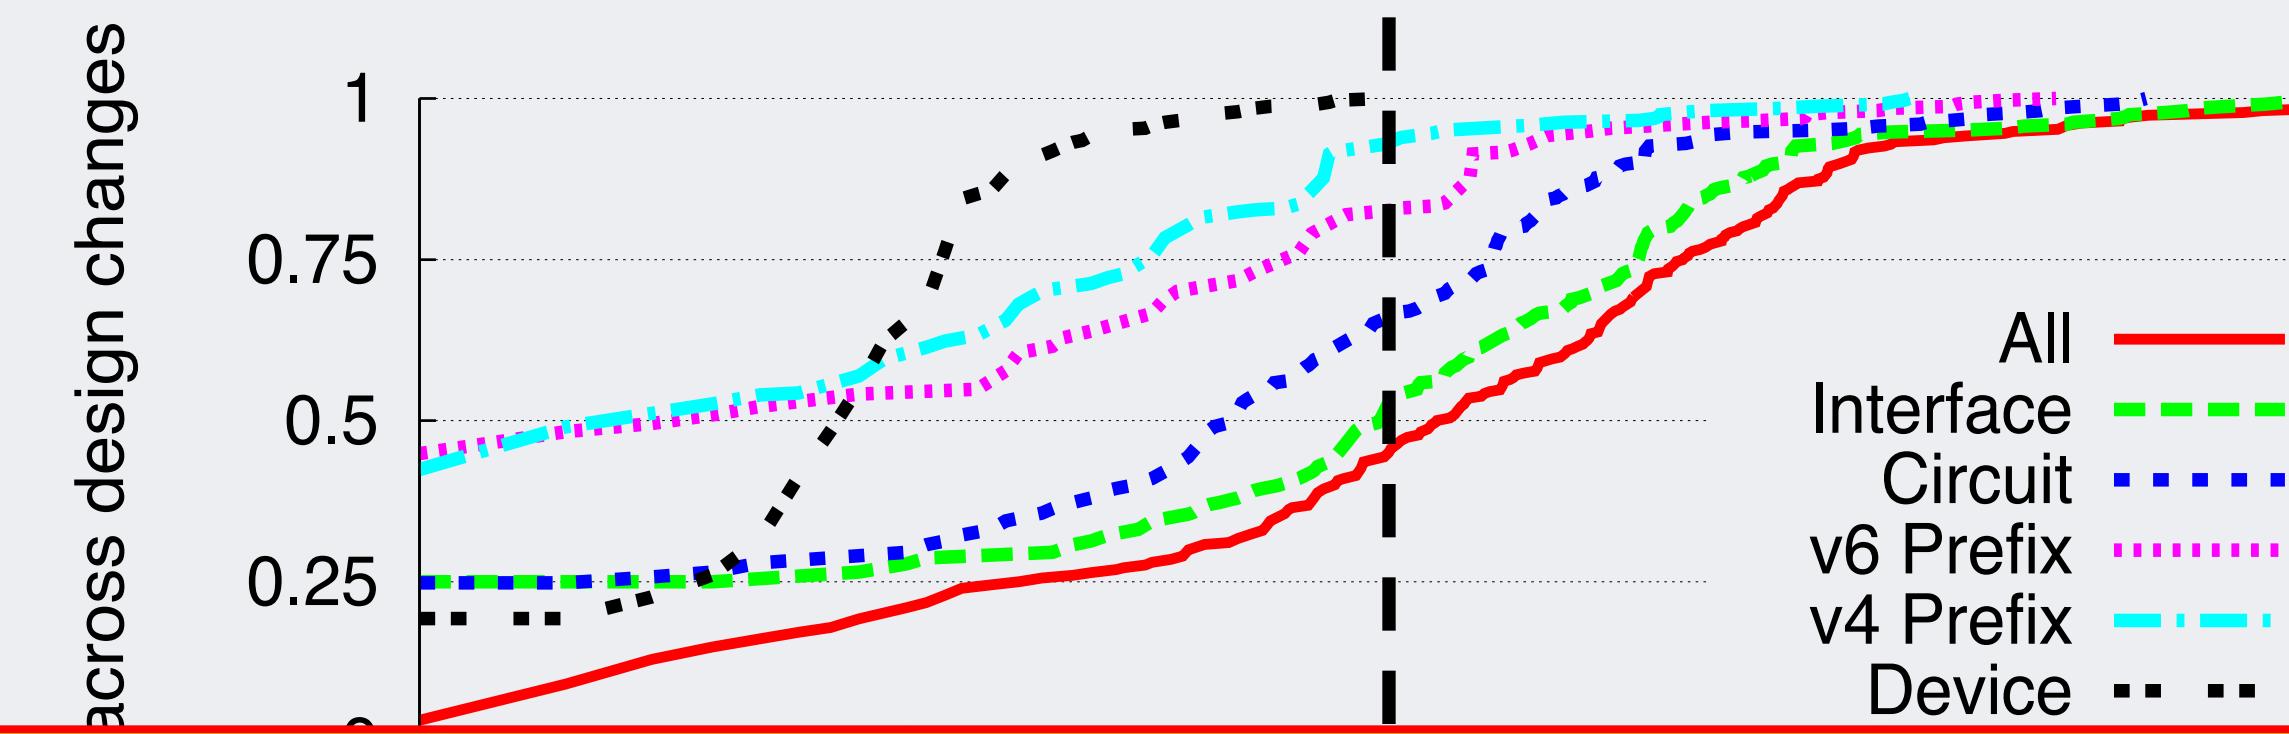
Backbone



Design Changes

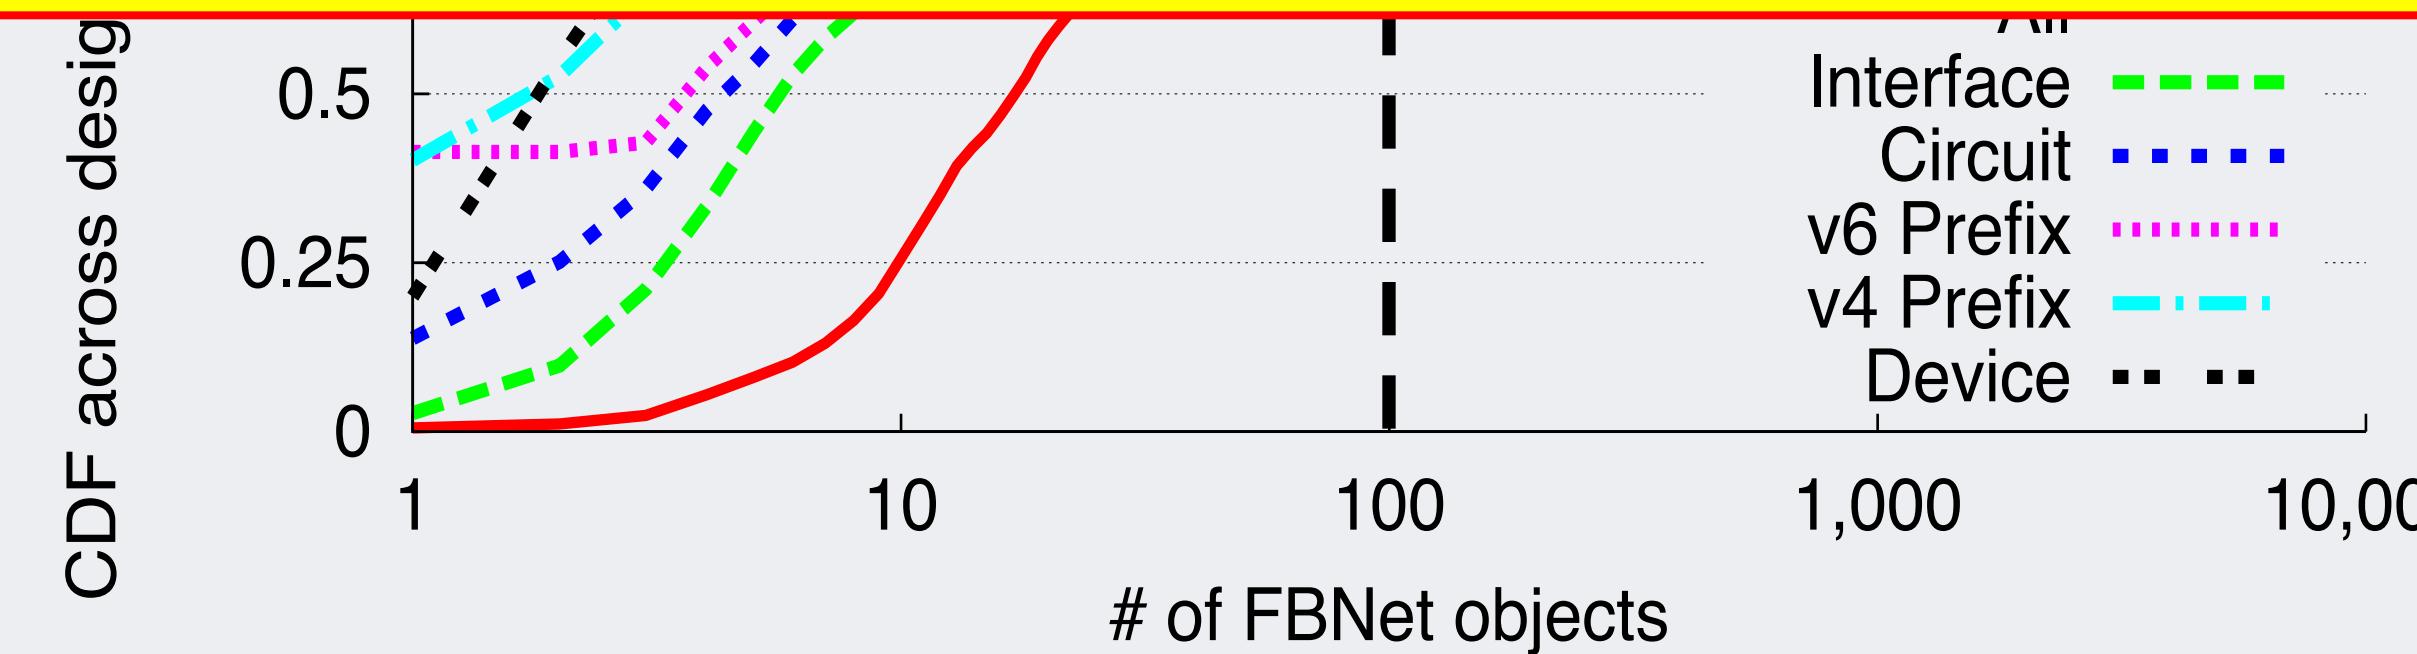
How many FBNet object are changed per design change?

POP/DC



- POP/DC: bigger design changes
- Backbone: smaller design changes

Backbone



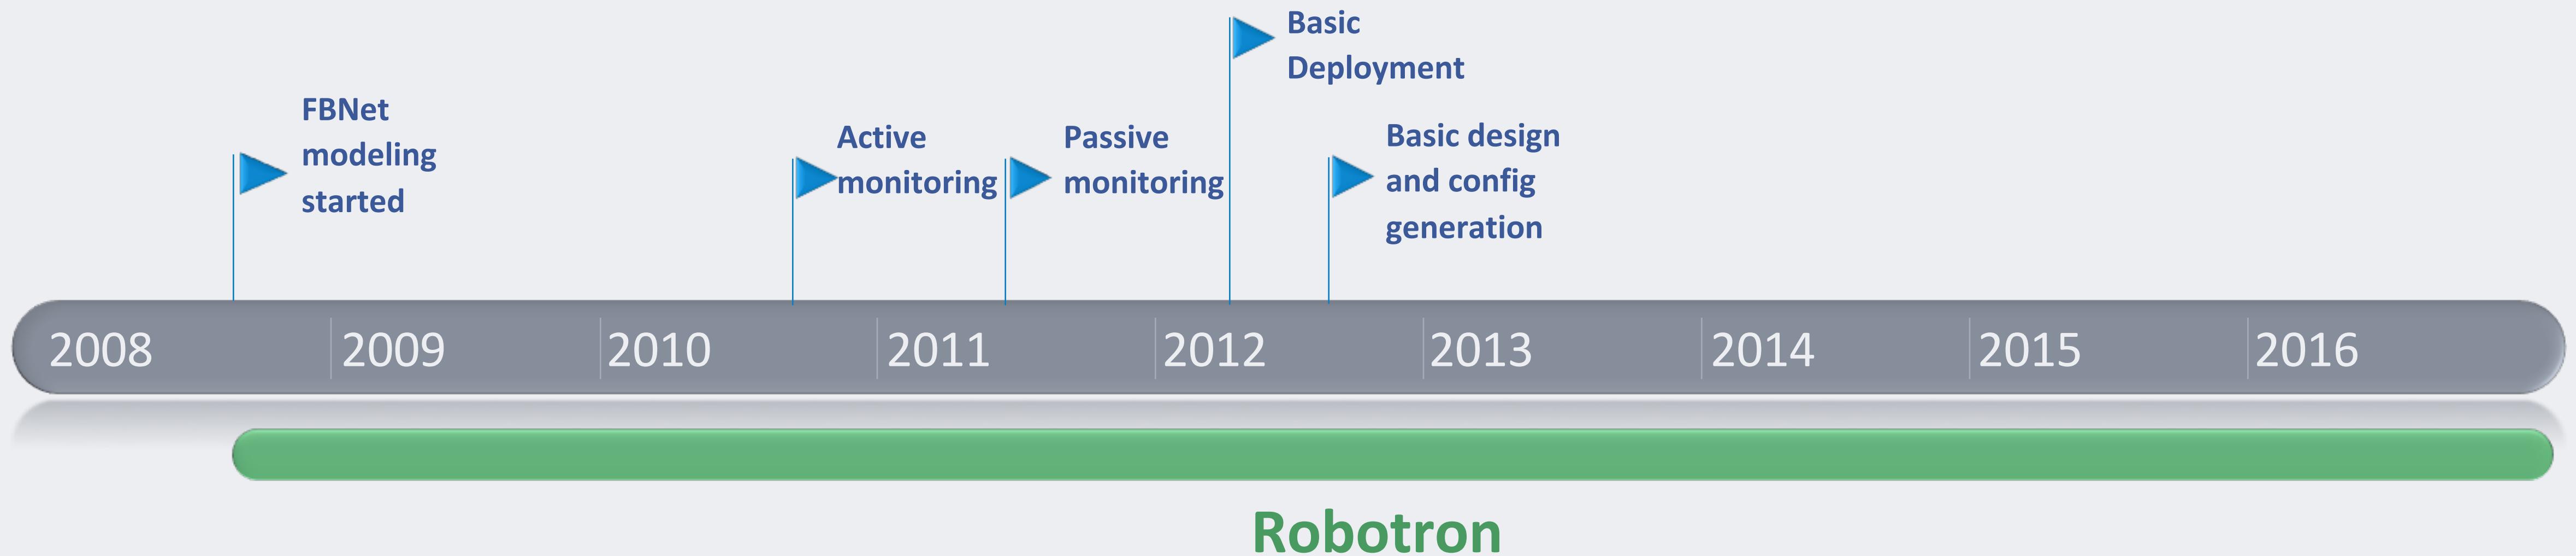
Configuration Changes

What's the frequency and size of configuration change?

- Median number of config lines changed per week
 - POP/DC devices: 500 lines
 - Backbone: few bigger config changes
- Avg
 - Backbone: many smaller config changes
 - POP/DC devices: 2.55
 - Backbone devices: 12.46

Evolution of Robotron

Bottom-up, experience driven



Experience: Modeling is laborious

Problem Scenario: new eBGP session configuration

- A new eBGP session needed a proper import policy
 - Robot saturation of the link
 - Most of the time spent on modeling
- Lesson: Modeling is hard
 - Open problem: Lack of a network model widely accepted by vendors

Experience: Coupling changes is key

Problem Scenario: POP cluster switch turnup

1. An engineer forgot to generate and deploy configuration files.
2. The configuration was deployed.
3. The network failed.
 - Lesson: Network design, config generation and deployment should be tightly coupled
 - Open problem:
 - Atomicity
 - Conflict resolution

Experience: Fallback is important

Problem Scenario: Robotron-less management

- Engineered bypass mechanism is needed
- SSE
- Management
- Logging
- Needed services
- How to safely revert such activities?
- Passively curtail with config monitoring

Conclusion

- First work sharing experience on a production network management system
- Open research problems:
 - Network modeling
 - Atomicity and conflict resolution across management tasks
 - Make network management system work with manual fallback mechanisms

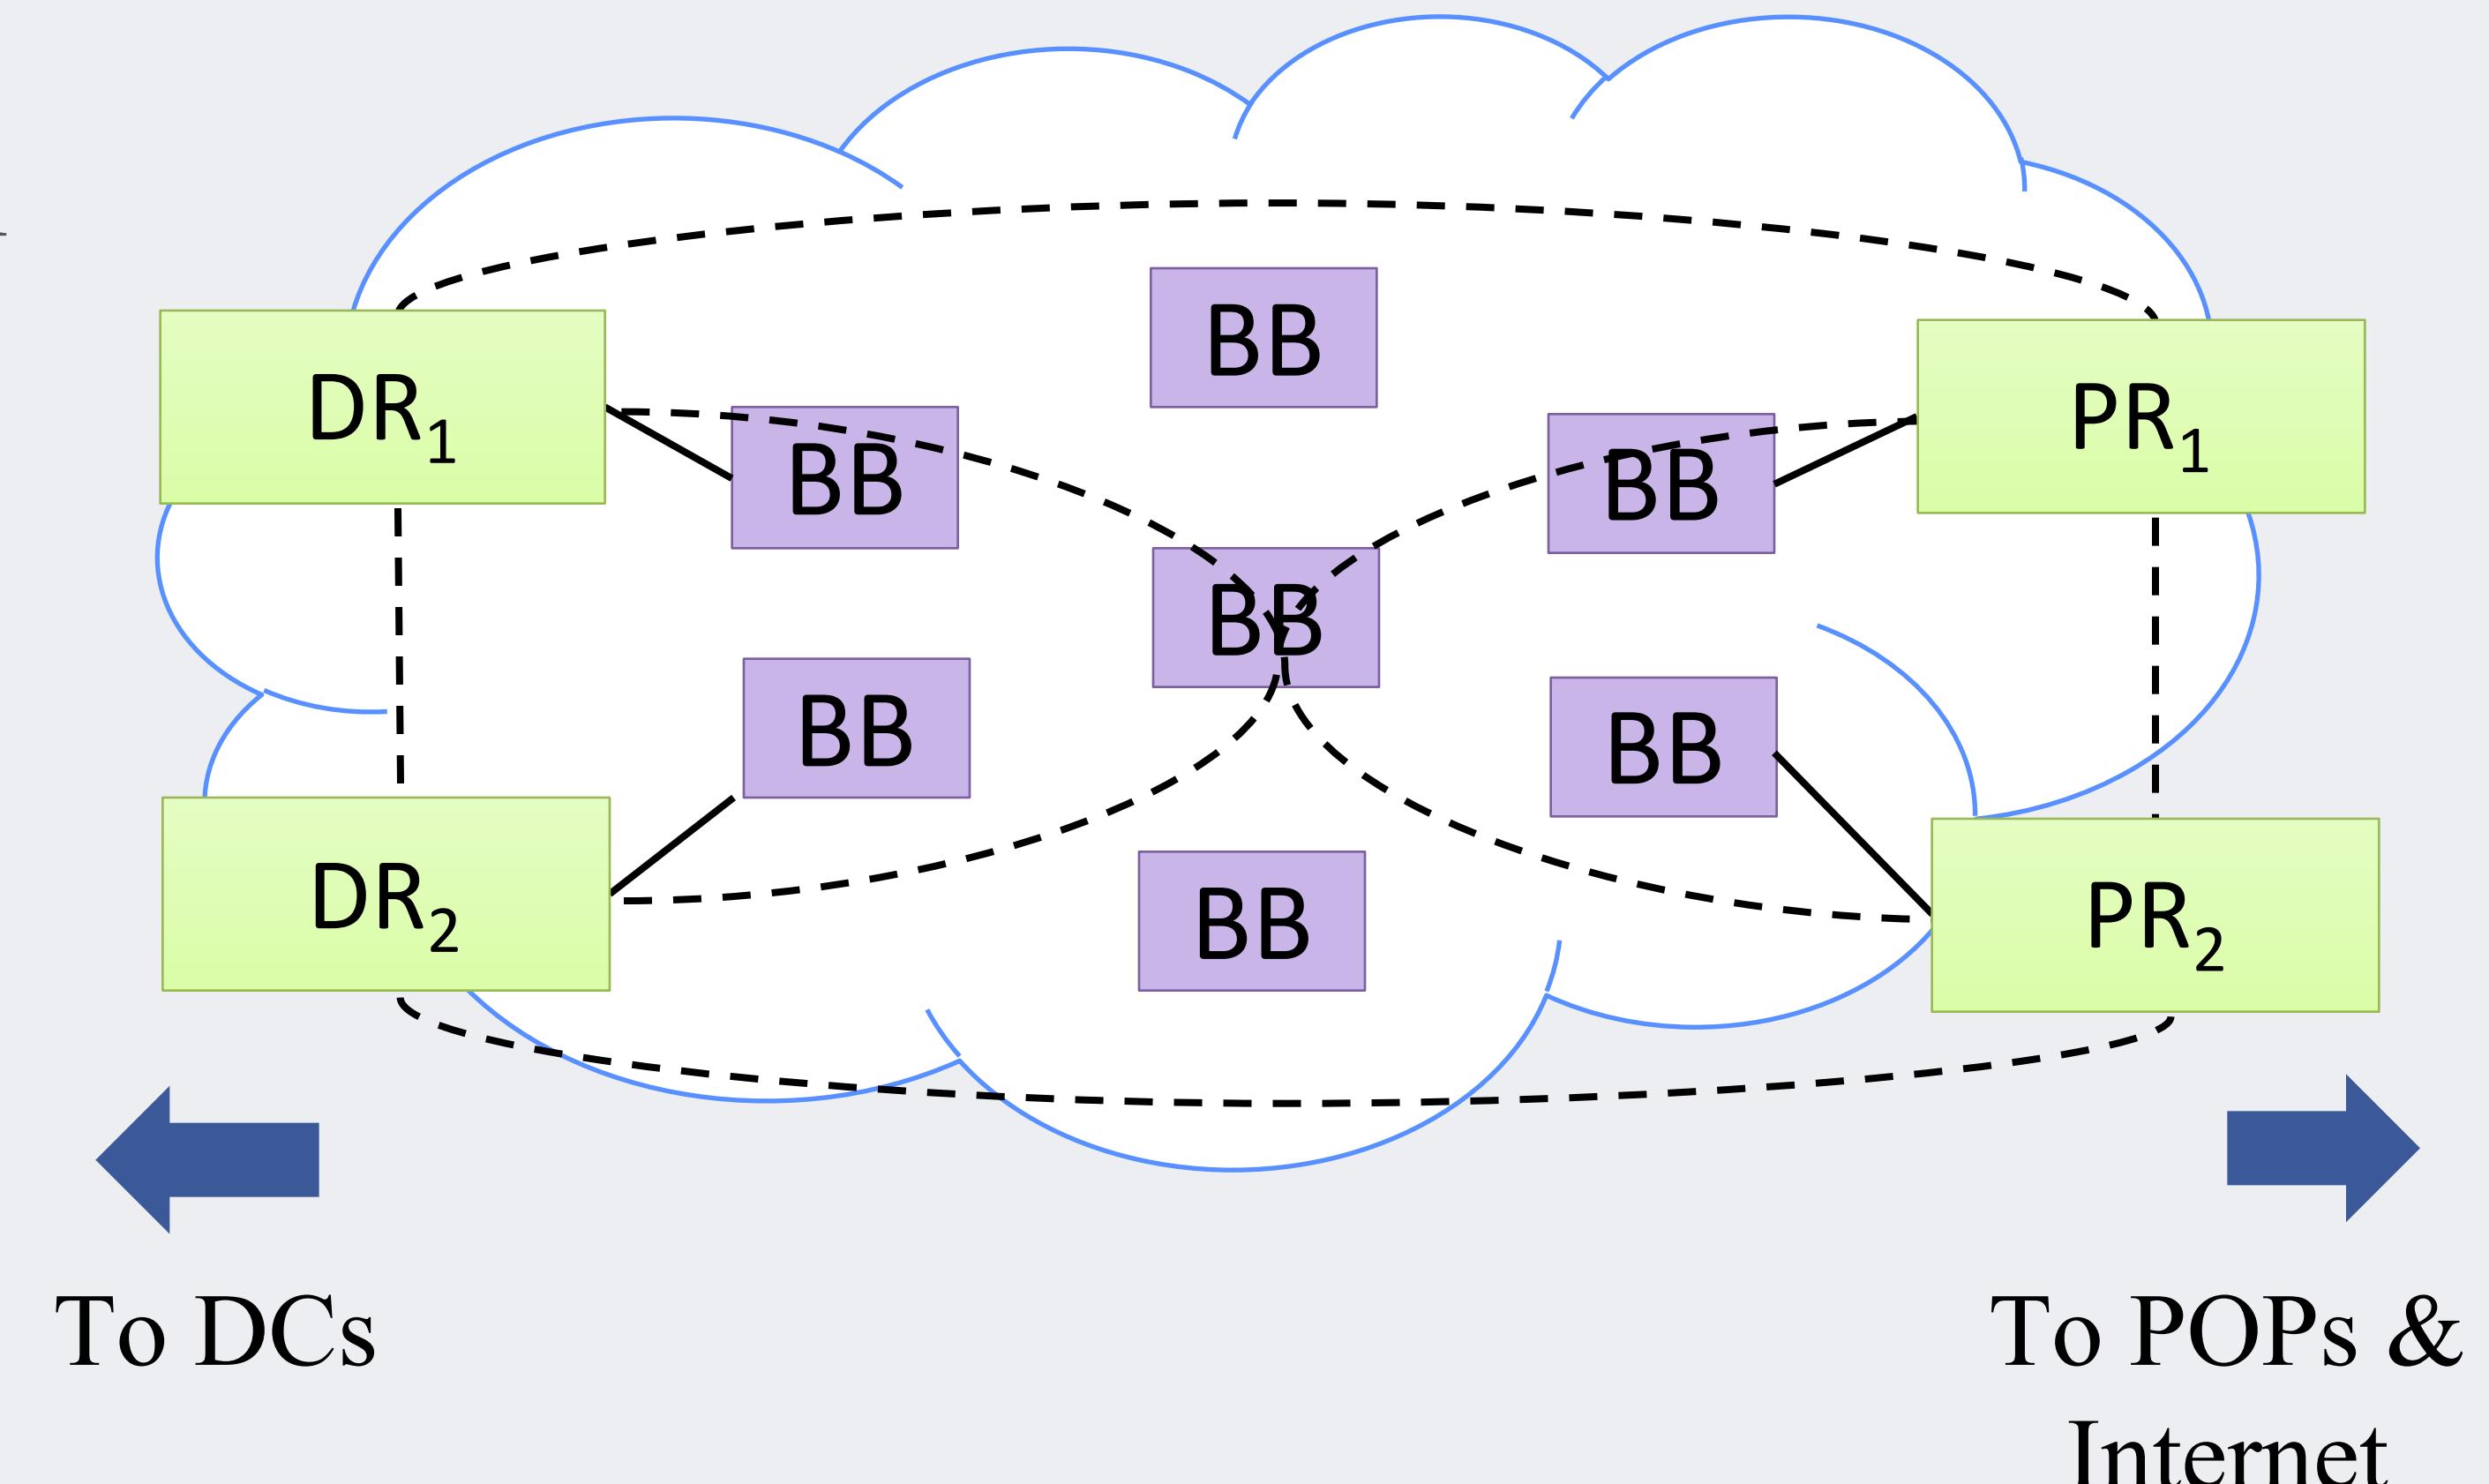
Questions?

- robotron@fb.com
- Poster session on Thursday

Overview of Facebook's Network

Backbone: Interconnecting POPs/DCs

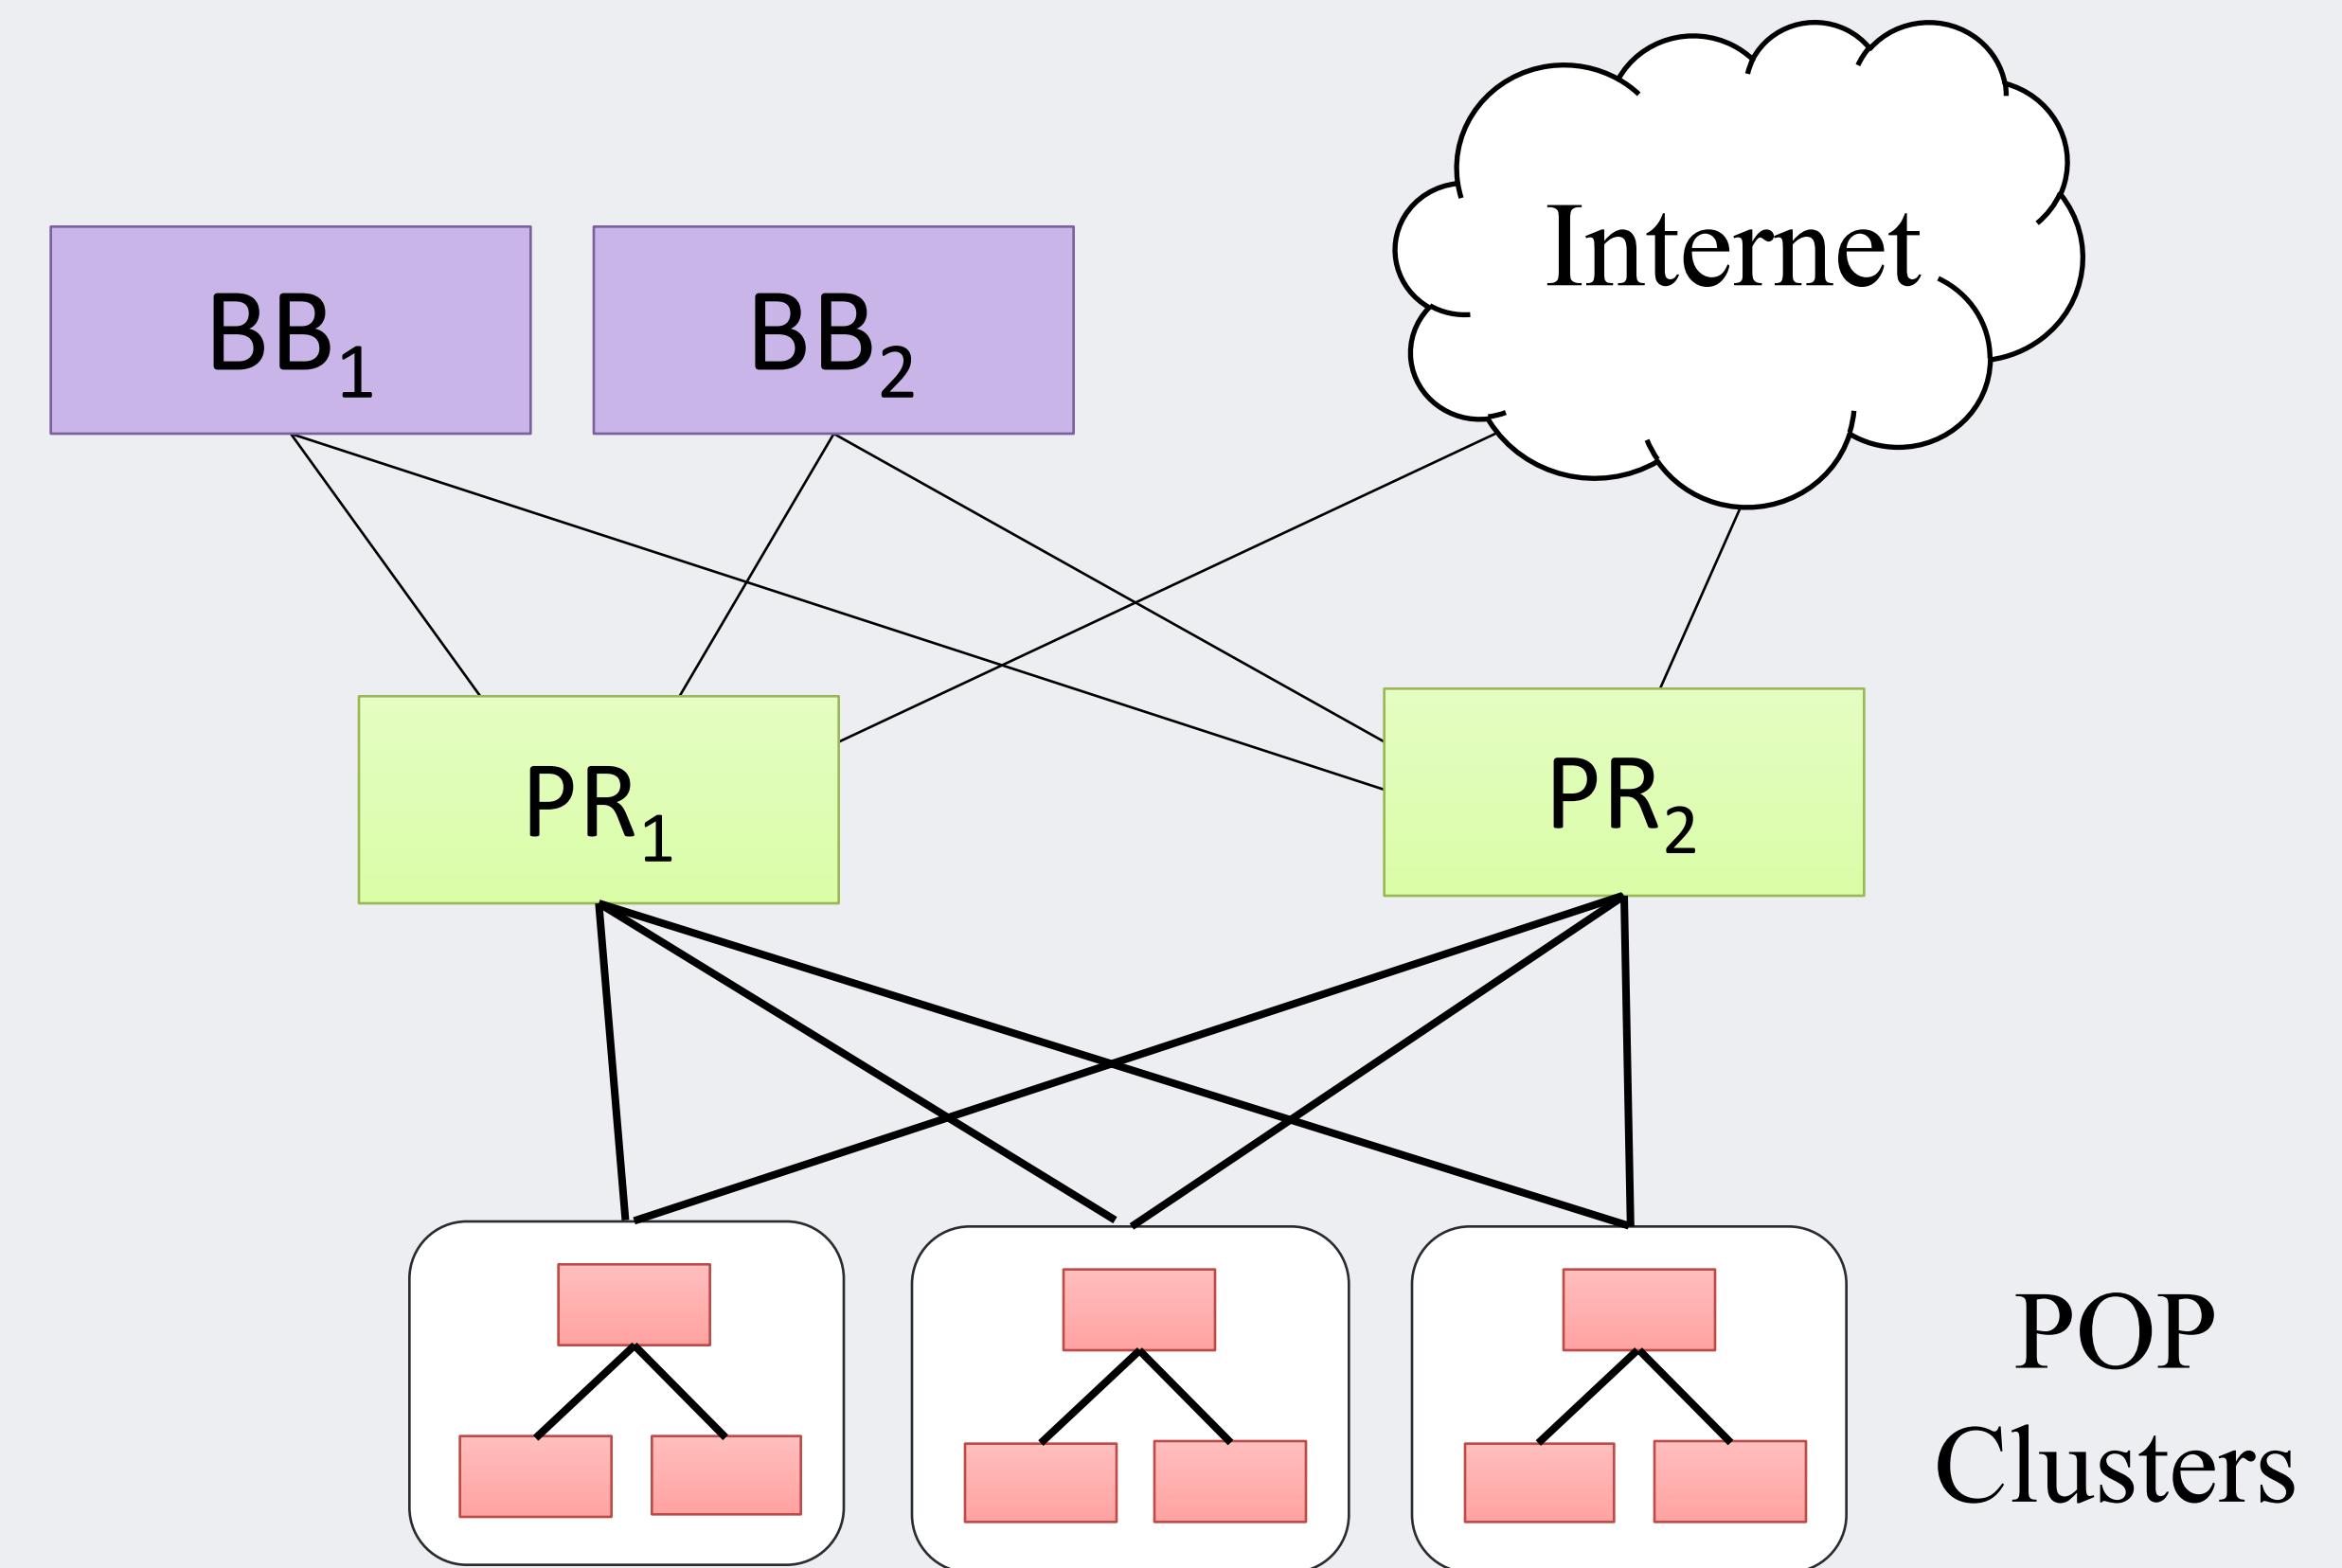
- Irregular, demand-driven topology
- PRs/DRs form an iBGP mesh
- Common tasks:
 - Add/migrate circuits
 - Add/remove BBs/PRs/DRs



Overview of Facebook's Network

Point of Presence (POP)

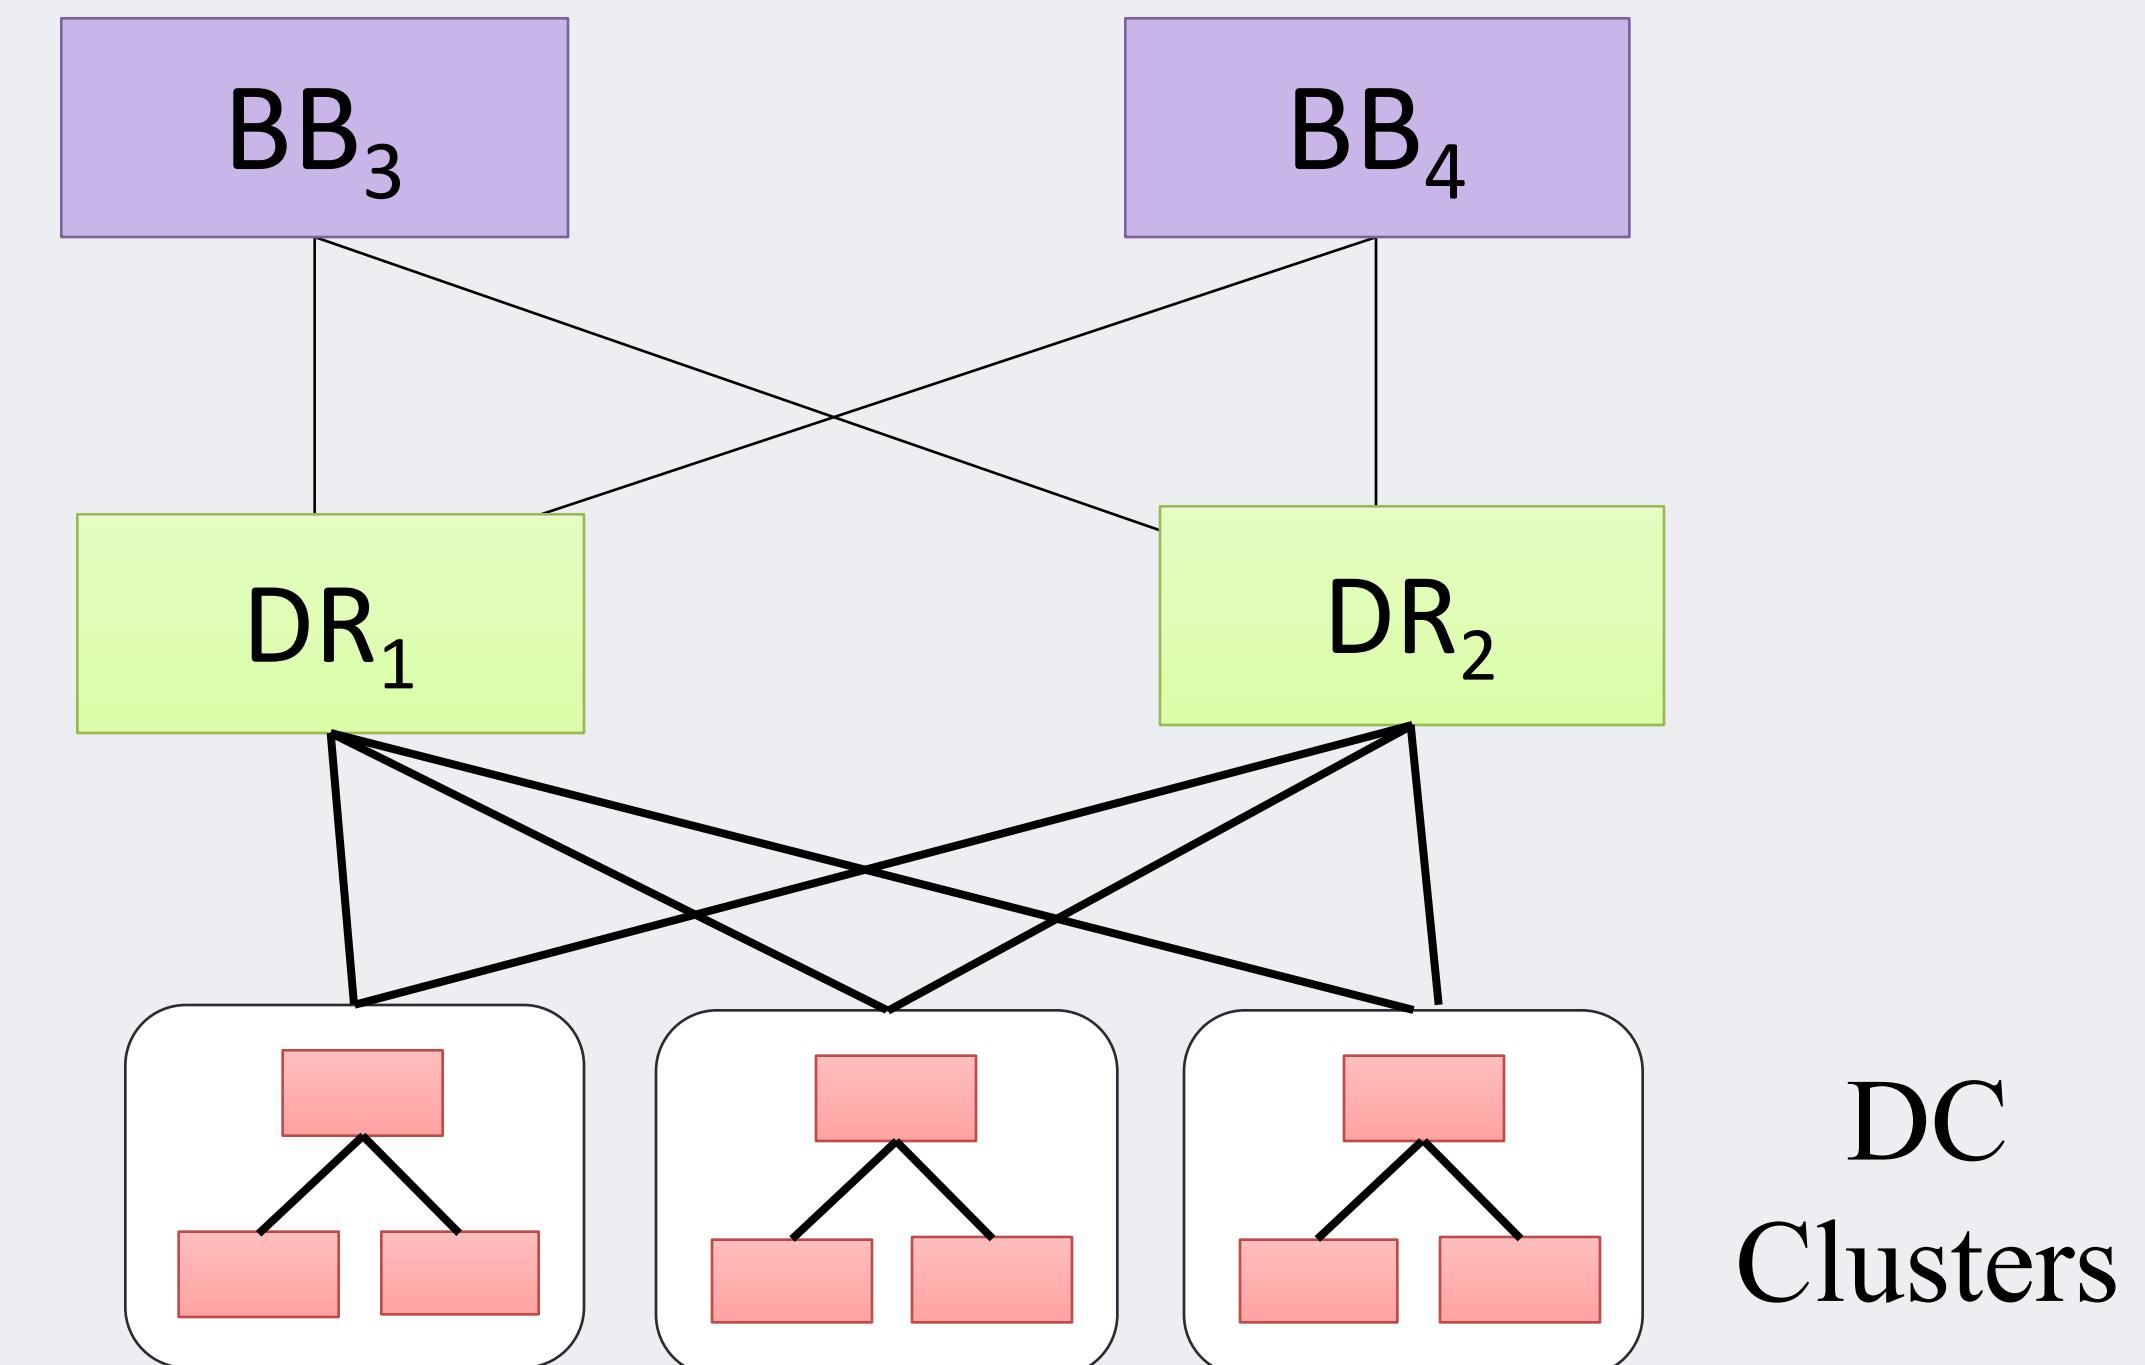
- Standardized topology
- Services: LB (Proxygen), Cache
- Common tasks
 - Build/upgrade a cluster
 - Provisioning new peering circuits



Overview of Facebook's Network

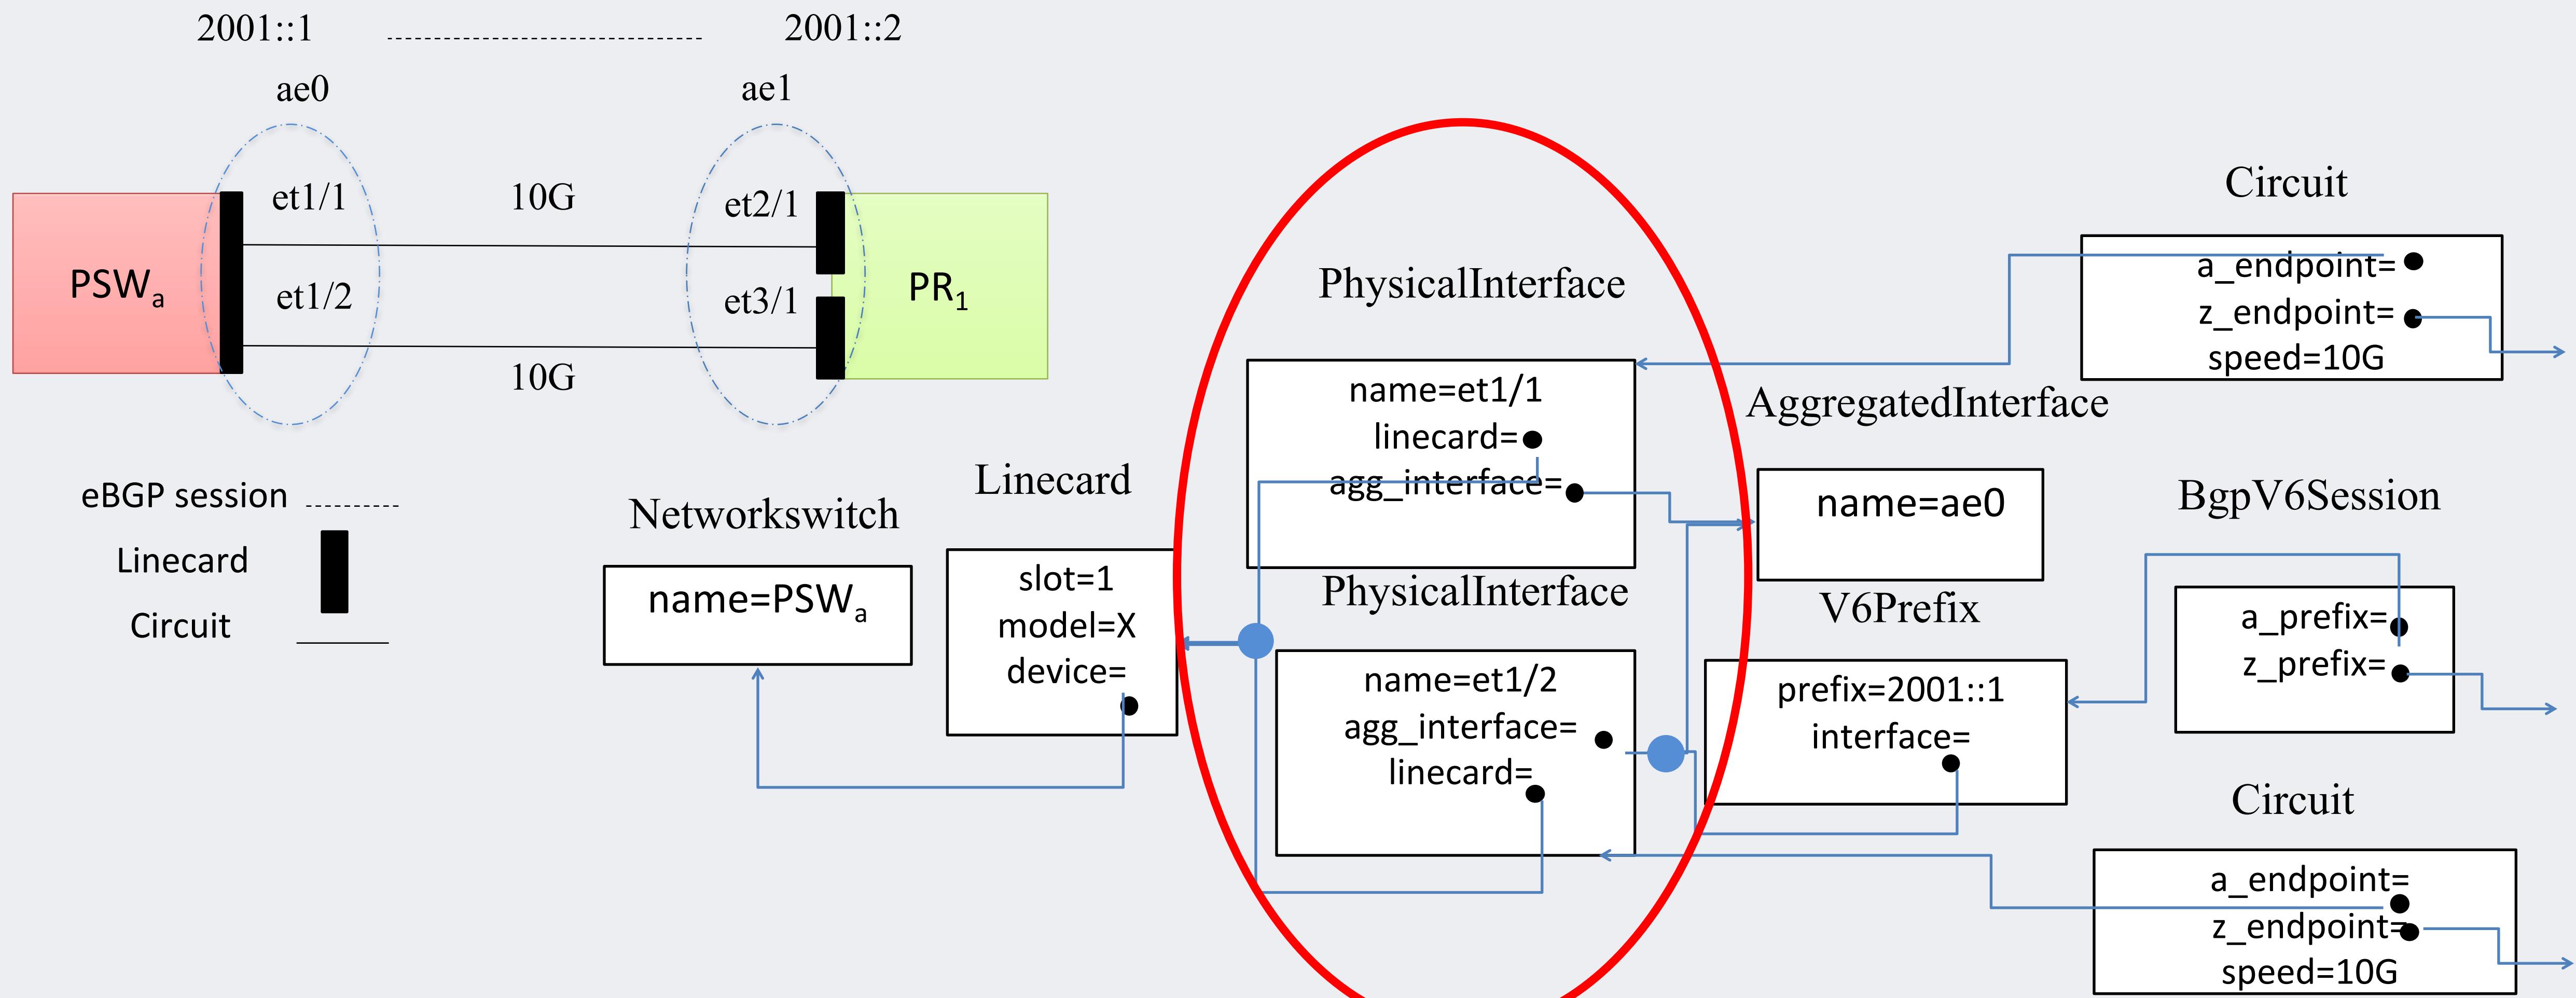
Data Center

- Standardized topology
- Services: Web, Cache (TAO), Database
- Common tasks
 - Build/decomm a cluster
 - Cluster capacity upgrade

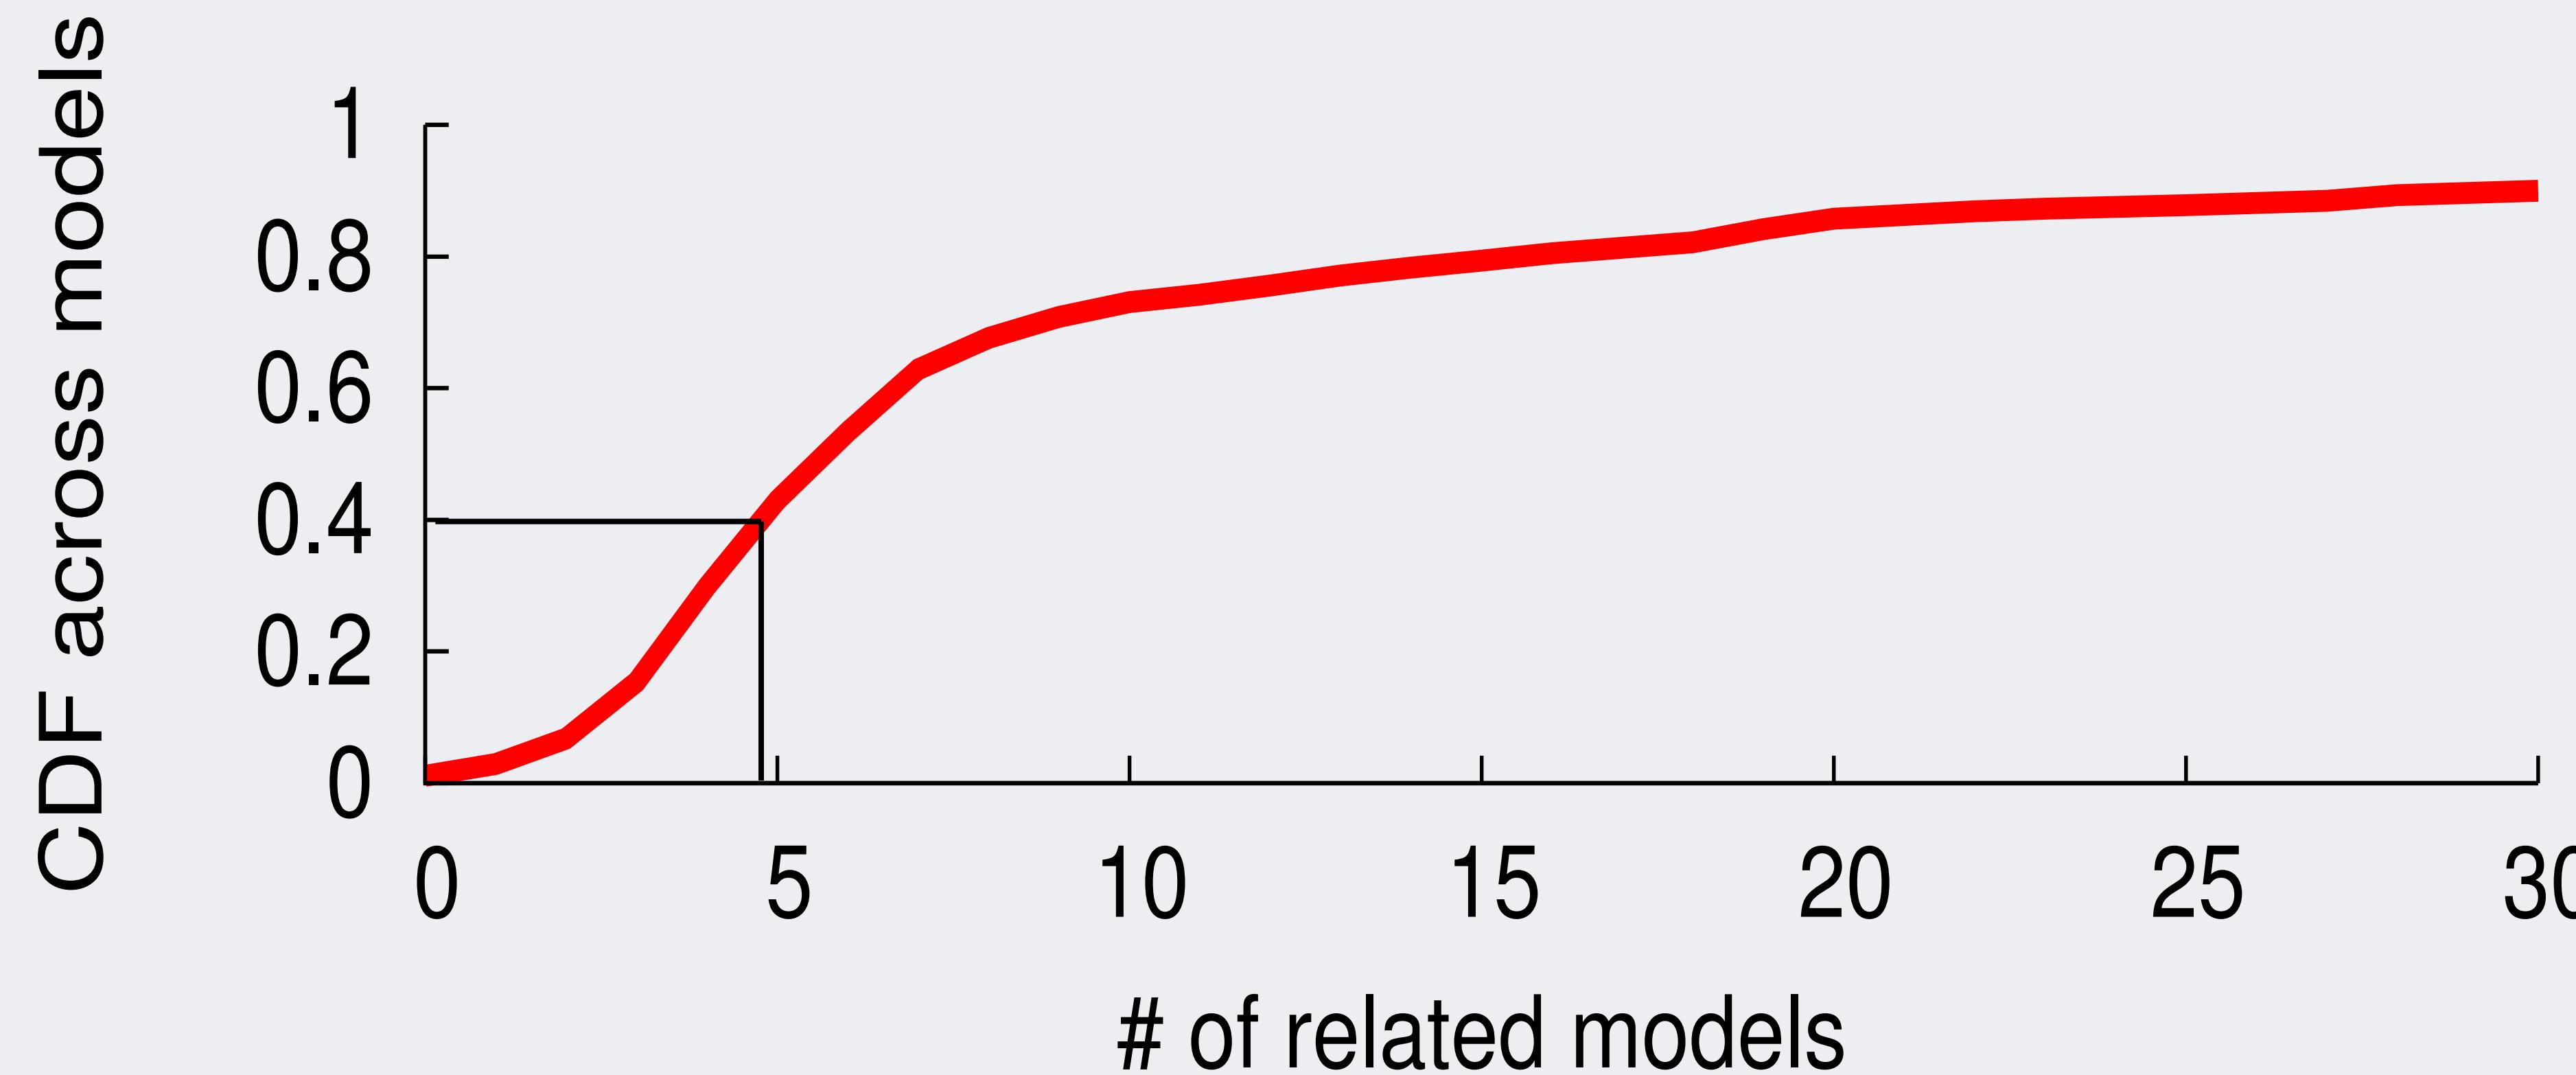


FBNet: Modeling the Network

Object, Value, and Relationship



Dependencies between FBNet models



Experience: Fallback is needed

Problem Scenario: manual changes to devices

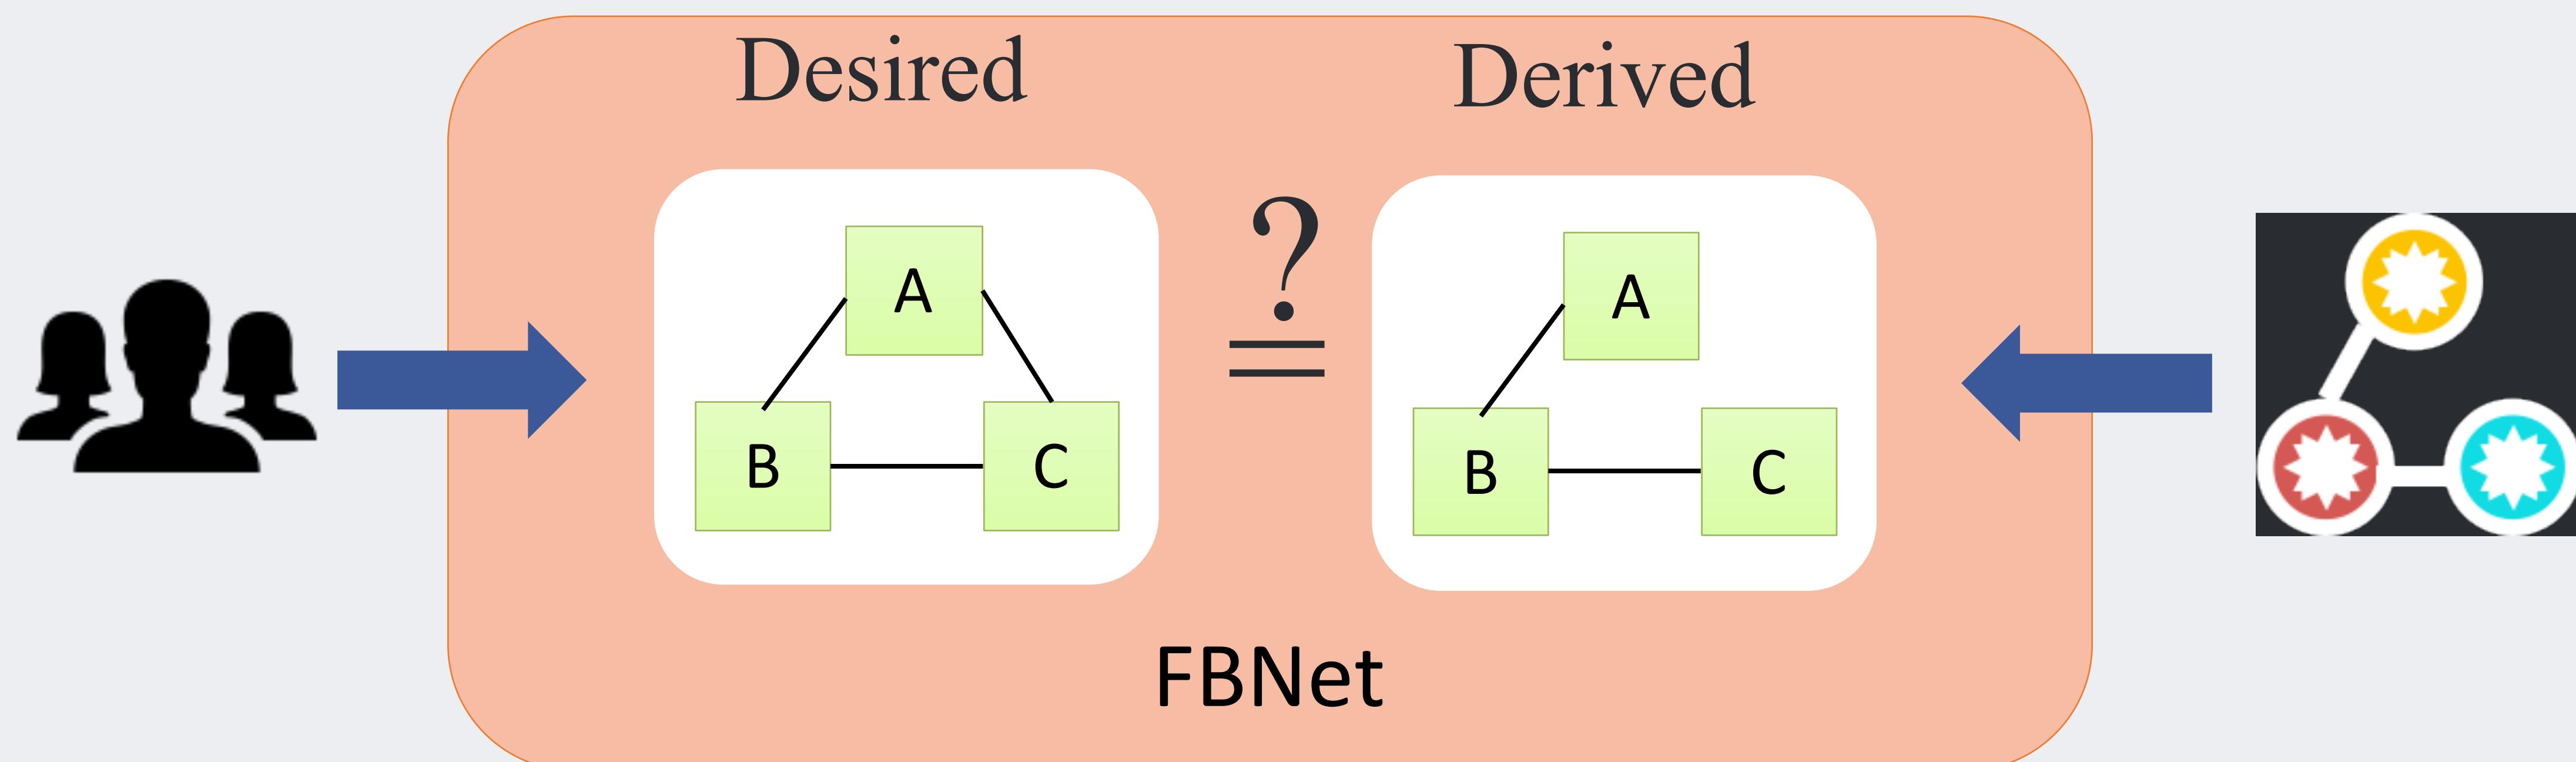
- Manual config changes on devices are error-prone
- Ideal: All changes made through Robotron
- Reality: Robotron has latency, bugs and missing features. Quick fixes needed upon emergency
- Alternatives to discourage manual changes:
 - Config monitoring
 - Automatic config override after emergency window

Related Work

- Bottom-up config analysis:
[Benson11, Sung09, Kim11, ...]
- Abstraction-driven design and config generation:
 - Top down config optimization: [Condor, Sun13]
 - Centralized platform for network management: [Onix, Statesman]
 - Template based config generation: [Enck09]
 - Config modeling: [OpenConfig, DMTF]

FBNet: Modeling the Network

Desired versus Derived



Deployment

Device configs → Devices

- New device: full config replacement
- Existing devices: Incremental “Live” updates
 - Dryrun, Atomic, Phased, etc

Monitoring

Is the network healthy?

- Passive monitoring
- Active monitoring
- Config monitoring

facebook